

SUPERVISED AND UNSUPERVISED MACHINE LEARNING TECHNIQUES
FOR TEXT DOCUMENT CATEGORIZATION

by

Arzucan Özgür

B.S. in Computer Engineering, Boğaziçi University, 2002

Submitted to the Institute for Graduate Studies in
Science and Engineering in partial fulfillment of
the requirements for the degree of
Master of Science

Graduate Program in Computer Engineering
Boğaziçi University

2004

SUPERVISED AND UNSUPERVISED MACHINE LEARNING TECHNIQUES
FOR TEXT DOCUMENT CATEGORIZATION

APPROVED BY:

Prof. Ethem Alpaydın
(Thesis Supervisor)

Assist. Prof. Tunga Güngör

Assist. Prof. Şule Gündüz Öğüdücü

DATE OF APPROVAL: 23.07.2004

ACKNOWLEDGEMENTS

I would like to thank Prof. Ethem Alpaydın, for his contribution to my education, for his motivation and showing me the directions to follow in this study, and especially for giving his time and helping me constantly.

I thank Assist. Prof. Tunga Güngör and Assist. Prof. Şule Gündüz Öğüdücü, for participating in my thesis jury and giving me feedback.

I thank the PILAB members and all my other friends and teachers in the department for their support. Special thanks to Rabun Koşar for his practical helps and to Kubilay Atası especially for his courage and help during my thesis defense. I am thankful to İtir Barutçuoğlu for encouraging me, for her help about SVM and for all the other things I do not mention. I thank my friends Canan Pembe, Deniz Gayde, Arzu Gencer, Sezen Yüzbaş, Murat Aydın, Can Tekeli, Gül Çalık, Tülay Nuri, Hatice Sarmemet, Barış Koçak, İlke Karşı, and my cousin Selçuk Altay for their support. I am thankful to İlker Türkmen for being my best friend and much more and to Nimet Türkmen and Gürer Piker for always making me feel better.

Finally, I want to express my gratefulness to my mother and father for their endless love, support, encouragement, patience and self-sacrifice. I am thankful to my brother and colleague Mürvet Özgür for picking me late from school, for his ideas, motivation, and every thing else I can not express.

ABSTRACT

SUPERVISED AND UNSUPERVISED MACHINE LEARNING TECHNIQUES FOR TEXT DOCUMENT CATEGORIZATION

Automatic organization of documents has become an important research issue since the explosion of digital and online text information. There are mainly two machine learning approaches to enhance this task: supervised approach, where pre-defined category labels are assigned to documents based on the likelihood suggested by a training set of labelled documents; and unsupervised approach, where there is no need for human intervention or labelled documents at any point in the whole process.

In this study we compare and evaluate the performance of the leading supervised and unsupervised techniques for document organization by using different standard performance measures and five standard document corpora. We conclude that among the unsupervised techniques we have evaluated, k -means and bisecting k -means perform the best in terms of time complexity and the quality of the clusters produced. On the other hand, among the supervised techniques support vector machines achieve the highest performance while naive Bayes performs the worst. Finally, we compare the supervised and the unsupervised techniques in terms of the quality of the clusters they produce. In contrast to our expectations, we observe that although k -means and bisecting k -means are unsupervised they produce clusters of higher quality than the naive Bayes supervised technique. Furthermore, the overall similarities of the clustering solutions obtained by the unsupervised techniques are higher than the supervised ones. We discuss that the reason may be due to the outliers in the training set and we propose to use unsupervised techniques to enhance the task of pre-defining the categories and labelling the documents in the training set.

ÖZET

BELGE SINIFLANDIRMA İÇİN GÖZETİMLİ VE GÖZETİMSİZ ÖĞRENME ALGORİTMALARI

Bilgisayar ve elektronik teknolojilerinin gelişmesi, İnternet ve Web'in yaygınlaşmasıyla elektronik belgelerin miktarı her geçen gün artmaktadır. Bu elektronik veritabanlarında ilgili verilere daha hızlı, kolay, ve doğru bir şekilde erişebilmek için belgelerin otomatik olarak sınıflandırılması önem kazanmıştır. Otomatik sınıflandırma için temelde iki yapay öğrenme yaklaşımı vardır: gözetimli öğrenme ve gözetimsiz öğrenme. Gözetimli öğrenmede, önceden sınıfların bilinmesi ve bu sınıflara ait belgelerden oluşan bir öğrenme kümesi gerekir. Gözetimsiz öğrenmede ise sınıfların önceden bilinmesine ve herhangi bir aşamada insan yardımına ihtiyaç yoktur.

Bu çalışmada otomatik belge sınıflandırma için gözetimli ve gözetimsiz temel yöntemleri ele alıyoruz. Bu temel yöntemlerin beş standart veritabanı üzerindeki başarımlarını farklı kıstaslara dayanarak inceliyor, gözetimli ve gözetimsiz öğrenme yaklaşımlarını birbiriyle kıyaslıyoruz. Bu çalışma sonucunda gözetimsiz yöntemler içinde k -means ve bisecting k -means'in belge öbeklenmesi için daha elverişli olduğunu gördük. Gözetimli yöntemler arasında en iyi başarımları destek vektör makinaları elde ediyor. Gözetimsiz yöntemler olmalarına rağmen k -means ve bisecting k -means gözetimli bir yöntem olan naive Bayes'den daha kaliteli öbekler oluşturuyor. Gözetimsiz yöntemlerin oluşturduğu öbeklerin toplam benzerliği gözetimli yöntemlerininkinden genellikle daha yüksek. Bu sonuç öğrenme kümesinde hatalı bazı belgelerin olmasından kaynaklanıyor olabilir. Bu nedenle sınıfların belirlenmesi ve öğrenme kümesinin oluşturulması aşamasında gözetimsiz yöntemlerden faydalanılmasını öneriyoruz.

TABLE OF CONTENTS

ACKNOWLEDGEMENTS	iii
ABSTRACT	iv
ÖZET	v
LIST OF FIGURES	viii
LIST OF TABLES	xiv
LIST OF SYMBOLS/ABBREVIATIONS	xvi
1. INTRODUCTION	1
1.1. Document Classification	1
1.2. Document Clustering	3
1.3. Motivation	4
1.4. Thesis Organization	5
2. DOCUMENT PREPROCESSING AND REPRESENTATION	7
2.1. Parsing the Documents and Case-folding	8
2.2. Removing Stopwords	8
2.3. Stemming	9
2.4. Term Weighting	11
2.4.1. Boolean Weighting	11
2.4.2. Term Frequency Weighting	12
2.4.3. Term Frequency \times Inverse Document Frequency Weighting	12
2.4.4. TF \times IDF Weighting with Length Normalization	13
2.5. Dimensionality Reduction	13
2.5.1. Information Gain (IG)	13
2.5.2. Mutual Information (MI)	14
2.5.3. Chi-Square Statistic	15
2.5.4. Term Strength (TS)	16
2.5.5. Document Frequency (DF) Thresholding	17
2.6. Document Similarity Measure	18
3. UNSUPERVISED TECHNIQUES FOR DOCUMENT CLUSTERING	19
3.1. Partitional Clustering Techniques	19

3.1.1.	K-Means Clustering	19
3.1.2.	Bisecting K-Means	20
3.2.	Hierarchical Clustering Techniques	21
3.2.1.	Divisive Hierarchical Clustering	21
3.2.2.	Agglomerative Hierarchical Clustering	22
3.2.2.1.	Single-link	22
3.2.2.2.	Complete-link	23
3.2.2.3.	Average-link	23
4.	SUPERVISED TECHNIQUES FOR DOCUMENT CLASSIFICATION . . .	24
4.1.	K Nearest Neighbor Classification	24
4.2.	Naive Bayes Approach	25
4.2.1.	Multinomial Model	26
4.2.2.	Multivariate Bernoulli Model	27
4.3.	Support Vector Machines	28
5.	EXPERIMENT RESULTS	30
5.1.	Document Data Sets	30
5.2.	Evaluation of the Clustering Techniques	33
5.2.1.	Evaluation Metrics	33
5.2.1.1.	Overall Similarity	34
5.2.1.2.	Purity	34
5.2.1.3.	Entropy	35
5.2.1.4.	F-measure	35
5.2.2.	Results and Discussion	36
5.3.	Evaluation of the Classification Techniques	53
5.3.1.	Evaluation Metrics	53
5.3.1.1.	Micro-averaged F_1 -measure	55
5.3.1.2.	Macro-averaged F_1 -measure	55
5.3.2.	Results and Discussion	56
6.	COMPARISON OF THE SUPERVISED AND UNSUPERVISED METHODS	62
7.	CONCLUSIONS AND FUTURE WORK	66
	APPENDIX A: STATISTICS OF THE CLUSTERING ALGORITHMS	70
	REFERENCES	94

LIST OF FIGURES

Figure 1.1.	ML approaches used in this thesis for document categorization . . .	6
Figure 2.1.	Portion of the stopwords list used	9
Figure 2.2.	Sample of words and their corresponding stems found by Porter's Stemming Algorithm	10
Figure 3.1.	Inter-cluster similarity defined by single-link, complete-link, and average-link	22
Figure 4.1.	Support vector machines find the hyperplane h that separates pos- itive and negative training examples with maximum margin. Sup- port vectors are marked with circles	28
Figure 5.1.	Class names of Reuters-21578 data set	32
Figure 5.2.	Class names of the Wap data set	33
Figure 5.3.	Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete- link, average-link and divisive clustering algorithms over Classic3 data set	40
Figure 5.4.	Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete- link, average-link and divisive clustering algorithms over Wap data set	41

Figure 5.5.	Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Reuters-21578 data set	42
Figure 5.6.	Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over LA1 data set	43
Figure 5.7.	Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Hitech data set	44
Figure 5.8.	Term distributions of Classic3, Wap, Reuters-21578, LA1, and Hitech data sets	45
Figure 5.9.	Per cent of documents in Classic3, Wap, Reuters-21578, LA1, and Hitech document collections whose nearest neighbor belongs to a different class	46
Figure 5.10.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by k -means for Classic3 data set	48
Figure 5.11.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Classic3 data set	49

Figure 5.12.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Classic3 data set	50
Figure 5.13.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Classic3 data set	51
Figure 5.14.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Classic3 data set	52
Figure 5.15.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Classic3 data set	53
Figure 5.16.	Effect of the degree of the polynomial kernel of SVM to the F_1 -measure scores	59
Figure 5.17.	Effect of the variance of the RBF kernel of SVM to the F_1 -measure scores	60
Figure A.1.	The performance, cluster-class distribution, and most descriptive 5 features of the clustering solution obtained by k -means for Hitech	70
Figure A.2.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Hitech data set	71
Figure A.3.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Hitech data set	72

Figure A.4.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Hitech data set	73
Figure A.5.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Hitech data set	74
Figure A.6.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Hitech data set	75
Figure A.7.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by k -means for LA1 data set	76
Figure A.8.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for LA1 data set	77
Figure A.9.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for LA1 data set	78
Figure A.10.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for LA1 data set	79
Figure A.11.	The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for LA1 data set	80

Figure A.12. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for LA1 data set	81
Figure A.13. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by k -means for Reuters-21578 data set	82
Figure A.14. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Reuters-21578 data set	83
Figure A.15. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Reuters-21578 data set	84
Figure A.16. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Reuters-21578 data set	85
Figure A.17. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Reuters-21578 data set	86
Figure A.18. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Reuters-21578 data set	87
Figure A.19. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by k -means for Wap data set	88

Figure A.20. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Wap data set	89
Figure A.21. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Wap data set	90
Figure A.22. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Wap data set	91
Figure A.23. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Wap data set	92
Figure A.24. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Wap data set	93

LIST OF TABLES

Table 5.1.	Summary description of document sets	30
Table 5.2.	Micro-averaged F_1 -measure results for SVM with polynomial kernel of degrees 1, 2, 3, 4, 5, 6, and 7	56
Table 5.3.	Macro-averaged F_1 -measure results for SVM with polynomial kernel of degrees 1, 2, 3, 4, 5, 6, and 7	57
Table 5.4.	Micro-averaged F_1 -measure results for SVM with RBF kernel with $\gamma = 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, 1.4$	57
Table 5.5.	Macro-averaged F_1 -measure results for SVM with RBF kernel with $\gamma = 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, 1.4$	57
Table 5.6.	Micro-averaged F_1 -measure results for the supervised techniques .	61
Table 5.7.	Macro-averaged F_1 -measure results for the supervised techniques .	61
Table 6.1.	Quality of the clusters ($k=3$) obtained by the unsupervised and the supervised techniques for Classic3	62
Table 6.2.	Quality of the clusters ($k=6$) obtained by the unsupervised and the supervised techniques for Hitech	63
Table 6.3.	Quality of the clusters ($k=6$) obtained by the unsupervised and the supervised techniques for LA1	63
Table 6.4.	Quality of the clusters ($k=90$) obtained by the unsupervised and the supervised techniques for Reuters-21578	64

Table 6.5.	Quality of the clusters ($k=20$) obtained by the unsupervised and the supervised techniques for Wap	64
------------	--	----

LIST OF SYMBOLS/ABBREVIATIONS

A_{ct}	Number of documents that belong to category c and contain term t
B_{ct}	Number of documents that do not belong to category c but contain term t
B_{it}	1 if term t appears in document d_i , 0 otherwise
C	Cluster
$ C $	Number of documents in cluster C
c	Category
\mathbf{c}	Centroid vector
C_{ct}	Number of documents that belong to category c but do not contain term t
C_i	Cluster i
c_i	Category i
\mathbf{c}_i	Feature vector of centroid i
C_j	Cluster j
c_j	Category j
D	Set of documents
d	Degree of the polynomial kernel of the support vector machines
\mathbf{d}	A document vector
\mathbf{d}'	A document vector
D_{ct}	Number of documents that do not belong to category c and do not contain term t
d_i	Document i
\mathbf{d}_i	Feature vector of document i
$ d_i $	Number of terms in document i
d_j	Document j
E_j	Entropy of cluster j
f	Target function
f'	Classifier function

f_{ij}	Observed frequency of the cell in row i and column j
\hat{f}_{ij}	Expected frequency of the cell in row i and column j
FN_i	The number of documents that are not assigned to category i by the classifier but which actually belong to category i
FP_i	The number of documents that do not belong to category i but are assigned to category i by the classifier incorrectly
I_j	Internal cluster similarity of cluster j
k	Requested number of clusters
M	Number of categories
N	Number of documents in the document corpus
N_i	Number of documents in the corpus where term i appears
n_i	Size of cluster i
n_{ij}	Number of documents with class label i in cluster j
N_j	Number of documents in the corpus where term j appears
n_j	Size of cluster j
P_j	Purity of cluster j
S	Set of predefined categories
T	Term space
$ T $	Number of terms in the document collection
t	term
t_i	The i^{th} term occurring in document d_i
t_k	The k^{th} term occurring in document d_i
tf_i	Raw frequency of term i in the specified document
tf_j	Raw frequency of term j in the specified document
TP_i	The number of documents assigned correctly to category i
w_i	Weight of term i in document vector \mathbf{d}
v_t	The t^{th} term in the term space $ T $
\mathbf{x}	Test document vector
β	Degree of importance in the range $[0, +\infty]$ given to π and ρ in F_β -Measure
χ^2	Chi-square distribution

γ	Variance of the RBF kernel of SVM
π	Precision
π_i	Precision for category i
ρ	Recall
ρ_i	Recall for category i
DF	Document Frequency
FN	False Negatives
FP	False Positives
HTML	Hyper Text Markup Language
IDF	Inverse Document Frequency
IG	Information Gain
IR	Information Retrieval
k -NN	k -nearest neighbor
LLSF	Linear Least Squares Fit
MI	Mutual Information
ModApte	Modified Apte
MPT	Most Prevailing Topic, i.e, the topic that has the greatest number of documents in the given cluster
NB	Naive Bayes
NLP	Natural Language Processing
NNet	Neural Networks
RBF	Radial Basis Function
SGML	Standard Generalized Markup Language
SVM	Support Vector Machines
TF	Term Frequency
TP	True Positives
TREC	Text Retrieval Conference
TS	Term Strength
WWW	World Wide Web

1. INTRODUCTION

The amount of electronic text information available such as electronic publications, digital libraries, electronic books, email messages, news articles, and Web pages is increasing rapidly. However, as the volume of online text information increases the challenge of extracting relevant knowledge increases as well. The need for tools that enhance people find, filter, and manage these resources has grown. Thus, automatic organization of text document collections has become an important research issue. A number of machine learning techniques have been proposed to enhance automatic organization of text data. These techniques can be grouped in two main categories as supervised (document classification) and unsupervised (document clustering).

1.1. Document Classification

Text classification, also known as text categorization or topic spotting, is a supervised learning task, where pre-defined category labels are assigned to documents based on the likelihood suggested by a training set of labelled documents. Until this machine learning approach to text categorization, the most popular approach was knowledge engineering. In knowledge engineering, expert knowledge is used to define manually a set of rules on how to classify documents under the pre-defined categories. It is discussed in [1] that the machine learning approach to document classification leads to time and cost savings in terms of expert manpower without loss in accuracy. In the problem of text classification we have a set D of documents and a set S of pre-defined categories. The aim is to assign a boolean value to each $\langle d_i, c_j \rangle$ pair, where $d_i \in D$ and $c_j \in S$. A value of *true* assigned to $\langle d_i, c_j \rangle$ stands for the decision of assigning document d_i to category c_j . Analogously, value of *false* assigned to $\langle d_i, c_j \rangle$ stands for the decision of not assigning document d_i to category c_j . To state more formally, the task is to approximate the unknown target function $f : D \times S \rightarrow \{true, false\}$, that describes the way the documents should actually be classified, by the classifier function $f' : D \times S \rightarrow \{true, false\}$ such that number of decisions of f and f' that do not coincide is minimized.

Many learning algorithms such as k -nearest neighbor (k -NN) [2, 3, 4], Support Vector Machines (SVM) [5], neural networks (NNet) [6, 7], linear least squares fit (LLSF) [8], and naive Bayes (NB) [9, 10] have been applied to text classification. A comparison of these techniques is presented by Yang and Liu in [8]. They conclude that all these techniques perform comparably when each category contains over 300 documents. However, when the number of positive training documents per category is less than 10, SVM, k -NN, and LLSF outperform significantly NNet and NB.

Text categorization has many interesting application areas such as document organization, text filtering and hierarchical categorization of Web pages. Document organization is the task of structuring documents of a corporate document base into folders, which may be hierarchical or flat. For instance, advertisements incoming to a newspaper office may be classified into categories such as Cars, Real Estate, Computers, and so on before publication. Similarly, Larkey has used document categorization to organize patents into categories to enhance their search in [11]. Text filtering is the process of classifying a dynamic collection of text documents into two disjoint categories as relevant and irrelevant. An example is a newsfeed system where news articles incoming to a newspaper from a news agency such as Reuters are filtered [1]. If it is a sports newspaper, delivery of news articles not related to sport are blocked. Similarly, an email filter may be trained to classify incoming messages as spam or not spam, and block the delivery of spam messages [12, 13]. Hierarchically categorizing web pages or sites facilitates searching and browsing operations. Rather than posing a generic query to a general purpose search engine, it is easier and more effective to first navigate in the hierarchy of categories and restrict the search to the particular categories of interest. To classify documents hierarchically, generally the classification problem is subdivided into smaller classification problems. Hierarchical classification of documents is addressed by Koller and Sahami in [10] and by Dumais and Chen in [14].

1.2. Document Clustering

Unlike document classification, document clustering is an unsupervised learning task, which does not require pre-defined categories and labelled documents. The aim of text clustering is to group text documents such that intra-group similarities are high and inter-group similarities are low. Document clustering has many application areas. In Information Retrieval (IR), it has been used to improve precision and recall, and as an efficient method to find similar documents. More recently, document clustering has been used in automatically generating hierarchical groupings of documents by Koller and Sahami in [10] and in document browsing to organize the results returned by a search engine by Zamir *et al.* in [15].

Machine learning algorithms for clustering can be categorized into two main groups as hierarchical clustering algorithms and partitional clustering algorithms [16]. Hierarchical algorithms produce nested partitions of data by splitting (divisive approach) or merging (agglomerative approach) clusters based on the similarity among them. Divisive algorithms start with one cluster of all data points and at each iteration split the most appropriate cluster until a stopping criterion such as a requested number k of clusters is achieved. Conversely, in agglomerative algorithms each item starts as an individual cluster and at each step, the most similar pair of clusters are merged. Agglomerative hierarchical clustering algorithms can be further categorized as single-link, complete-link, and average-link according to the way they define cluster similarity. While agglomerative hierarchical clustering is a commonly used hierarchical approach for document clustering, divisive approach has not been studied much as an approach for document clustering. Evaluation of hierarchical clustering algorithms for document data sets is presented by Zhao and Karypis in [17].

Partitional clustering algorithms group the data into un-nested non-overlapping partitions that usually locally optimize a clustering criterion. Popular partitional clustering techniques applied to the domain of text documents are k -means and its variant bisecting k -means. A comparison of agglomerative hierarchical techniques with k -means and bisecting k -means is performed by Steinbach *et al.* in [18] and it has

been shown that average-link algorithm generally performs better than single-link and complete-link algorithms for the document data sets used in the experiments. Next, average-link algorithm is compared with k -means and bisecting k -means and it has been concluded that bisecting k -means performs better than average-link agglomerative hierarchical clustering algorithm and k -means in most cases for the data sets used in the experiments.

1.3. Motivation

Various supervised machine learning techniques have been applied to document classification. However, text classification requires the extra effort to predefine the categories and to assign category labels to the documents in the training set. This can be very tedious in large and dynamic text databases such as the WWW. Another phenomenon that poses a challenge to document categorization is inter-indexer inconsistency discussed by Sebastiani in [1]. This phenomenon states that two human experts may disagree when deciding under which category to categorize a given document. For instance a news story about Bill Clinton and Monika Lewinsky may be classified under the category Politics, or under the category Gossip, or under the both categories, or under neither of the categories depending on the subjective judgement of the human indexer [1].

The dynamic nature of most text databases makes it challenging to pre-define the categories and the subjectivity in assigning documents to categories has lead us to believe that by nature text organization should be an unsupervised task rather than a supervised one. Therefore, we concentrated on text clustering which is an unsupervised task where no human intervention at any point in the whole process and no labelled documents are provided. There are many challenges in using the existing machine learning techniques in the domain of text documents. We can list them as follows:

- Number of documents to be clustered is usually very large.
- The feature space is usually very large.
- It is usually very difficult to determine the number of clusters in advance.

- Some documents may belong to more than one cluster (overlapping clusters).
- Shape of clusters may be arbitrary.
- The process should be online considering the dynamic structure of text databases such as the WWW.

Our aim in this study is to compare and evaluate the performance of the commonly used supervised and unsupervised techniques for text document organization. We propose that unsupervised techniques can be used to give feedback to the human indexers to enhance the task of pre-defining categories and preparing a labelled training set. This study will form the basis for developing a hybrid approach of supervised and unsupervised paradigm to the domain of text documents by also considering the challenges stated above.

1.4. Thesis Organization

The outline of this thesis is as follows. In Chapter 2 we discuss how we preprocess and represent documents so that machine learning algorithms can be applied to them. We overview the unsupervised clustering algorithms and the supervised classification algorithms that we evaluate in Chapter 3 and Chapter 4 respectively. Figure 1.1 displays the taxonomy of the ML techniques, that we evaluated for document clustering and classification in this study. In Chapter 5 we describe the standard document data sets we have used in the experiments, our experimental methodology, evaluation metrics and the results we have obtained. In Chapter 6 we perform a comparative study of the supervised techniques with the unsupervised ones. We conclude and outline future directions of research in Chapter 7.

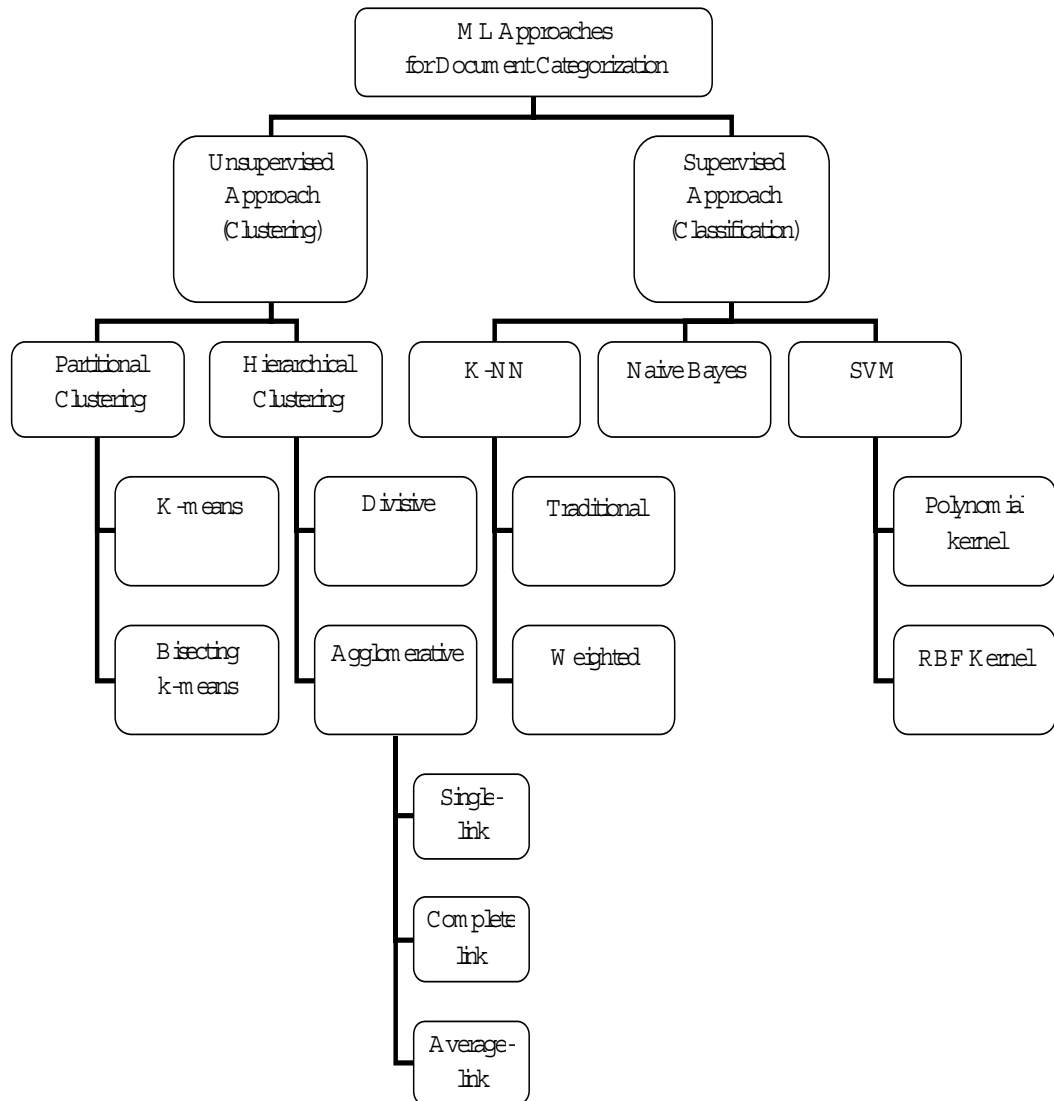


Figure 1.1. ML approaches used in this thesis for document categorization

2. DOCUMENT PREPROCESSING AND REPRESENTATION

In order to cluster or classify text documents by applying machine learning techniques, documents should first be preprocessed. In the preprocessing step, the documents should be transformed into a representation suitable for applying the learning algorithms. The most widely used method for document representation is the vector-space model introduced by Salton *et al.* [19], which we have also decided to employ. In this model, each document is represented as a vector \mathbf{d} . Each dimension in the vector \mathbf{d} stands for a distinct term in the term space of the document collection.

A term in the document collection can stand for a distinct single-word, a stemmed word or a phrase. Phrases consist of multiple words such as “data mining” or “mobile phone” and constitute a different context than when used separately. Phrases can be extracted by using statistical or Natural Language Processing (NLP) techniques. By statistical methods phrases can be extracted by considering the frequently appearing sequences of words in the document collection [20]. A research on extracting phrases by using NLP techniques for text categorization is discussed by Fuernkranz *et al.* [21]. Phrases can also be extracted by manually defining the phrases for a particular domain such as done to filter spam mail in [13]. However, this does not fulfill our requirement to organize documents in generic domains such as the Web.

In vector space representation, defining terms as distinct single words is referred to as “bag of words” representation. Some researchers state that using phrases rather than single words to define terms produce more accurate classification results [20, 21]; whereas others argue that using single words as terms does not produce worse results [22, 23]. As “bag of words” representation is the most frequently used method for defining terms and it is computationally more efficient than the phrase representation, we have chosen to adapt this method to define terms of the feature space.

One challenge emerging when terms are defined as single words is that the feature space becomes very high dimensional. In addition, words which are in the same context such as biology and biologist are defined as different terms. So, in order to define words that are in the same context with the same term and consequently to reduce dimensionality we have decided to define the terms as stemmed words. To stem the words, we have chosen to use Porter's Stemming Algorithm [24], which is the most commonly used algorithm for word stemming in English.

Preprocessing and document representation phase, which is implemented in Microsoft Visual C++ 6.0, consists of the following steps:

- Parsing the documents and case-folding
- Removing stopwords
- Stemming
- Term weighting
- Dimensionality reduction

These steps will be described briefly in the following sections.

2.1. Parsing the Documents and Case-folding

In this step, all the HTML or SGML mark-up tags and non-alpha characters are removed from the documents in the document corpora. Case-folding, which stands for converting all the characters in a document into the same case, is performed by converting all the characters into lower-case. Tokens consisting of alpha characters are extracted.

2.2. Removing Stopwords

There are words in English such as pronouns, prepositions and conjunctions that are used to provide structure in the language rather than content. These words, which are encountered very frequently and carry no useful information about the content and

a	alone	anyways	b	between
a's	along	anywhere	be	beyond
able	already	apart	became	both
about	also	appear	because	brief
above	although	appreciate	become	but
according	always	appropriate	becomes	by
accordingly	am	are	becoming	c
across	among	aren't	been	c'mon
actually	amongst	around	before	c's
after	an	as	beforehand	came
afterwards	and	aside	behind	can
again	another	ask	being	can't
against	any	asking	believe	cannot
ain't	anybody	associated	below	cant
all	anyhow	at	beside	cause
allow	anyone	available	besides	causes
allows	anything	away	best	certain
almost	anyway	awfully	better	certainly

Figure 2.1. Portion of the stopword list used

thus the category of documents, are called stopwords. Removing stopwords from the documents is very common in information retrieval. We have decided to eliminate the stopwords from the documents, which will lead to a drastic reduction in the dimensionality of the feature space. The list of 571 stopwords used in the Smart system is used [19]. This stopword list is obtained from [25]. Figure 2.1 shows a portion of the stopword list.

2.3. Stemming

In order to define words that are in the same context with the same term and consequently to reduce dimensionality, we have decided to define the terms as stemmed words. To stem the words, we have chosen to use Porter's Stemming Algorithm [24],

Word: ponies	Stem: poni
Word: caress	Stem: caress
Word: cats	Stem: cat
Word: feed	Stem: fe
Word: agreed	Stem: agre
Word: plastered	Stem: plaster
Word: motoring	Stem: motor
Word: sing	Stem: sing
Word: conflated	Stem: conflat
Word: troubling	Stem: troubl
Word: sized	Stem: size
Word: hopping	Stem: hop
Word: tanned	Stem: tan
Word: falling	Stem: fall
Word: fizzed	Stem: fizz
Word: failing	Stem: fail
Word: filing	Stem: file
Word: happy	Stem: happi

Figure 2.2. Sample of words and their corresponding stems found by Porter's Stemming Algorithm

which is the most commonly used algorithm for word stemming in English. In this way for instance, we reduce the similar terms “computer”, “computers”, and “computing” to the word stem “comput”. Implementation of Porter's Stemming Algorithm in C is downloaded from [26]. This algorithm is embedded to the preprocessing system. Figure 2.2 displays a sample of words and the stems produced by Porter's Stemming Algorithm. After stemming, terms that are shorter than two characters are also removed as they do not carry much information about the content of a document.

2.4. Term Weighting

We represent each document vector \mathbf{d} as

$$\mathbf{d} = (w_1, w_2, \dots, w_{|T|}) \quad (2.1)$$

where w_i is the weight of i^{th} term of document \mathbf{d} and $|T|$ is the number of distinct terms in the document collection. There are various term weighting approaches most of which are based on the following observations [27]:

- The relevance of a word to the topic of a document is proportional to the number of times it appears in the document.
- The discriminating power of a word between documents is less, if it appears in most of the documents in the document collection.

A comparative study of different term weighting approaches in automatic text retrieval is presented by Salton and Buckley in [28]. The term weighting approach we have applied and some other standard term weighting functions are discussed in the following subsections. We define:

tf_i as the raw frequency of term i in document \mathbf{d} ;

N as the total number of documents in the document corpus;

N_i as the number of documents in the corpus where term i appears; and

$|T|$ as the number of distinct terms in the document collection (after stopword removal and stemming is performed).

2.4.1. Boolean Weighting

Boolean weighting is the simplest method for term weighting. In this approach, the weight of a term is assigned to be 1 if the term appears in the document and it is

assigned to be 0 if the term does not appear in the document.

$$w_i = \begin{cases} 1 & \text{if } tf_i > 0 \\ 0 & \text{otherwise} \end{cases} \quad (2.2)$$

2.4.2. Term Frequency Weighting

Term Frequency (TF) weighting is also a simple method for term weighting.

$$w_i = tf_i \quad (2.3)$$

In this method, the weight of a term in a document is equal to the number of times the term appears in the document, i.e. to the raw frequency of the term in the document.

2.4.3. Term Frequency \times Inverse Document Frequency Weighting

Boolean weighting and TF weighting do not consider the frequency of the term throughout all the documents in the document corpus. Term Frequency \times Inverse Document Frequency (TF \times IDF) weighting is the most common method used for term weighting that takes into account this property. In this approach, the weight of term i in document \mathbf{d} is assigned proportionally to the number of times the term appears in the document, and in inverse proportion to the number of documents in the corpus in which the term appears.

$$w_i = tf_i \cdot \log \left(\frac{N}{N_i} \right) \quad (2.4)$$

TF \times IDF weighting approach weights the frequency of a term in a document with a factor that discounts its importance if it appears in most of the documents, as in this case the term is assumed to have little discriminating power.

2.4.4. TF×IDF Weighting with Length Normalization

In this approach, to account for documents of different lengths each document vector is normalized so that it is of unit length.

$$w_i = \frac{tf_i \cdot \log\left(\frac{N}{N_i}\right)}{\sqrt{\sum_{j=1}^{|T|} \left[tf_j \cdot \log\left(\frac{N}{N_j}\right)\right]^2}} \quad (2.5)$$

Salton and Buckley discuss that TF×IDF weighting with length normalization generally performs better than the other techniques [28]. Therefore, we applied this weighting approach in our study.

2.5. Dimensionality Reduction

There are various methods applied for dimensionality reduction in document categorization. Some common examples are Information Gain (IG), Mutual Information (MI), Chi-Square Statistic, Term Strength (TS), and Document Frequency (DF) Thresholding. We discuss these techniques briefly in the following subsections.

2.5.1. Information Gain (IG)

Information gain measures the number of bits of information gained for category prediction when the presence or absence of a term in a document is known. When the set of possible categories is $\{c_1, c_2, \dots, c_m\}$, the *IG* for each unique term t is calculated as follows [4]:

$$IG(t) = - \sum_{i=1}^m P(c_i) \cdot \log P(c_i) + P(t) \cdot \sum_{i=1}^m P(c_i|t) \cdot \log P(c_i|t) + P(\bar{t}) \cdot \sum_{i=1}^m P(c_i|\bar{t}) \cdot \log P(c_i|\bar{t}) \quad (2.6)$$

As seen from Equation 2.5, *IG* calculates the decrease in entropy when the feature is given vs. absent. $P(c_i)$ is the prior probability of category c_i . It can be estimated from the fraction of documents in the training set belonging to category c_i . $P(t)$ is

the prior probability of term t . It can be estimated from the fraction of documents in the training set in which term t is present. Likewise, $P(\bar{t})$ can be estimated from the fraction of documents in the training set in which term t is absent. Terms whose IG s are less than some predetermined threshold are removed from the feature space.

2.5.2. Mutual Information (MI)

Mutual information is a technique frequently used in statistical language modelling of word associations and related applications [29]. MI between term t and category c is defined to be [4]:

$$MI(t, c) = \log \frac{P(t \wedge c)}{P(t) \times P(c)} \quad (2.7)$$

It is estimated by using [4]:

$$MI(t, c) \approx \log \frac{A_{ct} \times N}{(A_{ct} + C_{ct}) \times (A_{ct} + B_{ct})} \quad (2.8)$$

Here, A_{ct} is the number of times t and c co-occur, B_{ct} is the number of times t occurs without c , C_{ct} is the number of times c occurs without t , and N is the total number of documents. When t and c are independent $MI(t, c)$ is equal to *zero*.

We can write equation 2.7 in the following equivalent form:

$$MI(t, c) = \log P(t|c) - \log P(t) \quad (2.9)$$

It is seen from equation 2.9 that, for terms that have an equal conditional probability, rare terms will have a higher MI value than common terms. So, MI technique has the drawback that MI values are not comparable among terms with large frequency gaps.

Category specific MI scores for a term t can be combined into a global MI score

for that term in the following two ways[4]:

$$MI_{avg}(t) = \sum_{i=1}^m P(c_i) \times MI(t, c_i) \quad (2.10)$$

or

$$MI_{max}(t) = \max_{i=1}^m \{MI(t, c_i)\} \quad (2.11)$$

Terms that have lower MI values than a predetermined threshold are eliminated.

2.5.3. Chi-Square Statistic

Chi-square (χ^2) statistic is a measure of association. In statistics chi-square measure is formulated as:

$$\chi^2 = \sum_i \sum_j \frac{(f_{ij} - \hat{f}_{ij})^2}{\hat{f}_{ij}} \quad (2.12)$$

Here, f_{ij} is the observed frequency of the cell in row i and column j and \hat{f}_{ij} is the expected frequency of that cell. A more detailed discussion of the χ^2 statistic can be found in [30].

The χ^2 statistic measures the degree of dependence between a certain term and a certain category. That is, it measures to what degree a certain term is indicative of membership or non-membership of a document in a certain category [31]. The χ^2 statistic is reformulated and used for the task of document categorization by Yang and Pedersen [4], Ng *et al.* [6], and Spitters [31] as follows:

$$\chi^2(t, c) = \frac{N \times (A_{ct}D_{ct} - C_{ct}B_{ct})^2}{(A_{ct} + C_{ct}) \times (B_{ct} + D_{ct}) \times (A_{ct} + B_{ct}) \times (C_{ct} + D_{ct})} \quad (2.13)$$

Here, we have a 2×2 contingency table. The first row stands for the number of documents that contain term t , the second row stands for the number of documents

that do not contain term t , the first column stands for the number of documents that belong to category c , and the second column stands for the number of documents that do not belong to category c . So, A_{ct} is the number of documents that belong to category c and contain term t , B_{ct} is the number of documents that do not belong to category c but contain term t , C_{ct} is the number of documents that belong to category c but do not contain term t , D_{ct} is the number of documents that do not belong to category c and do not contain term t , and N is the total number of documents in the corpus. Two different measures can be computed based on the χ^2 statistic [4]:

$$\chi_{avg}^2(t) = \sum_{i=1}^m P(c_i) \times \chi^2(t, c_i) \quad (2.14)$$

or

$$\chi_{max}^2(t) = \max_{i=1}^m \{\chi^2(t, c_i)\} \quad (2.15)$$

Terms that have lower χ^2 values than a predetermined threshold are eliminated.

2.5.4. Term Strength (TS)

Term strength method, estimates term importance based on how commonly a term is likely to appear in closely related documents [4]. The first step in this method is to use a training set of documents to find document pairs which have a similarity larger than a predetermined threshold. In the next step TS is calculated based on the estimated conditional probability that a term appears in the second document given that it appears in the first one. Suppose, d_i and d_j are any pair of distinct but related documents. Then the TS of term t is defined to be [4]:

$$TS(t) = P(t \in d_j | t \in d_i) \quad (2.16)$$

Unlike IG , MI , and χ^2 statistic, TS is an unsupervised dimensionality reduction technique where document categories are not used. It is based on document clustering and assumes that documents with many shared words are related and the terms that are

heavily shared among these related documents are relatively informative.

2.5.5. Document Frequency (DF) Thresholding

Document frequency (DF) of a term is the number of documents that term appears. In this technique, the document frequency of each unique term is computed and terms whose document frequencies are less than a predetermined threshold are eliminated. The basic assumption behind this technique is that rare terms are either non-informative for document categorization or they do not have much weight in global performance. This technique can also lead to improvement in categorization accuracy in case rare terms are noise terms. However, DF is usually not used for aggressive term elimination because there is another widely accepted assumption in information retrieval that low- DF terms are distinctive and thus relatively informative and for this reason should not be removed aggressively [4].

A comparative study of feature selection in text categorization is presented by Yang and Pedersen in [4]. It has been reported that IG and χ^2 statistic performed the best. However, DF , the simplest and the most efficient method in terms of computational complexity, performed similar to IG and χ^2 statistics. It has been suggested that DF can be reliably used instead of IG and χ^2 statistics when computation performances of the latter two are too expensive.

Another point to consider is that IG , MI and χ^2 statistics are supervised techniques and use information about term-category associations. As our main focus is on unsupervised techniques for document organization, these methods are not suitable to be applied in our study. To reduce the dimensionality of the data, we apply DF Thresholding. We define the DF threshold as 1 and hence remove the terms that appear in only one document.

2.6. Document Similarity Measure

To use a clustering or classification algorithm, a similarity measure between two documents must be defined. In our study we use the widely used cosine similarity measure to calculate the similarity of two documents. This measure is defined as [18]:

$$\cos(\mathbf{d}_1, \mathbf{d}_2) = \frac{\mathbf{d}_1 \bullet \mathbf{d}_2}{\|\mathbf{d}_1\| \|\mathbf{d}_2\|} \quad (2.17)$$

that is, it is the dot product of \mathbf{d}_1 and \mathbf{d}_2 divided by the lengths of \mathbf{d}_1 and \mathbf{d}_2 .

3. UNSUPERVISED TECHNIQUES FOR DOCUMENT CLUSTERING

In unsupervised clustering, we have unlabelled collection of documents. The aim is to cluster the documents without additional knowledge or intervention such that documents within a cluster are more similar than documents between clusters. Traditional clustering techniques can be categorized into two major groups as partitional and hierarchical. In this chapter we discuss these groups and their main representatives.

3.1. Partitional Clustering Techniques

Partitional algorithms produce un-nested, non-overlapping partitions of documents that usually locally optimize a clustering criterion. The general methodology is as follows: given the number of clusters k , an initial partition is constructed; next the clustering solution is refined iteratively by moving documents from one cluster to another. In the following sub-sections we discuss the most popular partitional algorithm k -means, and its variant bisecting k -means which has been applied to cluster documents by Steinbach *et al.* in [18] and has been shown to generally outperform agglomerative hierarchical algorithms.

3.1.1. K-Means Clustering

The idea behind the k -means algorithm, discussed by Hartigan [32], is that each of k clusters can be represented by the mean of the documents assigned to that cluster, which is called the *centroid* of that cluster. It is discussed by Berkhin in [33] that there are two versions of k -means algorithm known. The first version is the *batch* version and is also known as Forgy's algorithm [34]. It consists of the following two-step major iterations:

- (i) Reassign all the documents to their nearest centroids

- (ii) Recompute centroids of newly assembled groups

Before the iterations start, firstly k documents are selected as the initial centroids. Iterations continue until a stopping criterion such as no reassignments occur is achieved.

In our experiments we used the second version of k -means algorithm, which is known as *online* or *incremental* version. It is discussed by Steinbach *et al.* [18] and Berkhin [33] that online k -means performs better than the batch version in the domain of text document collections. Initially, k documents from the corpus are selected randomly as the initial centroids. Then, iteratively documents are assigned to their nearest centroid and centroids are updated incrementally, i.e., after each assignment of a document to its nearest centroid. Iterations stop, when no reassignments of documents occur.

We define the centroid vector \mathbf{c} of cluster C of documents as follows:

$$\mathbf{c} = \frac{\sum_{d \in C} \mathbf{d}}{|C|} \quad (3.1)$$

So, \mathbf{c} is obtained by averaging the weights of the terms of the documents in C . Analogously, we define the similarity between a document \mathbf{d} and a centroid vector \mathbf{c} by cosine similarity measure as

$$\cos(\mathbf{d}, \mathbf{c}) = \frac{\mathbf{d} \bullet \mathbf{c}}{\|\mathbf{d}\| \|\mathbf{c}\|} \quad (3.2)$$

Note that although documents are of unit length, centroid vectors are not necessarily of unit length.

3.1.2. Bisecting K-Means

Although bisecting k -means is actually a divisive clustering algorithm that achieves a hierarchy of clusters by repeatedly applying the basic k -means algorithm, we discuss it in this section as it is a variant of k -means.

In each step of bisecting k -means a cluster is selected to be split and it is split into two by applying basic k -means for $k = 2$. The largest cluster, that is the cluster containing the maximum number of documents, or the cluster with the least overall similarity can be chosen to be split. We performed experiments in both ways and observed that they perform similarly. So, in the experiment results section we reveal only the results of the case when the largest cluster is selected to be split.

3.2. Hierarchical Clustering Techniques

Hierarchical clustering algorithms produce a cluster hierarchy named a dendrogram [33]. These algorithms can be categorized as divisive (top-down) and agglomerative (bottom-up) [16, 33]. We discuss these approaches in the following sub-sections.

3.2.1. Divisive Hierarchical Clustering

Divisive algorithms start with one cluster of all documents and at each iteration split the most appropriate cluster until a stopping criterion such as a requested number k of clusters is achieved.

A method to implement a divisive hierarchical algorithm is described by Kaufman and Rousseeuw in [35]. In this technique in each step the cluster with the largest diameter is split, i.e. the cluster containing the most distant pair of documents. As we use document similarity instead of distance as a proximity measure, the cluster to be split is the one containing the least similar pair of documents. Within this cluster the document with the least average similarity to the other documents is removed to form a new singleton cluster. The algorithm proceeds by iteratively assigning the documents in the cluster being split to the new cluster if they have greater average similarity to the documents in the new cluster. To our knowledge, divisive hierarchical clustering in this sense has not been applied to document corpora. This method is not robust to outliers and in our experiments we observe that documents in the cluster being split generally tend to remain in the larger old cluster and for small number of clusters k , clustering quality is not comparable with the other algorithms we evaluated. So, we

made a slight modification to this algorithm. In our version we select the least similar pair of documents in the cluster being split and remove them to form two new singleton clusters. The rest of the documents in the cluster are assigned iteratively to one of the new clusters by taking the average similarity as criterion.

3.2.2. Agglomerative Hierarchical Clustering

Agglomerative clustering algorithms start with each document in a separate cluster and at each iteration merge the most similar clusters until the stopping criterion is met. They are mainly categorized as single-link, complete-link and average-link depending on the method they define inter-cluster similarity. Figure 3.1 illustrates the idea:

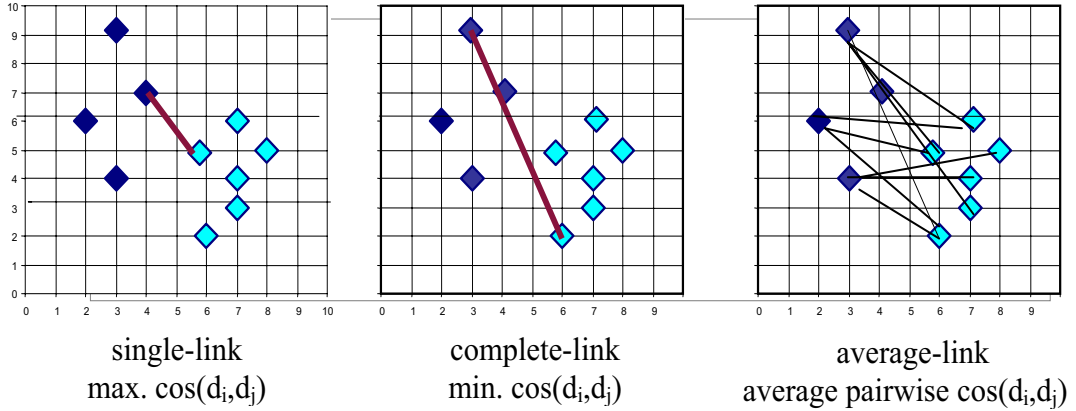


Figure 3.1. Inter-cluster similarity defined by single-link, complete-link, and average-link

3.2.2.1. Single-link. The single-link method defines the similarity of two clusters C_i and C_j as the similarity of the two most similar documents $\mathbf{d}_i \in C_i$ and $\mathbf{d}_j \in C_j$:

$$similarity_{single-link}(C_i, C_j) = \max_{\mathbf{d}_i \in C_i, \mathbf{d}_j \in C_j} |\cos(\mathbf{d}_i, \mathbf{d}_j)| \quad (3.3)$$

3.2.2.2. Complete-link. The complete-link method defines the similarity of two clusters C_i and C_j as the similarity of the two least similar documents $\mathbf{d}_i \in C_i$ and $\mathbf{d}_j \in C_j$:

$$similarity_{complete-link}(C_i, C_j) = \min_{\mathbf{d}_i \in C_i, \mathbf{d}_j \in C_j} |cos(\mathbf{d}_i, \mathbf{d}_j)| \quad (3.4)$$

3.2.2.3. Average-link. The average-link method defines the similarity of two clusters C_i and C_j as the average of the pairwise similarities of the documents from each cluster:

$$similarity_{average-link}(C_i, C_j) = \frac{\sum_{\mathbf{d}_i \in C_i, \mathbf{d}_j \in C_j} |cos(\mathbf{d}_i, \mathbf{d}_j)|}{n_i n_j} \quad (3.5)$$

where n_i and n_j are sizes of clusters C_i and C_j respectively.

4. SUPERVISED TECHNIQUES FOR DOCUMENT CLASSIFICATION

Supervised algorithms assume that the category structure or hierarchy of a text database is already known. They require a training set of labelled documents and return a function that maps documents to the pre-defined class labels. As discussed previously, knowing the category structure in advance and generation of correctly labelled training set are very challenging or even impossible in large and dynamic text databases. In this section we discuss the most popular supervised algorithms k -NN, naive Bayes, and support vector machines, that we have evaluated.

4.1. K Nearest Neighbor Classification

K -NN (k -nearest neighbor) classification is a popular instance-based learning method [36] that has been shown to be a strong performer in the task of text categorization [3, 8].

The algorithm works as follows: First, given a test document \mathbf{x} , the k nearest neighbors among the training documents are found. The category labels of these neighbors are used to estimate the category of the test document. In the traditional approach, the most common category label among the k -nearest neighbors is assigned to the test document.

Weighted k -NN is a refinement to the traditional approach. In weighted k -NN, the contribution of each of the k nearest neighbors is weighted according to its similarity to the test document \mathbf{x} . Then, for each category, the similarity of the neighbors belonging to that category are summed to obtain the score of the category for \mathbf{x} . That is, the score of category c_j for the test document \mathbf{x} is

$$score(c_j, \mathbf{x}) = \sum_{\mathbf{d}_i \in N(\mathbf{x})} cos(\mathbf{x}, \mathbf{d}_i) \cdot y(\mathbf{d}_i, c_j) \quad (4.1)$$

where \mathbf{d}_i is a training document; $N(x)$ is the set of the k training documents nearest to \mathbf{x} ; $\cos(\mathbf{x}, \mathbf{d}_i)$ is the cosine similarity between the test document \mathbf{x} and the training document \mathbf{d}_i ; and $y(\mathbf{d}_i, c_j)$ is a function whose value is 1 if \mathbf{d}_i belongs to category c_j and 0 otherwise. The test document \mathbf{x} is assigned to the category with the highest score.

In our study, we evaluated both the traditional k -NN and weighted k -NN for varying k parameter values and report the results of both approaches for the best k value in the experiment results chapter.

4.2. Naive Bayes Approach

The naive Bayes (NB) classifier is a probabilistic model that uses the joint probabilities of terms and categories to estimate the probabilities of categories given a test document [36]. The naive part of the classifier comes from the simplifying assumption that all terms are conditionally independent of each other given a category. Because of this independence assumption, the parameters for each term can be learned separately and this simplifies and speeds the computation operations compared to non-naive Bayes classifiers.

There are two common event models for NB text classification, discussed by McCallum and Nigam in [9], *multinomial model* and *multivariate Bernoulli model*. In both models classification of test documents is performed by applying the Bayes' rule [36]:

$$P(c_j|\mathbf{d}_i) = \frac{P(c_j) \cdot P(\mathbf{d}_i|c_j)}{P(\mathbf{d}_i)} \quad (4.2)$$

where \mathbf{d}_i is a test document and c_j is a category. The posterior probability of each category c_j given the test document \mathbf{d}_i , i.e. $P(c_j|\mathbf{d}_i)$, is calculated and the category with the highest probability is assigned to \mathbf{d}_i . In order to calculate $P(c_j|\mathbf{d}_i)$, $P(c_j)$ and $P(\mathbf{d}_i|c_j)$ have to be estimated from the training set of documents. Note that $P(\mathbf{d}_i)$ is same for each category so we can eliminate it from the computation. The category

prior probability, $P(c_j)$, can be estimated as follows:

$$\hat{P}(c_j) = \frac{\sum_{i=1}^N y(\mathbf{d}_i, c_j)}{N}, \quad (4.3)$$

where, N is number of training documents and $y(\mathbf{d}_i, c_j)$ is defined as follows:

$$y(\mathbf{d}_i, c_j) = \begin{cases} 1 & \text{if } \mathbf{d}_i \in c_j \\ 0 & \text{otherwise} \end{cases} \quad (4.4)$$

So, prior probability of category c_j is estimated by the fraction of documents in the training set belonging to c_j . $P(\mathbf{d}_i|c_j)$ parameters are estimated in different ways by the multinomial model and multivariate Bernoulli model. We present these models in the following two sub-sections.

4.2.1. Multinomial Model

In the multinomial model a document \mathbf{d}_i is an ordered sequence of term events, drawn from the term space T . The naive Bayes assumption is that the probability of each term event is independent of term's context, position in the document, and length of the document. So, each document \mathbf{d}_i is drawn from a multinomial distribution of terms with number of independent trials equal to the length of \mathbf{d}_i . The probability of a document \mathbf{d}_i given its category c_j can be approximated as:

$$P(\mathbf{d}_i|c_j) \approx \prod_{i=1}^{|d_i|} P(t_i|c_j), \quad (4.5)$$

where $|d_i|$ is the number of terms in document \mathbf{d}_i ; and t_i is the i^{th} term occurring in document \mathbf{d}_i . Thus the estimation of $P(\mathbf{d}_i|c_j)$ is reduced to estimating each $P(t_i|c_j)$ independently. The following Bayesian estimate is used for $P(t_i|c_j)$:

$$\hat{P}(t_i|c_j) = \frac{1 + TF(t_i, c_j)}{|T| + \sum_{t_k \in T} TF(t_k, c_j)} \quad (4.6)$$

Here, $TF(t_i, c_j)$ is the total number of times term t_i occurs in the training set documents belonging to category c_j . The summation term in the denominator stands for the total number of term occurrences in the training set documents belonging to category c_j . This estimator is called Laplace estimator and assumes that the observation of each word is a priori likely [37].

4.2.2. Multivariate Bernoulli Model

Multivariate Bernoulli model for naive Bayes classification is the event model we used and evaluated in our study. In this model a document is represented by a vector of binary features indicating the terms that occur and that do not occur in the document. Here, the document is the event and absence or presence of terms are the attributes of the event. The naive Bayes assumption is that the probability of each term being present in a document is independent of the presence of other terms in a document. To state differently, the absence or presence of each term is dependent only on the category of the document. Then, $P(\mathbf{d}_i|c_j)$, the probability of a document given its category is simply the product of the probability of the attribute values over all term attributes:

$$P(\mathbf{d}_i|c_j) = \prod_{t=1}^{|T|} (B_{it} \cdot P(v_t|c_j) + (1 - B_{it})(1 - P(v_t|c_j))), \quad (4.7)$$

where $|T|$ is the number of terms in the training set and B_{it} is defined as follows:

$$B_{it} = \begin{cases} 1 & \text{if term } t \text{ appears in document } \mathbf{d}_i \\ 0 & \text{otherwise} \end{cases} \quad (4.8)$$

Thus, a document can be seen as a collection of multiple independent Bernoulli experiments, one for each term in the term space. The probabilities of each of these term events are defined by the class-conditional term probabilities $P(v_t|c_j)$. We can estimate the probability of term v_t in category c_j as follows:

$$\hat{P}(v_t|c_j) = \frac{1 + \sum_{i=1}^N B_{it} \cdot y(\mathbf{d}_i, c_j)}{2 + \sum_{i=1}^N y(\mathbf{d}_i, c_j)}, \quad (4.9)$$

where, N is number of training documents and $y(\mathbf{d}_i, c_j)$ is defined as in equation 4.4.

Different from the multinomial model, the multivariate Bernoulli model does not take into account the number of times each term occurs in the document, and it explicitly includes the non-occurrence probability of terms that are absent in the document [9].

4.3. Support Vector Machines

Support Vector Machines (SVM) is a technique introduced by Vapnik in 1995, which is based on the Structural Risk Minimization principle [38]. It is designed for solving two-class pattern recognition problems. The problem is to find the decision surface that separates the positive and negative training examples of a category with maximum margin. Figure 4.1 illustrates the idea for linearly separable data points. A decision surface in a linearly separable space is a hyperplane. The dashed lines parallel to the solid line show how much the decision surface can be moved without leading to a misclassification of data. *Margin* is the distance between these parallel lines. Examples closest to the decision surface are called *support vectors*.




Figure 4.1. Support vector machines find the hyperplane h that separates positive and negative training examples with maximum margin. Support vectors are marked with circles

For the linearly separable case, the decision surface is a hyperplane that can be written as [8]:

$$\mathbf{w} \bullet \mathbf{d} + b = 0 \quad (4.10)$$

where \mathbf{d} is a document to be classified, and vector \mathbf{w} and constant b are learned from the training set. The SVM problem is to find \mathbf{w} and b that satisfy the following constraints [5]:

$$\text{Minimize } \|\mathbf{w}\|^2 \quad (4.11)$$

$$\text{so that } \forall i : y_i [\mathbf{w} \bullet \mathbf{d}_i + b] \geq 1 \quad (4.12)$$

Here, $i \in \{1, 2, \dots, N\}$, where N is the number of documents in the training set; and y_i equals $+1$ if document \mathbf{d}_i is a positive example for the category being considered and equals -1 otherwise. This optimization problem can be solved by using quadratic programming techniques [39].

SVM can be also used to learn non-linear decision functions such as polynomial of degree d or radial basis function (RBF) with variance γ . These kernel functions can be illustrated as follows:

$$K_{polynomial}(\mathbf{d}_1, \mathbf{d}_2) = (\mathbf{d}_1 \bullet \mathbf{d}_2 + 1)^d \quad (4.13)$$

$$K_{rbf}(\mathbf{d}_1, \mathbf{d}_2) = \exp(\gamma(\mathbf{d}_1 - \mathbf{d}_2)^2) \quad (4.14)$$

In our study we evaluated SVM with linear kernel, polynomial kernel with different degrees, and with RBF kernel with different γ parameters. For our experiments we used the *SVM^{light}* system implemented by Joachims [40].

5. EXPERIMENT RESULTS

We experimentally evaluated the performance of the partitional and hierarchical clustering techniques, and the supervised classification techniques for document organization on five different standard data sets. In this chapter we first describe the data sets we used in our experiments and our experimental methodology. Next we present and evaluate the experimental results for the unsupervised and supervised techniques separately.

5.1. Document Data Sets

In our experiments we used five standard document corpora widely used in automatic text organization research. Summary description of these document sets after preprocessing as described in Chapter 2 is presented in Table 5.1.

Table 5.1. Summary description of document sets

Data set	# of documents	# of classes	# of terms
Classic3	3,891	3	6,729
Hitech	1,530	6	10,919
LA1	2,134	6	14,363
Reuters-21578	12,902	90	12,772
Wap	1,560	20	8,061

Classic3 data set contains 1,398 CRANFIELD documents from aeronautical system papers, 1,033 MEDLINE documents from medical journals and 1,460 CISI documents from information retrieval papers.

The Hitech data set was derived from the San Jose Mercury newspaper articles which are delivered as part of the TREC collection [41]. The classes of this document corpora are computers, electronics, health, medical, research, and technology.

LA1 data set consists of documents from Los Angeles Times newspaper used in TREC-5 [41]. The categories correspond to the desk of the paper that each article appeared. The data set consists of documents from entertainment, financial, foreign, metro, national, and sports desks.

The documents in Reuters-21578 v1.0 document collection [42], which is considered as the standard benchmark for automatic document organization systems, have been collected from Reuters newswire in 1987. This corpus consists of 21,578 documents. 135 different categories have been assigned to the documents. The maximum number of categories assigned to a document is 14 and the mean is 1.24. Frequency of occurrence of categories varies widely. For instance the “earnings” category is assigned to 2,709 training documents, but 75 categories are assigned to less than 10 training documents. 21 categories are not assigned to any training documents [27].

We have obtained Reuters-21578 corpus from [42]. The documents in the corpus are in SGML format and we used JBuilder 6.0 to parse them. For our results to be comparable with the results of other studies, we used the modified Apte (ModApte) splitting method, which has been most frequently used to divide the corpus into training and test sets [42]. This splitting method uses a subset of 12,902 documents from the whole corpus. It assigns documents from April 7, 1987 and before to the training set and from April 8, 1987 and after to the test set. The documents are organized as follows:

- Documents with tag values LEWISSPLIT=”TRAIN“ and TOPICS=”YES“ are included in the training set. The training set consists of 9,603 documents.
- Documents with tag values LEWISSPLIT=”TEST“ and TOPICS=”YES“ are included in the test set. The test set consists of 3,299 documents.
- Documents with tag values LEWISSPLIT=”NOT-USED“ and TOPIC=”YES“ or TOPICS=”NO“ or TOPICS=”BYPASS“ are not used. Number of unused documents is 8,676.

In the Reuters-21578 data set we removed the classes that do not exist both in the

cocoa	money-supply	rice	iron-steel	palladium
grain	coffee	rubber	hog	nickel
wheat	ship	copra-cake	propane	lumber
corn	sugar	palm-oil	heat	jet
barley	trade	palmkernel	gas	instal-debt
oat	reserves	tea	jobs	dfl
sorghum	meal-feed	alum	lei	dmk
veg-oil	soy-meal	gold	yen	coconut-oil
lin-oil	rye	platinum	zinc	cpu
soy-oil	cotton	strategic-metal	orange	cotton-oil
sun-oil	carcass	tin	pet-chem	naphtha
soybean	livestock	rapeseed	fuel	nzdlr
oilseed	crude	groundnut-oil	wpi	rand
sunseed	nat-gas	rape-oil	potato	coconut
earn	cpi	dlr	lead	castor-oil
acq	gnp	l-cattle	groundnut	nkr
copper	money-fx	retail	income	sun-meal
housing	interest	ipi	bop	silver

Figure 5.1. Class names of Reuters-21578 data set

training set and in the test set remaining with 90 classes out of 135. In our experiments, we divide the training set into two and use one half as the training set to train the classifiers and the other half as the validation set to optimize the parameters. After the optimization phase, we again train the classifier with the whole training set of 9,603 documents and test the performance with the separate test set of 3,299 documents. We report the results for the test set of this corpus. The class names of the Reuters-21578 data set are listed in Figure 5.1.

Wap data set consists of 1,560 web pages from Yahoo! subject hierarchy collected and classified into 20 different classes for the WebACE project [43]. The class names of the Wap data set are listed in Figure 5.2.

art	music
business	online
cable	people
culture	politics
entertainment	review
film	sports
health	stage
industry	technology
media	television
multimedia	variety

Figure 5.2. Class names of the Wap data set

Documents in Classic3, Hitech, LA1, and Wap data sets are assigned to exactly one category, whereas some documents in Reuters-21578 data set are assigned to multiple categories. For Classic3, Hitech, LA1, and Wap data sets, which can be also obtained from [44], we performed 10-fold stratified cross-validation in our experiments. We report the average results of 10-folds over the test sets for these document collections.

5.2. Evaluation of the Clustering Techniques

5.2.1. Evaluation Metrics

There are two types of measures to evaluate cluster quality, *internal quality measure* and *external quality measure* [18]. Internal quality measure does not use external knowledge such as class label information to evaluate the produced clustering solution. On the other hand, external quality measure relies on labelled test document corpora. Its methodology is to compare the resulting clusters to labelled classes and measure the degree to which documents from the same class are assigned to the same cluster. To evaluate the quality of the unsupervised clustering algorithms, we use overall similarity, which is an internal quality measure, purity, which is an external quality measure, and

two widely used external quality measures in text mining: entropy and F-measure [18].

5.2.1.1. Overall Similarity. Overall similarity is an internal quality measure that uses weighted similarity of internal cluster similarities to measure the cohesiveness of the produced clusters. Internal cluster similarity I for cluster C_j can be computed as:

$$I_j = \frac{1}{n_j^2} \sum_{d \in C_j, d' \in C_j} \cos(\mathbf{d}, \mathbf{d}') \quad (5.1)$$

where n_j is number of documents in cluster j . We can rewrite I_j as:

$$I_j = \left(\frac{1}{n_j} \sum_{d \in C_j} \mathbf{d} \right) \bullet \left(\frac{1}{n_j} \sum_{d' \in C_j} \mathbf{d}' \right) = \mathbf{c} \bullet \mathbf{c} = \|\mathbf{c}\|^2 \quad (5.2)$$

So, I_j the average pairwise similarity between all points in cluster C_j is equal to the square of the length of the centroids of that cluster. Overall similarity of the clustering solution is:

$$\text{Overall Similarity} = \sum_j \frac{n_j}{N} I_j \quad (5.3)$$

where N is the total number of documents in the corpus.

5.2.1.2. Purity. Purity measures the extent to which each cluster contains documents from primarily one class. For a particular cluster j of size n_j , purity of this cluster is defined to be:

$$P_j = \frac{1}{n_j} \max_i n_{ji}, \quad (5.4)$$

where n_{ji} is number of documents of class i that are assigned to cluster j . So, P_j is the fraction of overall cluster size that the largest class of documents assigned to that cluster constitute. The overall purity of the clustering solution is obtained by the

weighted sum of individual cluster purities.

$$P = \sum_j \frac{n_j}{N} P_j \quad (5.5)$$

where N is total number of documents in the document collection. In general, the larger are the values of purity, the better is the clustering solution.

5.2.1.3. Entropy. Entropy measures the homogeneity of the clusters. A perfect clustering solution will be the one that leads to clusters that consist of documents from only one class. In that case the entropy will be zero. In general, the lower the entropy is, the more homogenous the clusters are. The total entropy E for a set of clusters is obtained by summing the entropies E_j of each cluster j weighted by its size:

$$E_j = - \sum_i P(i, j) \cdot \log P(i, j) \quad (5.6)$$

$$E = \sum_j \frac{n_j}{N} E_j \quad (5.7)$$

$P(i, j)$ is the probability that a document has class label i and is assigned to cluster j , n_j is size of cluster j and N is the total number of documents in the corpus.

5.2.1.4. F-measure. The F-measure cluster evaluation metric combines the precision and recall ideas from information retrieval. Each cluster is considered as if it were the result of a query and each class as if it were the desired set of documents for the query. Recall and precision for each cluster j and class i are calculated as follows:

$$Recall(i, j) = n_{ij}/n_i, \quad Precision(i, j) = n_{ij}/n_j \quad (5.8)$$

Here, n_{ij} is the number of documents with class label i in cluster j , n_i is the number of documents with class label i and n_j is the number of documents in cluster j . The

F-measure of cluster j and class i is calculated as follows:

$$F(i, j) = \frac{2\text{Recall}(i, j)\text{Precision}(i, j)}{\text{Recall}(i, j) + \text{Precision}(i, j)} \quad (5.9)$$

For an entire hierarchical clustering, the F-measure of any class is the maximum value it achieves at any node in the hierarchy tree. An overall value for the F-measure is calculated by taking the weighted average of all values for the F-measure:

$$F = \sum_i \frac{n_i}{N} \max_j F(i, j) \quad (5.10)$$

The F-measure values are in the interval (0,1) and larger F-measure values correspond to higher clustering quality.

5.2.2. Results and Discussion

Figures 5.3, 5.4, 5.5, 5.6, and 5.7 display the performance of k -means, bisecting k -means, divisive hierarchical; and single-link, complete-link, and average-link agglomerative hierarchical algorithms in terms of entropy, F-measure, purity, and overall similarity evaluation metrics over Classic3, Wap, Reuters-21578, LA1, and Hitech document sets, respectively.

For all the document collections, single-link algorithm performs considerably worse than the other algorithms. This algorithm assigns each document to the cluster of its nearest neighbor. However, any two documents may share many of the same terms and be nearest neighbors without belonging to the same topic (class). Figure 5.8 displays the distribution of terms of the data sets among the topics and Figure 5.9 displays for each data set the per cent of documents whose nearest neighbor is of a different topic. For instance, in the Classic3 data set about 50 per cent of the terms occur in only one topic and the rest are shared among two and three topics. In the Classic3 data set, only about two per cent of documents which are nearest neighbors belong to different topics, while in the other data set more than 30 per cent of the documents which are nearest neighbors belong to different topics. These properties make

Classic3 a relatively easy data set. The topics of Classic3 data set are disjoint from each other, whereas some of the topics of the other four data sets are more general and overlap with each other. For instance the health and medical topics; and computers, electronics, and technology topics of Hitech data set overlap with each other. When we look at Figure 5.8, we see that about 30 per cent of the terms in the Hitech data set are shared among two topics, about 16 per cent are shared among three topics, and nearly 18 per cent of the terms are shared among all the six topics. A similar discussion holds also for LA1, Reuters-21578, and Wap data sets. For instance, Wap data set contains documents from the very general people and variety topics. Likewise, the topics television, media, film; and the topics business and industry overlap with each other. In Reuters-21578 data set the topics veg-oil, lin-oil, sun-oil, cotton-oil, and castor-oil; the topics gold, platinum, silver, nickel, copper, and zinc; and the topics corn, wheat, and grain share many common terms.

In the previous paragraph we discussed the general properties of document collections, that documents may share many common terms and be close to each other, even be nearest neighbors but belong to different topics. These properties are the ones that make the task of document organization challenging. These properties may account for the bad performance of the single-link algorithm in the domain of text documents. However, single-link algorithm achieves comparable F-measure performance to complete-link on the Classic3 data set and it achieves better F-measure performance than online k -means for the same data set for number of clusters greater than 30. This data set consists of three classes and the decrease in F-measure performance for large number of clusters may be due to considerable decrease in recall. The similar discussion for the F-measure performance of k -means also holds for LA1 and Hitech data sets. Note that F-measure values for hierarchical clustering algorithms do not depend on the number of clusters as its calculation is done considering the whole hierarchy tree. This is not the case for partitional algorithms.

Among the agglomerative hierarchical clustering algorithms, average-link performs the best. We have explained intuitively above the reason for the poor performance of single-link. On the other hand, complete-link algorithm is based on the

assumption that all the documents in the cluster are very similar to each other. This assumption does not account for the high dimensional diverse nature of text document domain, where each distinct word is considered as a different feature and context knowledge such as synonyms, hyponyms, and hypernyms are not considered. Average-link algorithm overcomes these problems by relying on more global properties to measure cluster similarity. In this algorithm, similarity of two clusters is measured by taking into account all the documents in both clusters.

Agglomerative algorithms are more common hierarchical clustering techniques than the divisive ones since at each stage of a divisive algorithm we should consider all possible ways of splitting data. This poses a great overhead for large data sets. However, as we discussed in Chapter 3, Kaufman and Rousseeuw present a divisive method in [35] that considers only a subset of the all possible partitions. We discussed in Chapter 3 that we have modified this method slightly so that it becomes more robust to outliers and our results reveal that it achieves similar performance to average-link agglomerative hierarchical clustering algorithm.

When entropy, purity, and overall similarity metrics are considered online k -means and bisecting k -means perform better than divisive hierarchical; and complete-link and single-link agglomerative hierarchical clustering algorithms. They either achieve better or similar performance to average-link algorithm.

In terms of the F-measure, bisecting k -means performs better than k -means and similar to average-link and divisive hierarchical clustering algorithms. Calculating the cosine similarity of a document to a cluster centroid is equivalent to calculating the average similarity of the document to all the documents in that cluster. Hence, like average-link and divisive hierarchical algorithms, k -means and its bisecting variant also rely more on global properties to make decisions. A property of agglomerative and divisive hierarchical algorithms that degrades their performance is that, in contrast to k -means and bisecting k -means they do not revisit intermediate clusters for the purpose of improving them, once they are constructed. K -means on the other hand has the drawback that its performance depends very much on the parameter k and the initial

selection of centroids.

An important superiority of k -means and bisecting k -means is their $O(N)$ time complexity compared to the $O(N^2)$ time complexity of agglomerative and divisive hierarchical clustering algorithms. This especially becomes an important criterion when N , the number of documents, is large.

In order to visualize better the quality of the clustering solutions obtained by the algorithms, we present the performance, cluster-class distribution, and most descriptive five terms of each cluster together with the percentage of average similarity between the documents in the cluster each term explains for online k -means, bisecting k -means; divisive hierarchical; and single-link, complete-link, and average-link agglomerative hierarchical algorithms over Classic3 data set in Figures 5.10, 5.11, 5.12, 5.13, 5.14, and 5.15, respectively.

We can conclude that k -means and bisecting k -means could successfully discriminate the three topics *cisi*, *cran*, and *med*. Also, the most-descriptive five terms of each cluster can be interpreted as successful representative keywords for the topic constituting the vast majority of that cluster.

On the other hand, the remaining algorithms tend to produce unbalanced clusters and could not discriminate the topics from each other successfully. For instance, average-link algorithm could discriminate the topic *med* (medicine) in the second cluster, however the first cluster is very inhomogeneous and the third cluster contains only one document. Similar arguments hold for divisive and complete-link algorithms. We can see from Figure 5.15 that single-link algorithm is faced with the chaining effect such that the first cluster contains only one document, the second cluster contains only three documents, while the third cluster contains all the rest documents.

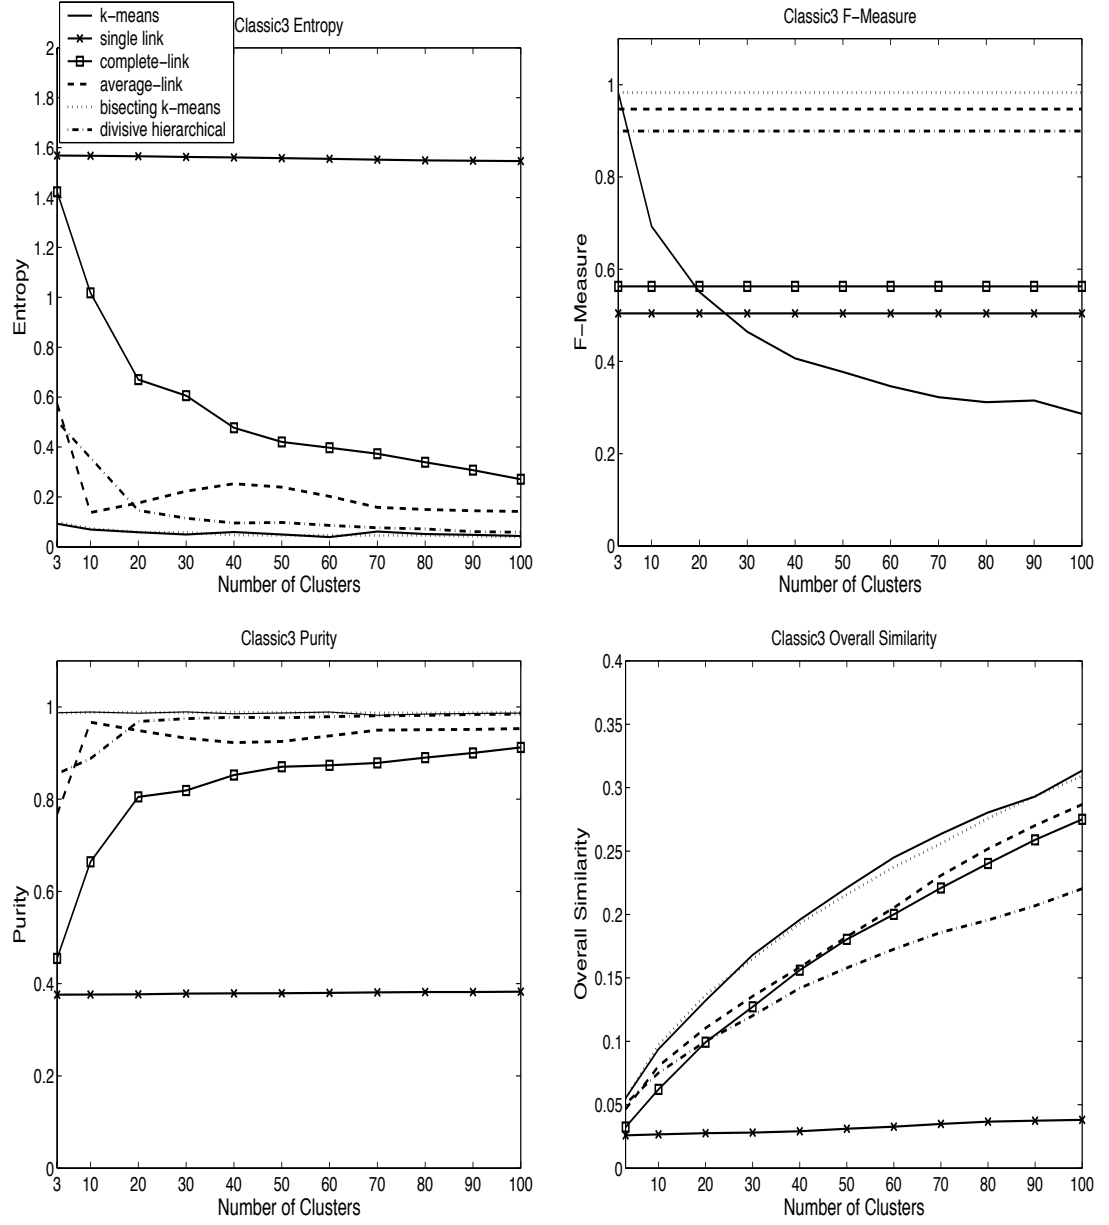


Figure 5.3. Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Classic3 data set

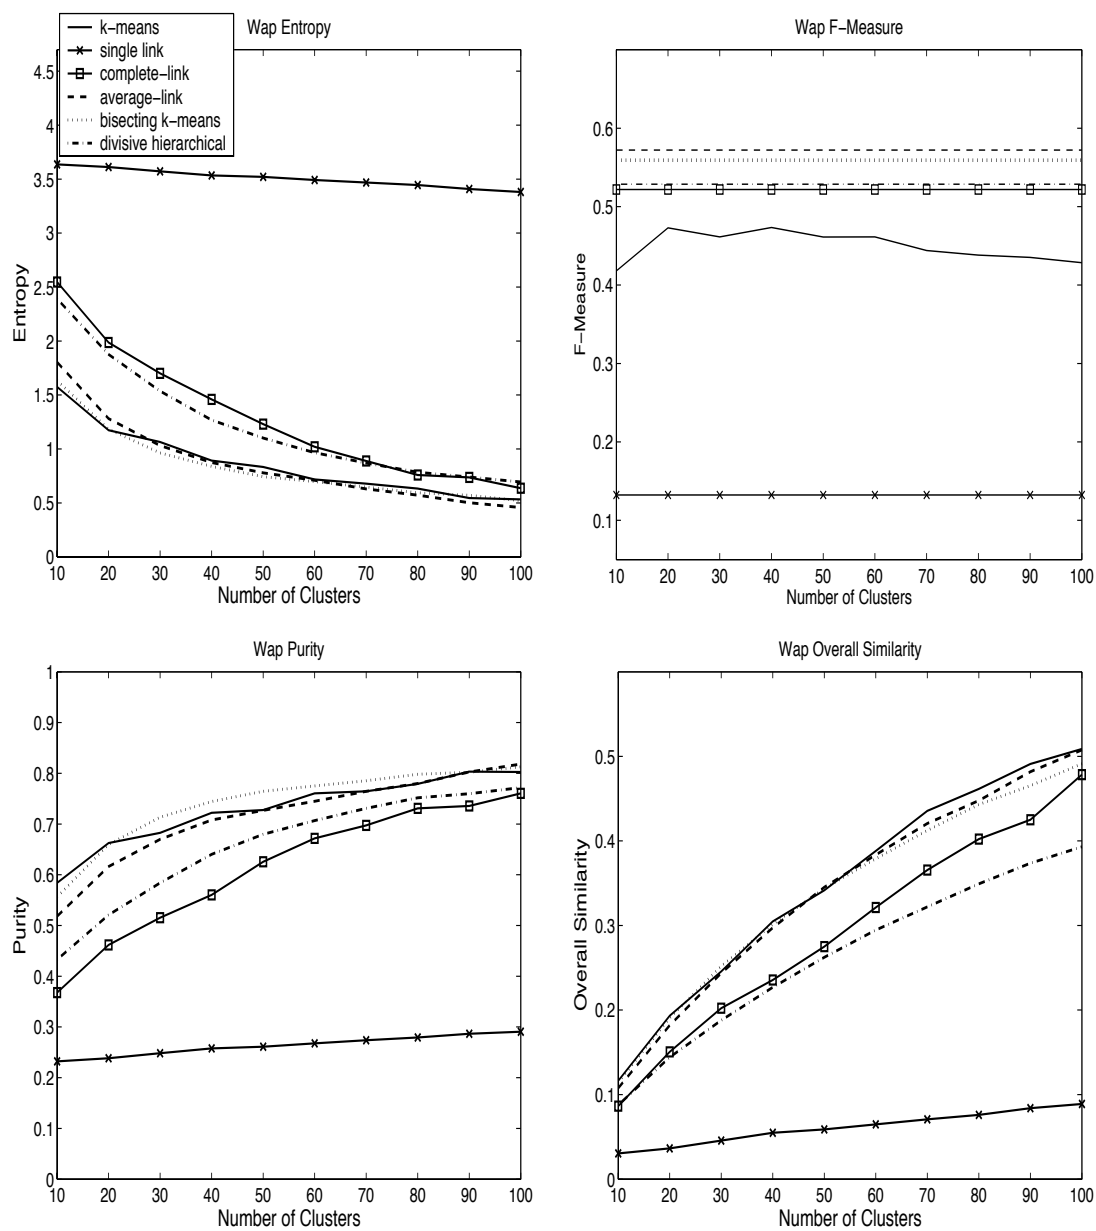


Figure 5.4. Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Wap data set

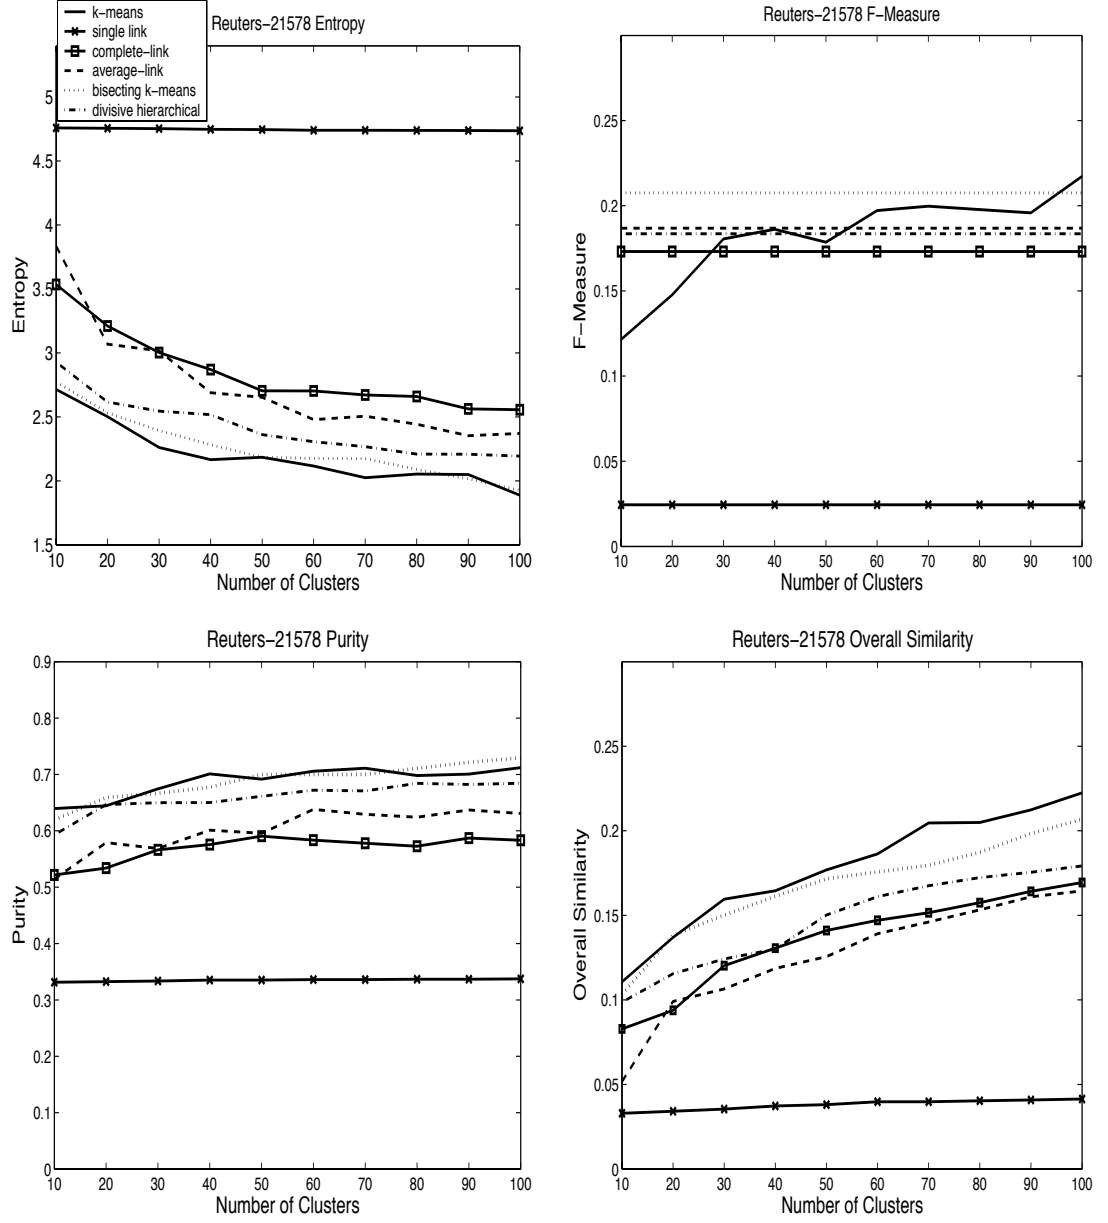


Figure 5.5. Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Reuters-21578 data set

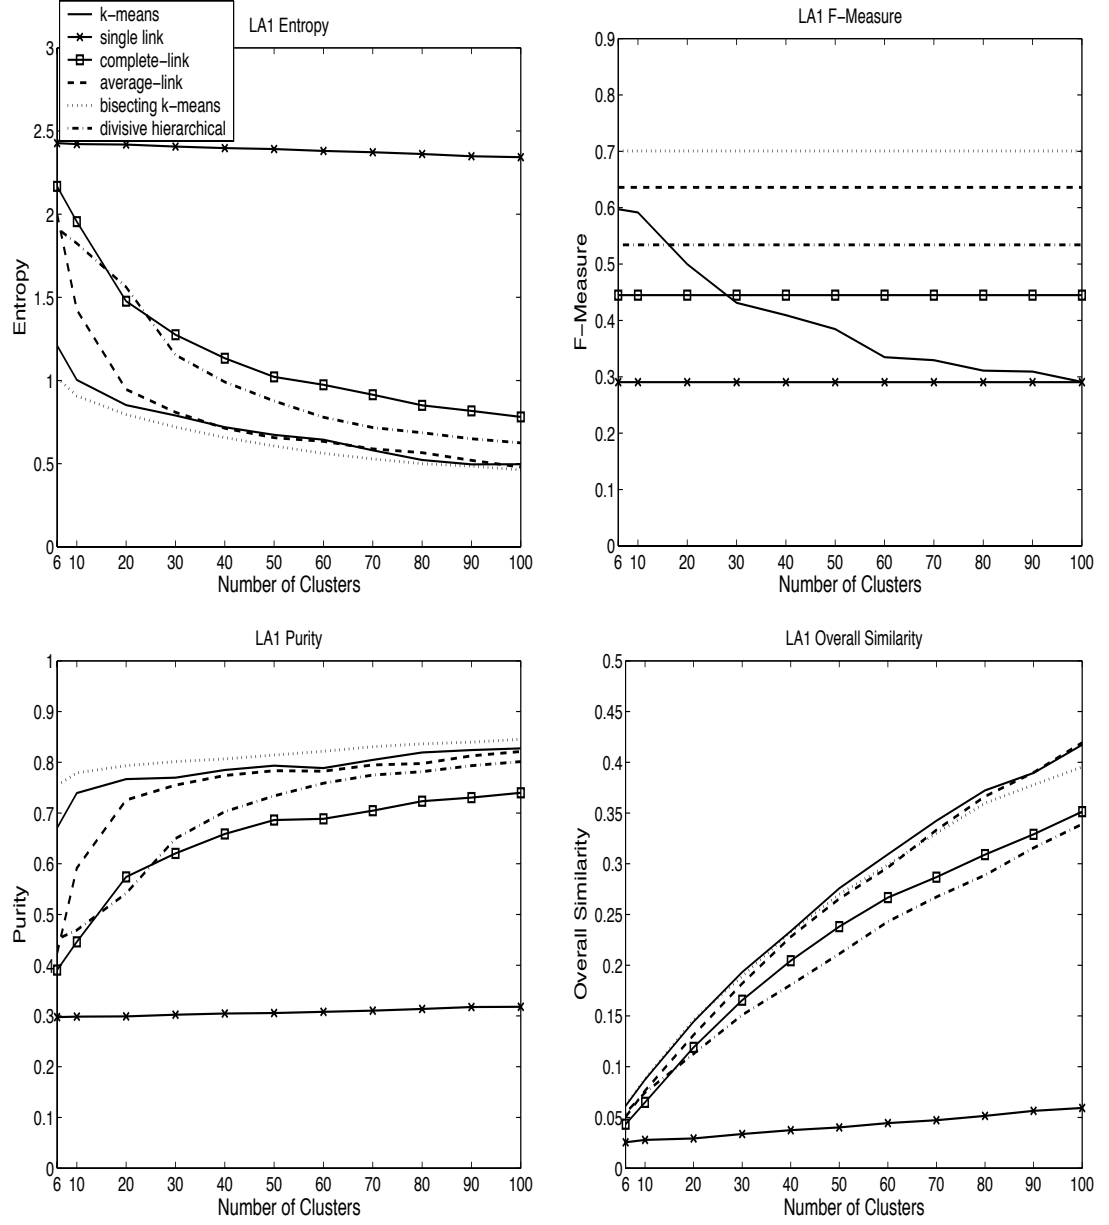


Figure 5.6. Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over LA1 data set

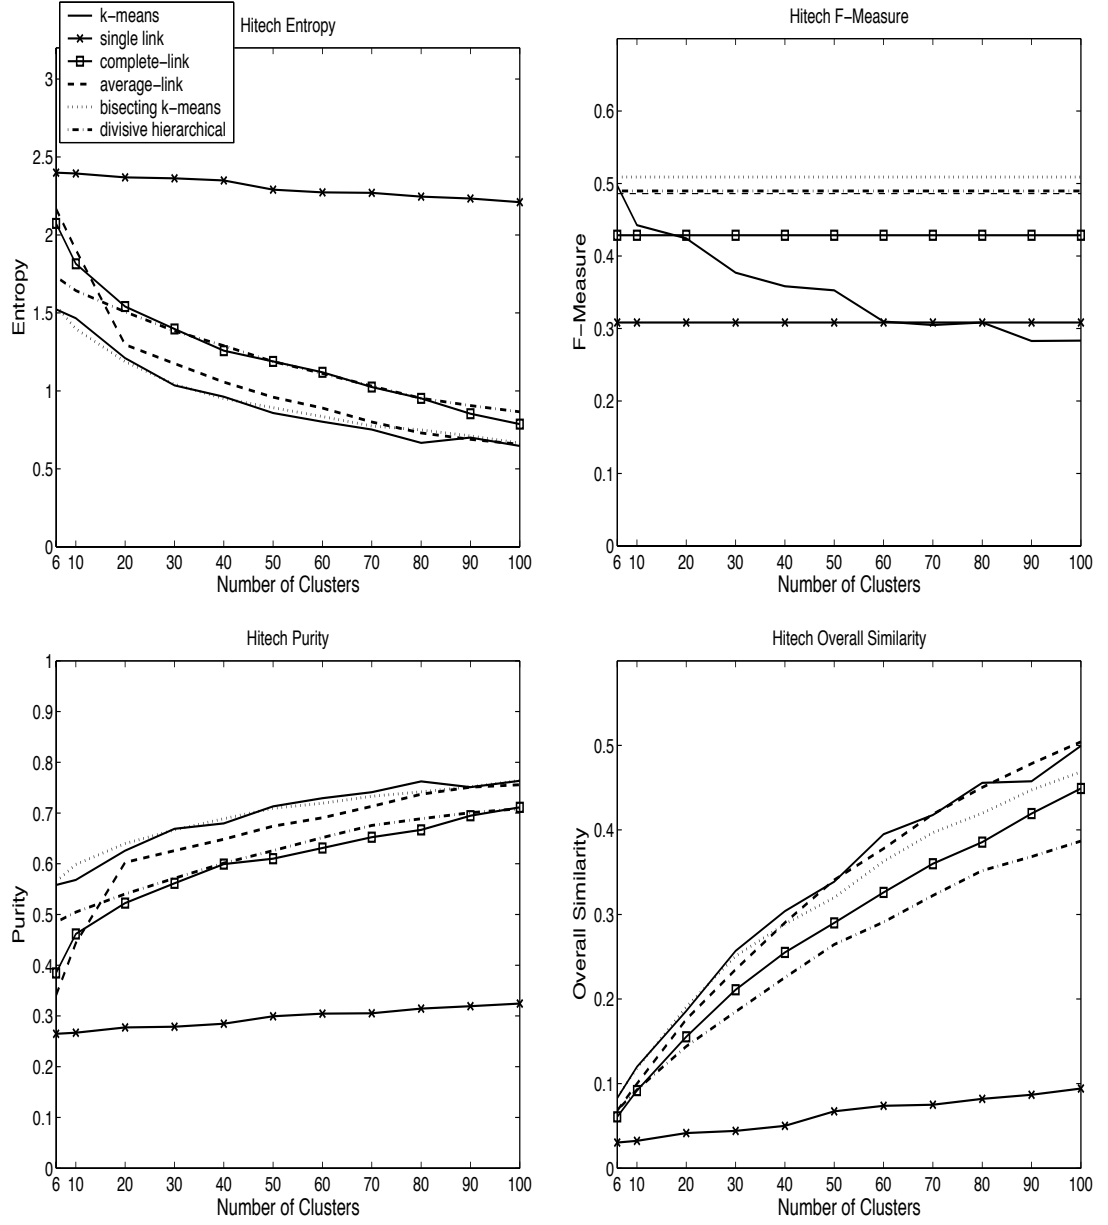


Figure 5.7. Comparison of entropy, F-measure, purity, and overall similarity values for online k -means, bisecting k -means, single-link, complete-link, average-link and divisive clustering algorithms over Hitech data set

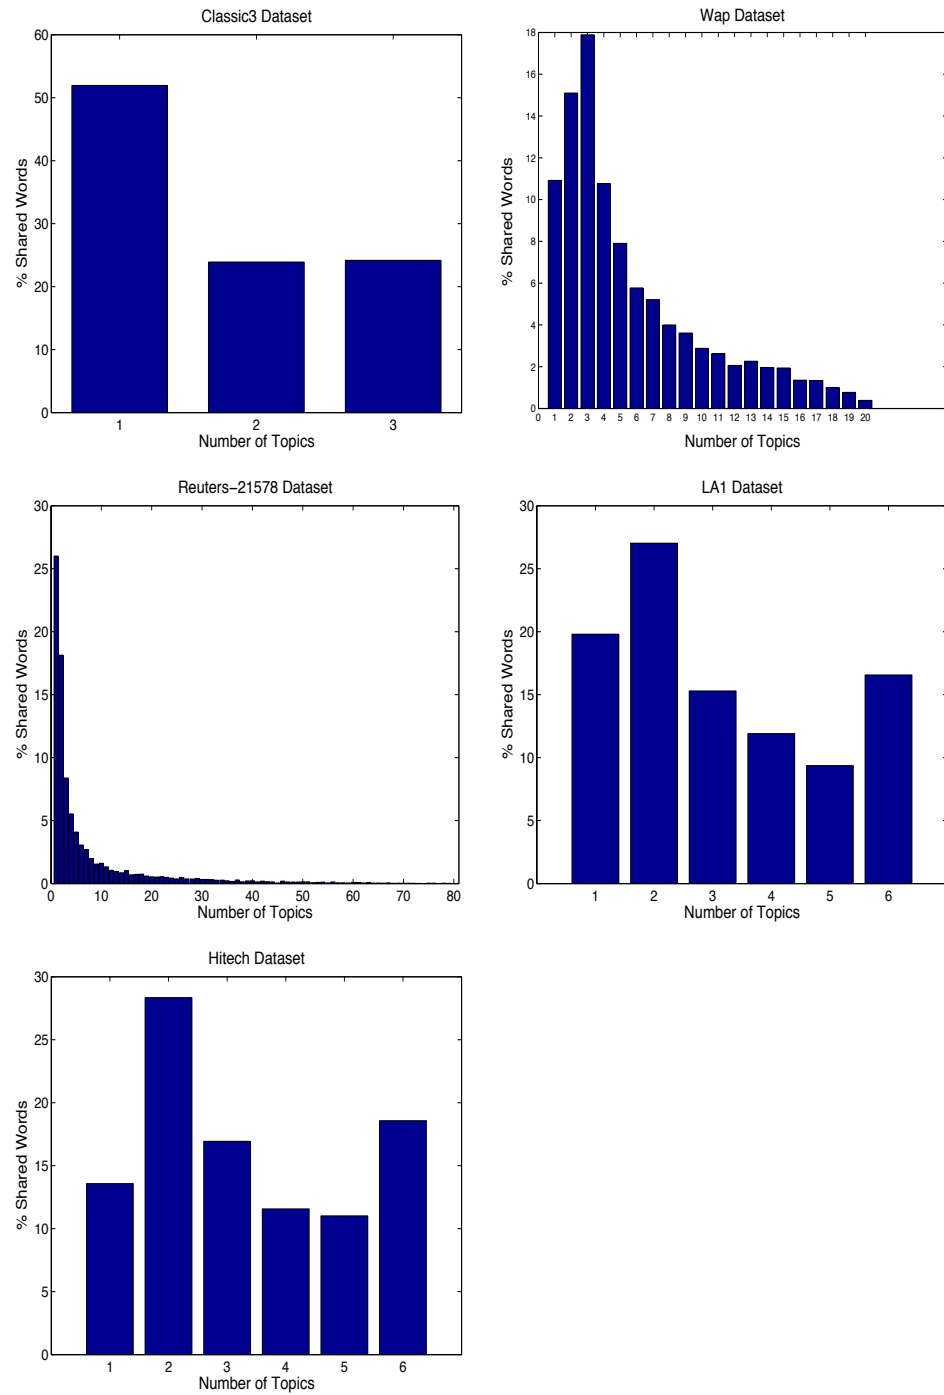


Figure 5.8. Term distributions of Classic3, Wap, Reuters-21578, LA1, and Hitech data sets

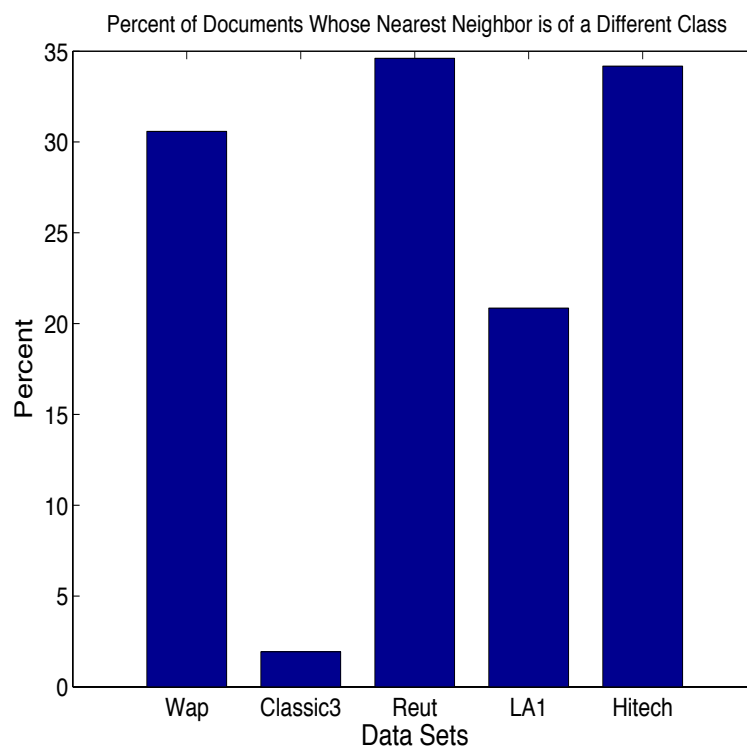


Figure 5.9. Per cent of documents in Classic3, Wap, Reuters-21578, LA1, and Hitech document collections whose nearest neighbor belongs to a different class

In Appendix A, we present the same statistics for Hitech, LA1, Reuters-21578, and Wap data sets as well. We display these statistics for number of clusters equal to number of topics for Hitech and LA1 data sets. As Reuters-21578 and Wap data sets contain large numbers of topics we present the statistics for ten clusters for these data sets so that they are more comprehensible. It can be said that the algorithms display similar behaviors for each data set. We can conclude that the wrong decisions taken by agglomerative and divisive hierarchical algorithms in the initial steps degrade their performance drastically, as they do not revisit intermediate clusters for the purpose of improving them, once they are constructed, in contrast to k -means and bisecting k -means. So, for the domain of text documents our results reveal that k -means and its hierarchical variant bisecting k -means are more appropriate.

Statistics of k-means (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.0398921 Overall F-measure = 0.982755							
Overall Entropy = 0.0946597 Overall Purity = 0.988692							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	1490	0.038	0.18	0.98	1455	17	18
1	1015	0.022	0.032	1	2	1	1012
2	1386	0.055	0.045	1	3	1380	3
CID Most Descriptive 5 Features							
0	librari: 8.3% inform: 5.8% system: 3.3% index: 2.1% document: 1.8%						
1	cell: 5.6% patient: 5.2% rat: 1.5% hormon: 1.3% growth: 1.3%						
2	flow: 4.3% boundari: 3.2% layer: 3.1% pressur: 2.6% wing: 2.2%						

Figure 5.10. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by *k*-means for Classic3 data set

Statistics of Bisecting k-means (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.039796 Overall F-measure = 0.97872 Overall							
Entropy = 0.118024 Overall Purity = 0.985094							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	1003	0.022	0.041	1	3	1	999
1	1397	0.054	0.11	0.99	0	1377	20
2	1491	0.038	0.18	0.98	1457	20	14
CID Most Descriptive 5 Features							
0	cell: 5.7% patient: 5.2% rat: 1.5% hormon: 1.3% growth: 1.3%						
1	flow: 4.4% boundari: 3.2% layer: 3.1% pressur: 2.8% wing: 2.1%						
2	librari: 8.3% inform: 5.7% system: 3.3% index: 2.1% document: 1.8%						

Figure 5.11. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Classic3 data set

Statistics of Divisive Hierarchical Clustering Alg. (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.0285019 Overall F-measure = 0.669711							
Overall Entropy = 1.13054 Overall Purity = 0.657929							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	18	0.17	0	1	18	0	0
1	1439	0.031	0.78	0.81	1166	22	251
2	2434	0.026	1.3	0.57	276	1376	782
CID Most Descriptive 5 Features							
0	bradford:	36%	journal:	9%	articl:	4%	edit: 3.7% refer: 3.4%
1	librari:	8%	inform:	4.9%	system:	4.2%	index: 2.4% document: 2%
2	flow:	3.2%	boundari:	2.2%	layer:	2.2%	pressur: 2.1% heat: 1.5%

Figure 5.12. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Classic3 data set

Statistics of Average Link (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.0271813 Overall F-measure = 0.741457							
Overall Entropy = 0.838165 Overall Purity = 0.627859							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	2864	0.029	1.1	0.5	1440	1393	31
1	1026	0.022	0.18	0.98	19	5	1002
2	1	1	0	1	1	0	0
CID Most Descriptive 5 Features							
0	librari: 2.9% inform: 2.1% flow: 2% boundari: 1.4% layer: 1.4%						
1	cell: 5.6% patient: 5.1% rat: 1.5% growth: 1.3% hormon: 1.3%						
2	warn: 30% discriminatori: 8.7% collat: 8% earli: 7.7% aacr: 7.5%						

Figure 5.13. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Classic3 data set

Statistics of Complete Link (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.0225543 Overall F-measure = 0.496991							
Overall Entropy = 1.53505 Overall Purity = 0.382935							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	7	0.21	0	1	7	0	0
1	85	0.11	0	1	85	0	0
2	3799	0.02	1.6	0.37	1368	1398	1033
CID Most Descriptive 5 Features							
0	taxonomi:	17%	dewei:	8.8%	domain:	6.7%	classif: 4.4% optim: 2%
1	librari:	52%	public:	3%	servic:	2.1%	educ: 1.7% school: 1.3%
2	flow:	1.8%	inform:	1.7%	librari:	1.3%	boundari: 1.2% system: 1.2%

Figure 5.14. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Classic3 data set

Statistics of Single Link (K=3) for Classic3 Dataset							
Number of: Documents = 3891 Topics = 3 Terms = 6729							
Overall Similarity = 0.0208245 Overall F-measure = 0.499871							
Overall Entropy = 1.56817 Overall Purity = 0.375482							
CID	Size	Sim	Entropy	Purity	cisi	cran	med
0	1	1	0	1	0	0	1
1	3	0.44	0.92	0.67	2	0	1
2	3887	0.02	1.6	0.38	1458	1398	1031
CID Most Descriptive 5 Features							
0	ceylon:	33%	vector:	16%	plasmodium:	15%	host: 5.6% natur: 4.4%
1	echo:	34%	delus:	13%	ancient:	3.7%	moral: 3.7% mental: 3.1%
2	librari:	2.3%	inform:	1.7%	flow:	1.7%	system: 1.2% boundari: 1.1%

Figure 5.15. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Classic3 data set

5.3. Evaluation of the Classification Techniques

5.3.1. Evaluation Metrics

To evaluate the performance of the supervised text classification techniques we use the commonly used F_1 -measure, which is first introduced by van Rijsbergen [45].

F_1 -measure is a special case of the more general F_β -measure defined as [1]:

$$F_\beta = \frac{(\beta^2 + 1)\pi\rho}{\beta^2\pi + \rho} \quad (5.11)$$

where π is precision and ρ is recall. β can take values in the range $[0, +\infty]$. It may be interpreted as the relative degree of importance given to π and ρ . If $\beta = 0$, $F_\beta = \pi$, whereas if $\beta = +\infty$ then $F_\beta = \rho$. F_1 -measure combines recall and precision with equal weight. It is equal to their harmonic mean:

$$F_1 = \frac{2\pi\rho}{\pi + \rho} \quad (5.12)$$

Precision for class i is the ratio of correct assignments by the system divided by the total number of assignments by the system to class i , whereas recall for class i is the ratio of correct assignments by the system to class i divided by the total number of correct assignments:

$$\pi_i = \frac{TP_i}{TP_i + FP_i} \quad (5.13)$$

$$\rho_i = \frac{TP_i}{TP_i + FN_i} \quad (5.14)$$

Here, TP_i (True Positives) is the number of documents assigned correctly to class i ; FP_i (False Positives) is the number of documents that do not belong to class i but are assigned to class i incorrectly by the classifier; and FN_i (False Negatives) is the number of documents that are not assigned to class i by the classifier but which actually belong to class i .

The overall F_1 measure score of the entire classification problem can be computed by two different types of averages, *micro-average* and *macro-average*.

5.3.1.1. Micro-averaged F_1 -measure . In micro-averaging, F_1 -measure is computed globally over all category decisions. ρ and π are obtained by summing over all individual decisions:

$$\pi = \frac{TP}{TP + FP} = \frac{\sum_{i=1}^M TP_i}{\sum_{i=1}^M (TP_i + FP_i)} \quad (5.15)$$

$$\rho = \frac{TP}{TP + FN} = \frac{\sum_{i=1}^M TP_i}{\sum_{i=1}^M (TP_i + FN_i)} \quad (5.16)$$

where M is the number of categories. Micro-averaged F_1 -measure is then computed as:

$$F_1(\text{micro-averaged}) = \frac{2\pi\rho}{\pi + \rho} \quad (5.17)$$

The micro-averaged F_1 -measure gives equal weight to each document . It is therefore considered as an average over all the document/category pairs. It is discussed by Yang and Liu in [8] that micro-averaged F_1 -measure tends to be dominated by the classifiers' performance on common categories.

5.3.1.2. Macro-averaged F_1 -measure . In macro-averaging, F_1 -measure is computed locally over each category first and then the average over all categories is taken. π and ρ are computed for each category as in equations 5.13 and 5.14. Then F_1 -measure for each category i is computed as:

$$F_{1_i} = \frac{2\pi_i\rho_i}{\pi_i + \rho_i} \quad (5.18)$$

The macro-averaged F_1 -measure is obtained by taking the average of F_1 -measure values for each category.

$$F_1(\text{macro-averaged}) = \frac{\sum_{i=1}^M F_{1_i}}{M} \quad (5.19)$$

where M is total number of categories. Macro-averaged F_1 -measure gives equal weight to each category, regardless of its frequency. It is discussed by Yang and Liu in [8] that macro-averaged F_1 -measure is influenced more by the classifiers performance on rare categories.

So, we provide both measurement scores to be more informative.

5.3.2. Results and Discussion

In this section we present the micro-averaged and macro-averaged F_1 -measure results of naive Bayes classifier relying on multivariate Bernoulli model, traditional k -NN, weighted k -NN, and SVM.

We used SVM^{light} package, version 5.00 [40], to train and test the SVM classifier. We performed experiments on SVM with linear kernel; polynomial kernel of degrees 2, 3, 4, 5, 6, and 7; and RBF kernel with variances 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, and 1.4. Tables 5.2 and 5.3 display micro-averaged and macro-averaged F_1 -measure results, respectively, for SVM with polynomial kernel. Tables 5.4 and 5.5 display micro-averaged and macro-averaged F_1 -measure results, respectively, for SVM with RBF kernel.

Table 5.2. Micro-averaged F_1 -measure results for SVM with polynomial kernel of degrees 1, 2, 3, 4, 5, 6, and 7

Data set	1	2	3	4	5	6	7
Classic3	0.970	0.994	0.994	0.994	0.994	0.992	0.991
Hitech	0.728	0.736	0.727	0.719	0.712	0.695	0.674
LA1	0.884	0.882	0.882	0.874	0.832	0.831	0.786
Reuters-21578	0.823	0.813	0.802	0.780	0.756	0.710	0.659
Wap	0.839	0.838	0.838	0.822	0.802	0.770	0.702

Table 5.3. Macro-averaged F_1 -measure results for SVM with polynomial kernel of degrees 1, 2, 3, 4, 5, 6, and 7

Data set	1	2	3	4	5	6	7
Classic3	0.968	0.994	0.994	0.994	0.994	0.992	0.991
Hitech	0.662	0.666	0.641	0.627	0.608	0.576	0.555
LA1	0.847	0.843	0.839	0.825	0.766	0.764	0.708
Reuters-21578	0.503	0.486	0.474	0.461	0.446	0.422	0.412
Wap	0.662	0.658	0.657	0.617	0.583	0.546	0.491

Table 5.4. Micro-averaged F_1 -measure results for SVM with RBF kernel with $\gamma = 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, 1.4$

Data set	0.2	0.4	0.6	0.8	1.0	1.2	1.4
Classic3	0.993	0.994	0.994	0.994	0.993	0.993	0.993
Hitech	0.736	0.731	0.725	0.721	0.721	0.716	0.711
LA1	0.882	0.878	0.876	0.873	0.864	0.857	0.849
Reuters-21578	0.814	0.806	0.794	0.777	0.758	0.733	0.704
Wap	0.839	0.834	0.832	0.821	0.813	0.801	0.781

Table 5.5. Macro-averaged F_1 -measure results for SVM with RBF kernel with $\gamma = 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, 1.4$

Data set	0.2	0.4	0.6	0.8	1.0	1.2	1.4
Classic3	0.994	0.994	0.994	0.994	0.993	0.993	0.993
Hitech	0.667	0.653	0.636	0.619	0.616	0.609	0.598
LA1	0.840	0.833	0.828	0.823	0.806	0.795	0.784
Reuters-21578	0.485	0.464	0.433	0.412	0.388	0.361	0.337
Wap	0.662	0.642	0.637	0.609	0.593	0.577	0.551

In the training phase the parameter C , which is the trade-off between training error and margin, is left to its default value. In the SVM^{light} system, the default value of C is defined to be equal to the reciprocal of the average Euclidean norm of training examples. We did not perform experiments with different C values as the results we obtained are satisfactory. From the results we can say that generally linear kernel or polynomial kernel of degree 2 lead to better performance. However, performance does not change drastically with the complexity of the model. Effect of the degree of the polynomial kernel and the variance of the RBF kernel of SVM to the F_1 -measure scores are displayed in Figures 5.16 and 5.17 respectively.

Tables 5.6 and 5.7 display the micro-averaged and macro-averaged F_1 -measure results of the supervised techniques we evaluated. For traditional k -NN and weighted k -NN we performed experiments for k values starting with 5 and incrementing by 5 till $k = 60$. We report the results for the optimum k value, and report that k value in parentheses as well. Likewise, for SVM we report the best results obtained and the model leading to those results in parentheses.

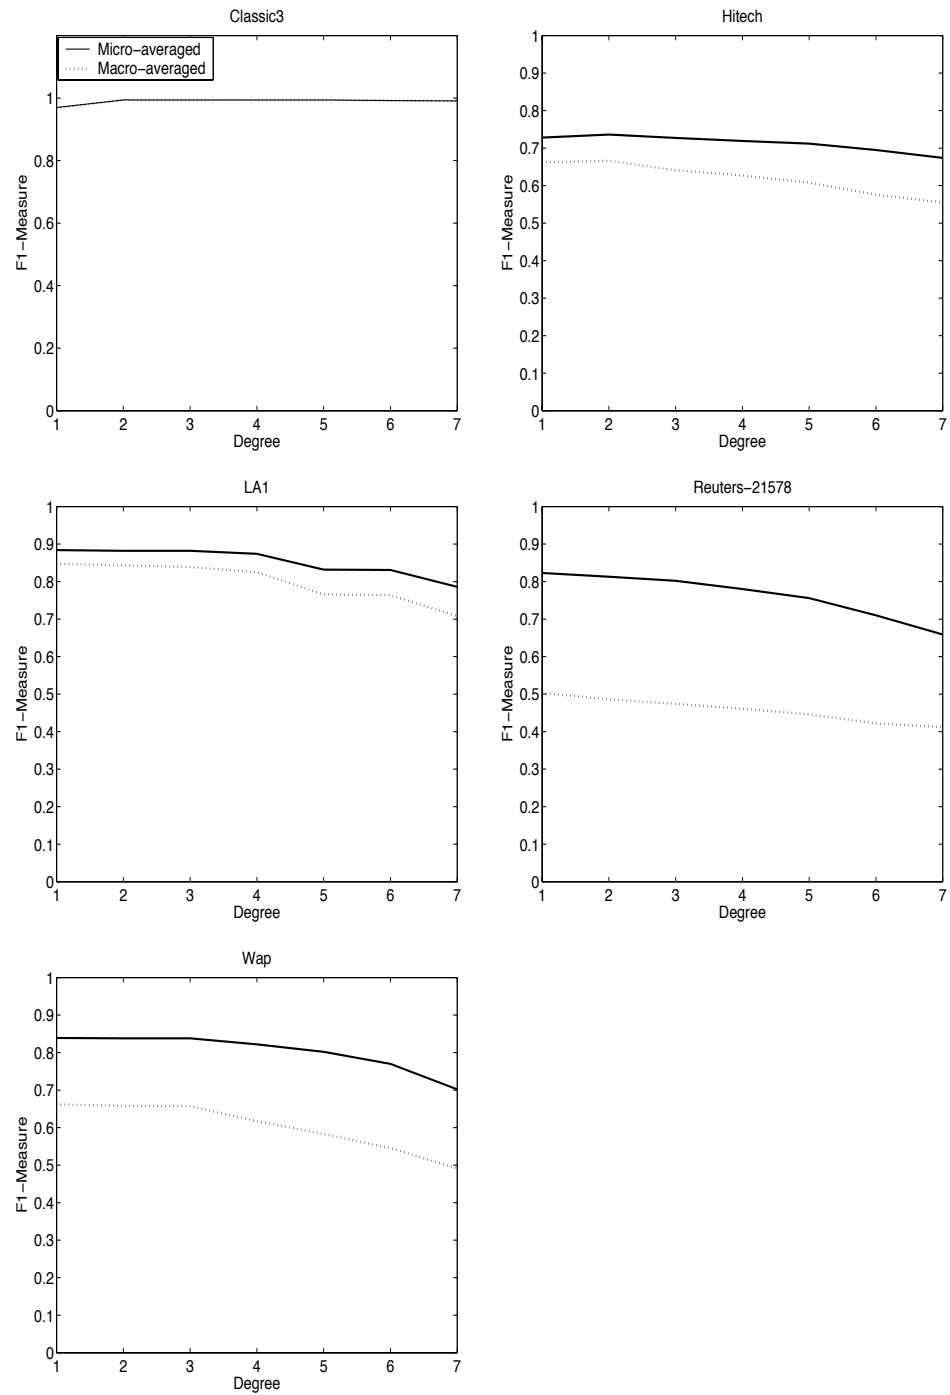


Figure 5.16. Effect of the degree of the polynomial kernel of SVM to the F_1 -measure scores

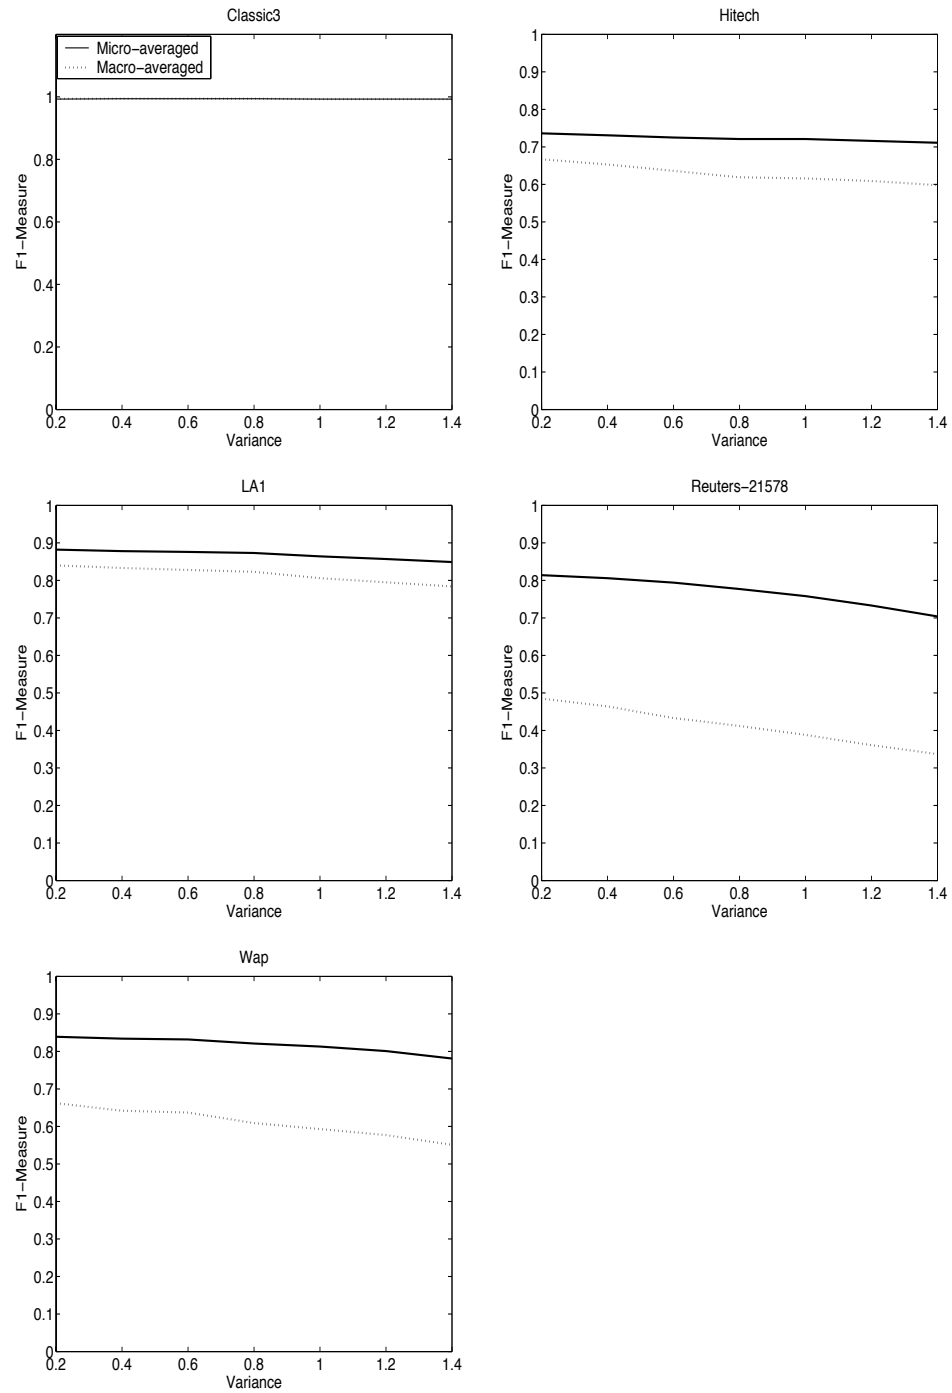


Figure 5.17. Effect of the variance of the RBF kernel of SVM to the F_1 -measure scores

Table 5.6. Micro-averaged F_1 -measure results for the supervised techniques

Data set	naive Bayes	k -NN (Trad.)	k -NN (Weigh.)	SVM
Classic3	0.987	0.987 (k=35)	0.989 (k=45)	0.994 (d=2)
Hitech	0.533	0.668 (k=25)	0.693 (k=20)	0.736 (d=2) ($\gamma=0.2$)
LA1	0.574	0.835 (k=20)	0.844 (k=50)	0.884 (d=1)
Reuters-21578	0.625	0.757 (k=30)	0.769 (k=30)	0.850 (d=1)
Wap	0.623	0.768 (k=35)	0.785 (k=45)	0.839 (d=1)

Table 5.7. Macro-averaged F_1 -measure results for the supervised techniques

Data set	naive Bayes	k -NN (Trad.)	k -NN (Weigh.)	SVM
Classic3	0.986	0.987 (k=35)	0.989 (k=45)	0.994 (d=2)
Hitech	0.383	0.585 (k=10)	0.612 (k=20)	0.667 ($\gamma=0.2$)
LA1	0.476	0.783 (k=20)	0.794 (k=20)	0.847 (d=1)
Reuters-21578	0.200	0.416 (k=30)	0.499 (k=30)	0.430 (d=1)
Wap	0.354	0.590 (k=5)	0.620 (k=15)	0.662 (d=1)

We can conclude that, for the relatively easy data set Classic3, each algorithm performs well and similarly. However, for the remaining data sets, in terms of micro-averaged F_1 -measure and macro-averaged F_1 -measure NB performs the worst and SVM performs the best. The performance of weighted k -NN is slightly better than traditional k -NN and comparable to SVM. We can also observe that, no matter what the underlying model is, SVM either outperforms the other algorithms or is comparable with them.

6. COMPARISON OF THE SUPERVISED AND UNSUPERVISED METHODS

Comparison of the unsupervised techniques [18] and comparison of the supervised techniques [8] have been performed in the literature. However to our knowledge the two approaches have not been compared with each other. In this chapter we compare the quality of the clusters produced by the two approaches. The classification solutions obtained by the supervised techniques are considered as if they were clustering solutions and they are evaluated by using the evaluation metrics that are used to evaluate the performance of the clustering algorithms. The clustering algorithms on the other hand are ran for number of clusters equal to the number of pre-defined categories in each document corpus. Tables 6.1, 6.2, 6.3, 6.4, and 6.5 display the results for Classic3, Hitech, LA1, Reuters-21578, and Wap data sets respectively. We do not present the results for the single-link algorithm as it is shown to be a very bad performer in the previous chapter.

Table 6.1. Quality of the clusters ($k=3$) obtained by the unsupervised and the supervised techniques for Classic3

Algorithm	Entropy	Purity	Overall Similarity	F-Measure
average-link	0.574	0.767	0.047	0.817
bisecting k -means	0.098	0.987	0.055	0.983
complete-link	1.423	0.455	0.032	0.520
divisive	0.503	0.855	0.051	0.865
k -means	0.093	0.987	0.055	0.984
k -NN (Traditional)	0.093	0.987	0.055	0.983
k -NN (Weighted)	0.078	0.989	0.055	0.986
NB	0.091	0.987	0.055	0.983
SVM	0.047	0.994	0.055	0.991

Table 6.2. Quality of the clusters ($k=6$) obtained by the unsupervised and the supervised techniques for Hitech

Algorithm	Entropy	Purity	Overall Similarity	F-Measure
average-link	2.167	0.341	0.068	0.308
bisecting k -means	1.544	0.565	0.084	0.483
complete-link	2.067	0.385	0.061	0.356
divisive	1.732	0.485	0.068	0.453
k -means	1.523	0.558	0.083	0.497
k -NN (Traditional)	1.394	0.670	0.078	0.593
k -NN (Weighted)	1.308	0.695	0.079	0.626
NB	1.832	0.536	0.061	0.485
SVM	1.181	0.738	0.079	0.666

Table 6.3. Quality of the clusters ($k=6$) obtained by the unsupervised and the supervised techniques for LA1

Algorithm	Entropy	Purity	Overall Similarity	F-Measure
average-link	2.000	0.424	0.050	0.372
bisecting k -means	1.016	0.755	0.062	0.674
complete-link	2.167	0.391	0.043	0.344
divisive	1.916	0.451	0.055	0.410
k -means	1.210	0.671	0.061	0.597
k -NN (Traditional)	0.831	0.835	0.062	0.763
k -NN (Weighted)	0.781	0.845	0.062	0.773
NB	1.660	0.574	0.049	0.525
SVM	0.618	0.884	0.062	0.816

Table 6.4. Quality of the clusters ($k=90$) obtained by the unsupervised and the supervised techniques for Reuters-21578

Algorithm	Entropy	Purity	Overall Similarity	F-Measure
average-link	2.353	0.637	0.161	0.171
bisecting k -means	2.016	0.721	0.198	0.188
complete-link	2.563	0.587	0.164	0.164
divisive	2.209	0.682	0.175	0.178
k -means	2.049	0.701	0.212	0.196
k -NN (Traditional)	2.200	0.739	0.108	0.247
k -NN (Weighted)	2.092	0.748	0.115	0.264
NB	3.267	0.627	0.067	0.111
SVM	1.982	0.879	0.192	0.297

Table 6.5. Quality of the clusters ($k=20$) obtained by the unsupervised and the supervised techniques for Wap

Algorithm	Entropy	Purity	Overall Similarity	F-Measure
average-link	1.282	0.616	0.182	0.453
bisecting k -means	1.183	0.660	0.191	0.467
complete-link	1.987	0.462	1.151	0.332
divisive	1.876	0.521	0.088	0.350
k -means	1.173	0.662	0.193	0.473
k -NN (Traditional)	1.026	0.769	0.142	0.580
k -NN (Weighted)	0.954	0.787	0.151	0.597
NB	1.552	0.658	0.119	0.463
SVM	0.706	0.846	0.165	0.652

Different from the unsupervised techniques, the supervised techniques use class label information in addition to the similarity information between documents. For this reason, it is expected that the clusters (groups) obtained by the supervised techniques are of higher quality compared to the unsupervised techniques. However, we can observe from the results that the best performers of the unsupervised techniques k -means and bisecting k -means achieve generally better performance than NB and not much worse performance than k -NN, which are supervised techniques, in terms of entropy, purity, overall similarity and F -measure. SVM achieves the highest performance in terms of entropy, purity, and F -measure. Another observation is that, compared with the supervised techniques the unsupervised techniques generally achieve higher overall similarity performance. This is due to the fact that they make decisions depending only on the similarity information between documents. On the other hand the supervised techniques use a labelled training set. This observation has made us think that there may be some outliers in the labelled training set that lead to decrease in the overall similarity of the clusters obtained and unsupervised techniques can be used to enhance the task of pre-defining categories and labelling documents in the training set.

7. CONCLUSIONS AND FUTURE WORK

The advances in technology of computers and electronics, the increasing popularity of the Internet and the WWW have lead to vast amounts of increase in electronic text information. In order to browse text databases and extract relevant information fast, efficiently, and correctly organizing digital text documents automatically has become an important research issue. There are two major machine learning approaches for document organization, supervised and unsupervised. In this study we present an experimental evaluation of these two major paradigms.

We discuss, different alternatives for document representation and explain our methodology in Chapter 2. We chose to use the “bag-of-words” representation, where each distinct stemmed word is defined as a term. Stopwords and mark-up tags are removed, words are stemmed by using Porter’s stemming algorithm, and dimensionality is reduced by the unsupervised DF Thresholding technique.

After the preprocessing phase, we implement and apply the leading unsupervised and supervised algorithms to five standard document corpora. Among the unsupervised algorithms, we evaluate the most popular partitionial clustering technique k -means, its variant bisecting k -means, the main agglomerative hierarchical clustering techniques; single-link, complete-link, and average-link, and divisive hierarchical clustering technique which to our knowledge has not been used to cluster documents previously. We conclude that among the hierarchical clustering algorithms the average-link algorithm achieves the best performance and the divisive algorithm achieves similar performance. The reason for the worse performance of the single-link and the complete-link algorithms is that, they depend on assumptions far from reality for the nature of text document collections. The single-link algorithm assumes that nearest neighbors belong to the same class and the complete-link algorithm assumes that documents in a cluster are very similar to each other. We presented a discussion in Chapter 5 that documents may share many common terms and be close to each other, even be nearest neighbors but belong to different categories. These properties are the ones that make

the task of document organization challenging. In the domain of text documents, the number of documents, N , is usually very large. In this case, k -means and bisecting k -means can be more favorable than agglomerative and divisive hierarchical clustering algorithms as they have $O(N)$ time complexity in contrast to the $O(N^2)$ time complexity of agglomerative and divisive hierarchical clustering algorithms. On the other hand, performance of k -means and bisecting k -means depends very much on the parameter k and the initial selection of centroids. In Chapter 5 and in Appendix A we illustrate the performance of the unsupervised techniques over the data sets by presenting some statistics where we observe that k -means and bisecting k -means usually better discriminate between the actual classes, while the hierarchical techniques may lead to rather inhomogeneous and unbalanced clusters. Especially the single-link algorithm has a tendency to suffer from the chaining effect.

We evaluate the performance of multivariate Bernoulli model of NB, traditional k -NN, weighted k -NN, and SVM supervised techniques. Our results reveal that, NB achieves the lowest performance, while SVM performs the best. Weighted k -NN achieves comparable performance to SVM, and traditional k -NN performs slightly worse than weighted k -NN. The reasons for the poor performance of NB may be the boolean vector representation of the documents and the unrealistic naive assumption that the probability of each term being present in a document is only dependent on the category of the documents and independent from the presence of other terms in the document.

As a conclusion we can say that, for unsupervised document clustering k -means and its variant bisecting k -means are more appropriate than the agglomerative and divisive hierarchical clustering techniques both in terms of time complexity and the quality of the clusters produced. Agglomerative and divisive hierarchical clustering algorithms generally produce unbalanced, inhomogeneous clusters. Although divisive hierarchical clustering is not as common as agglomerative approach it does not produce worse results. On the other hand for supervised document categorization the method with best performance is SVM and NB approach performs poorly.

In Chapter 6 we compared the unsupervised techniques with the supervised techniques in terms of the quality of the clusters they produce. We considered the classification solutions obtained by the supervised techniques as if they were clustering solutions and likewise evaluated the unsupervised algorithms for number of clusters equal to the pre-defined number of categories for each data set. From these experiments, different from our expectations we conclude that although k -means and bisecting k -means are unsupervised techniques they produce clusters of better quality than NB, and not much worse than k -NN and SVM. In addition, clusters produced by unsupervised techniques have generally greater overall similarity than the clusters produced by the supervised techniques. The reason may be due to outliers in the labelled documents. For the supervised techniques to be applied to document categorization, categories should be pre-defined and there should be a training set of labelled documents. We discuss that defining the categories in advance and preparing a labelled training set is a challenging, error prone, and subjective task especially in dynamic text environments such as the WWW. We discuss the inter-indexer inconsistency phenomenon where two different indexers may decide to label a document under different categories, depending on their subjective opinion. Therefore, we suggest to use unsupervised clustering before applying supervised classification to enhance pre-definition of categories and preparation of labelled training set. Given a document corpus a clustering solution can be obtained and the human indexer can be presented with most descriptive terms of each cluster representative as a suggestion for keywords of the cluster (category). The indexer can be informed about the documents in each cluster which have relatively small average pairwise similarity to the other documents in the cluster. This small average pairwise similarity may be an indication of overlapping clusters or outlier documents in the cluster.

As future work, we will work on a hybrid approach to organize text documents that incorporates the strengths of supervised and unsupervised paradigms. This approach will start in an unsupervised manner without pre-defined categories and labelled data and in a later phase will incorporate the supervised approach. Further future plans are to integrate context knowledge such as word synonyms, hypernyms, hyponyms, and phrases into the supervised and unsupervised techniques; to evaluate the techniques

for documents in Turkish and study the effects of language in the performance of the supervised and unsupervised techniques; and to gather the source codes we have implemented into a publicly available toolkit for document preprocessing and representation; classification, and clustering.

APPENDIX A: STATISTICS OF THE CLUSTERING ALGORITHMS

Statistics of K-Means Algorithm (K=6) for Hitech Dataset										
Number of: Documents=1530 Topics = 6 Terms = 10919										
Overall Similarity = 0.0436829 Overall F-measure = 0.498438										
Overall Entropy = 1.64171 Overall Purity = 0.581699										
Topics: 0:computer 1:electronics 2:medical 3:health 4:research										
5:technology										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	423	0.042	1.8	0.44	12	12	186	153	50	10
1	78	0.084	1.4	0.73	2	4	4	2	57	9
2	431	0.045	1.4	0.69	298	52	7	4	15	55
3	256	0.034	1.4	0.68	4	3	43	174	26	6
4	300	0.026	1.9	0.54	6	5	33	65	161	30
5	42	0.15	1.9	0.33	0	0	13	4	11	14
CID Most Descriptive 5 Features										
0	health:4.9%		aid:4.8%		care:4.5%		patient:3.2%		insur:2.3%	
1	space: 14%		telescop:5.8%		nasa:5.4%		astronaut:3.3%		star:3.2%	
2	comput: 10%		compani:4.7%		ibm:3.2%		appl:2.1%		quarter:1.7%	
3	exercis:3.3%		bush:2.6%		women:2.4%		heart:2.3%		studi:1.1%	
4	cancer:2.3%		studi: 2%		scientist:1.9%		anim:1.8%		research:1.7%	
5	hambrecht: 17%		quist: 17%		index: 15%		nurs:7.9%		growth:6.7%	

Figure A.1. The performance, cluster-class distribution, and most descriptive 5 features of the clustering solution obtained by k -means for Hitech

Statistics of Bisecting K-Means Algorithm (K=6) for Hitech Dataset										
Number of: Documents = 1530 Topics = 6 Terms = 10919										
Overall Similarity = 0.0427151 Overall F-measure = 0.452209										
Overall Entropy = 1.70651 Overall Purity = 0.532026										
Topics: 0:computer 1:electronics 2:medical 3:health 4:research										
5:technology										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	333	0.045	1.4	0.69	231	25	3	5	16	53
1	428	0.042	1.7	0.45	5	1	193	155	52	22
2	112	0.055	1.6	0.68	4	9	8	6	76	9
3	162	0.036	2.1	0.37	7	3	32	51	60	9
4	363	0.027	1.8	0.5	4	4	44	183	109	19
5	132	0.082	1.8	0.54	71	34	6	2	7	12
CID Most Descriptive 5 Features										
0	comput: 13%	ibm:4.8%	appl:3.3%	compani: 3.1%	system:1.3%					
1	health:4.8%	aid:4.6%	care:4.4%	patient: 3.4%	insur:2.2%					
2	space:8.2%	farlei:6.6%	telescop:4.5%	nasa: 3.9%	earth:3.1%					
3	cancer:6.4%	bush:5.5%	gene:3.3%	cell: 3.1%	diseas:2.2%					
4	studi:2.2%	exercis:1.9%	women:1.3%	water: 1.2%	research:0.85%					
5	quarter:8.6%	stock:5.4%	compani: 5%	million: 4.5%	share:3.5%					

Figure A.2. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Hitech data set

Statistics of Divisive Hierarchical Algorithm (K=6) for Hitech Dataset										
Number of: Documents = 1530 Topics = 6 Terms = 10919										
Overall Similarity = 0.0321126 Overall F-measure = 0.354704										
Overall Entropy = 2.10168 Overall Purity = 0.401307										
Topics: 0:computer 1:electronics 2:medical 3:health 4:research 5:technology										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	988	0.024	2.4	0.32	318	76	166	125	186	117
1	490	0.032	1.5	0.54	0	0	99	266	121	4
2	5	0.28	0.72	0.8	0	0	1	0	4	0
3	13	0.16	2	0.31	2	0	0	4	4	3
4	32	0.15	1.3	0.63	0	0	20	7	5	0
5	2	0.73	0	1	2	0	0	0	0	0
CID Most Descriptive 5 Features										
0	comput:4.1%		compani:2.3%		ibm:1.2%		million:1.1%		appl:0.8%	
1	aid: 4%		health:3.5%		patient: 3%		care:2.7%		infect:1.7%	
2	bird: 24%		dinosaur: 12%		parrot: 12%		hunter:6.4%		ancestor: 3%	
3	fax: 27%		label: 19%		wine:4.8%		printer:3.9%		grape:2.5%	
4	suicid: 18%		kevorkian: 18%		casolaro: 4%		banoff:3.3%		sutherland:3.2%	
5	novell: 59%		banyan: 20%		microsoft: 2%		digit:1.8%		acquisit:0.91%	

Figure A.3. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Hitech data set

Statistics of Average-link Agglomerative Hierarchical Algorithm
(K=6) for Hitech Dataset

Number of: Documents = 1530 Topics = 6 Terms = 10919

Overall Similarity = 0.0337664 Overall F-measure = 0.28479

Overall Entropy = 2.36008 Overall Purity = 0.277124

Topics : 0:computer 1:electronics 2:medical 3:health 4:research
5:technology

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	1496	0.022	2.4	0.26	321	76	276	396	317	110
1	16	0.22	0.95	0.63	0	0	10	6	0	0
2	1	1	0	1	1	0	0	0	0	0
3	2	0.56	0	1	0	0	0	0	2	0
4	14	0.89	0	1	0	0	0	0	0	14
5	1	1	0	1	0	0	0	0	1	0

CID Most Descriptive 5 Features

0	comput: 2%	compani:1.3%	health:1.1%	aid:0.93%	patient:0.88%
1	cicippio:21%	hostag: 13%	traci: 9%	sutherland:8.4%	steen: 4.9%
2	soul:37%	resign: 30%	kai: 11%	geoffrei:2.3%	reviv: 2%
3	tick:46%	lyme: 25%	diseas: 5%	fair:0.85%	waysid:0.79%
4	hambrecht:27%	quist: 27%	index: 23%	growth: 10%	fridai: 6%
5	spider:87%	estrada: 3.7%	insect:0.93%	guest:0.52%	woman:0.51%

Figure A.4. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Hitech data set

Statistics of Complete-link Agglomerative Hierarchical Algorithm
(K=6) for Hitech Dataset

Number of: Documents = 1530 Topics = 6 Terms = 10919

Overall Similarity = 0.0301916 Overall F-measure = 0.28724

Overall Entropy = 2.27326 Overall Purity = 0.322876

Topics: 0:computer 1:electronics 2:medical 3:health 4:research
5:technology

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	59	0.089	0.9	0.85	50	1	2	0	3	3
1	37	0.066	1.5	0.59	0	0	5	9	22	1
2	47	0.082	1.7	0.45	21	17	0	1	2	6
3	24	0.12	2	0.54	1	3	1	4	13	2
4	55	0.056	2.2	0.33	7	0	6	18	17	7
5	1308	0.022	2.4	0.28	243	55	272	370	263	105

CID Most Descriptive 5 Features

0	comput: 10%	pc:5.2%	disk:4.5%	box:3.9%	program:2.2%
1	anim:9.4%	speci:4.6%	bee:3.7%	collagen:3.1%	bird:2.3%
2	stock: 12%	electron:5.6%	trade:3.9%	exchang:2.4%	compani:2.4%
3	quak: 11%	earthquak:8.1%	fault:7.1%	volcano:3.3%	erupt:2.9%
4	pyramid: 7%	nuclear:6.5%	lab:2.7%	ey:2.4%	uranium:2.3%
5	comput:1.5%	health:1.4%	compani:1.2%	aid:1.2%	care:1.1%

Figure A.5. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Hitech data set

Statistics of Single-link Agglomerative Hierarchical Algorithm

(K=6) for Hitech Dataset

Number of: Documents = 1530 Topics = 6 Terms = 10919

Overall Similarity = 0.024555 Overall F-measure = 0.285581

Overall Entropy = 2.40469 Overall Purity = 0.266013

Topics: 0:computer 1:electronics 2:medical 3:health 4:research

5:technology

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	1525	0.021	2.4	0.26	322	76	284	402	318	123
1	1	1	0	1	0	0	0	0	1	0
2	1	1	0	1	0	0	1	0	0	0
3	1	1	0	1	0	0	0	0	1	0
4	1	1	0	1	0	0	1	0	0	0
5	1	1	0	1	0	0	0	0	0	1

CID Most Descriptive 5 Features

0	comput: 2%	compani:1.2%	health:1.1%	aid:0.91%	care:0.87%
1	goat:58%	fugit:4.2%	transmitt:3.5%	radioact:2.6%	poss: 2.3%
2	newman:76%	stephani:3.7%	lupu:2.8%	disabl:1.2%	didn:0.83%
3	spider:87%	estrada:3.7%	insect:0.93%	guest:0.52%	woman:0.51%
4	booth:17%	assassin: 15%	lincoln: 10%	grandfath: 5%	samuel: 4.3%
5	racket:73%	tenni:3.9%	graf:1.5%	wide:1.5%	game: 1.4%

Figure A.6. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Hitech data set

Statistics of k-means (K=6) for LA1 Dataset										
Number of: Documents = 2134 Topics = 6 Terms = 14363										
Overall Similarity = 0.0357047 Overall F-measure = 0.702847										
Overall Entropy = 1.02057 Overall Purity = 0.789128										
Topics: 0:financial 1:foreign 2:national 3:metro 4:sports										
5:entertainment										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	507	0.028	1.5	0.66	50	8	94	337	6	12
1	302	0.034	1.4	0.69	7	34	43	208	7	3
2	246	0.037	1.5	0.69	17	170	30	24	3	2
3	276	0.035	1.1	0.79	4	4	3	30	16	219
4	467	0.043	0.15	0.98	1	1	2	4	459	0
5	336	0.038	0.77	0.87	291	9	10	25	1	0
CID Most Descriptive 5 Features										
0	counti:	3.4%	bush:	2.4%	citi:	1.9%	court:	1.5%	reagan:	1.1%
1	police:	13%	fire:	2.8%	diego:	2.3%	arrest:	2.1%	car:	1.9%
2	soviet:	6.8%	israel:	2.5%	libya:	2.1%	airlin:	2%	union:	1.8%
3	art:	7.3%	aleen:	5.4%	macmin:	5.4%	music:	4.2%	film:	2%
4	game:	8.9%	scor:	5.7%	team:	3.3%	coach:	2.1%	plai:	2.1%
5	compani:	3.9%	million:	3.1%	bank:	2.7%	stock:	2.6%	market:	2.3%

Figure A.7. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by *k*-means for LA1 data set

Statistics of Bisecting k-means (K=6) for LA1 Dataset										
Number of: Documents = 2134 Topics = 6 Terms = 14363										
Overall Similarity = 0.0335053 Overall F-measure = 0.685949										
Overall Entropy = 1.08306 Overall Purity = 0.761481										
Topics: 0:financial 1:foreign 2:national 3:metro 4:sports										
5:entertainment										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	338	0.028	1.6	0.63	23	212	58	35	1	9
1	479	0.042	0.16	0.98	0	0	3	5	470	1
2	422	0.033	1.3	0.73	309	3	17	30	1	62
3	310	0.029	1.8	0.51	7	4	50	82	9	158
4	466	0.03	0.99	0.83	25	5	36	386	8	6
5	119	0.045	1.2	0.76	6	2	18	90	3	0
CID Most Descriptive 5 Features										
0	soviet:	5.3%	datelin:	2.7%	israel:	1.8%	union:	1.8%	airlin:	1.5%
1	game:	8.8%	scor:	5.6%	team:	3.3%	bowl:	2.1%	coach:	2.1%
2	compani:	3.4%	million:	2.5%	aleen:	2.3%	macmin:	2.3%	stock:	2%
3	bush:	5.7%	music:	2.6%	art:	2.3%	reagan:	2.3%	white:	1.2%
4	polic:	4.9%	counti:	4.8%	citi:	2.4%	court:	1.8%	diego:	1.8%
5	health:	4.6%	care:	2.5%	counti:	2%	medic:	1.9%	hospit:	1.7%

Figure A.8. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for LA1 data set

Statistics of Divisive Hierarchical Algorithm (K=6) for LA1										
Dataset										
Number of: Documents = 2134 Topics = 6 Terms = 14363										
Overall Similarity = 0.0247786 Overall F-measure = 0.296538										
Overall Entropy = 2.31287 Overall Purity = 0.329897										
Topics: 0:financial 1:foreign 2:national 3:metro 4:sports										
5:entertainment										
CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	1	1	0	1	1	0	0	0	0	0
1	263	0.037	1.5	0.52	17	0	2	96	137	11
2	44	0.068	1.5	0.61	7	0	2	27	8	0
3	1758	0.016	2.5	0.28	311	208	175	492	347	225
4	16	0.32	0.87	0.81	0	13	1	2	0	0
5	52	0.11	1.4	0.65	34	5	2	11	0	0
CID Most Descriptive 5 Features										
0	lem:34%		amp: 14%	receptor: 8.5%	network: 8.4%				mci:6.9%	
1	scor:9.3%		game:4.4%	orang: 3.3%	counti:3.1%	fullerton:2.1%				
2	health:8.7%		insur:6.3%	winfield: 4.2%	hous:2.9%	counti:2.9%				
3	polic:1.1%		san:0.8%	game:0.79%	diego:0.68%	nation:0.6%				
4	israel: 30%	palestinian:15%		israe: 7.3%	arab:5.2%	plo:4.5%				
5	bank: 32%		loan:14%	sav: 4.6%	mortgag:2.7%	billion:2.1%				

Figure A.9. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for LA1 data set

Statistics of Average-link Agglomerative Hierarchical Algorithm
(K=6) for LA1 Dataset

Number of: Documents = 2134 Topics = 6 Terms = 14363

Overall Similarity = 0.0256362 Overall F-measure = 0.438493

Overall Entropy = 1.90898 Overall Purity = 0.487816

Topics: 0:financial 1:foreign 2:national 3:metro 4:sports
5:entertainment

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	1610	0.017	2.3	0.36	357	213	178	582	50	230
1	3	0.37	0.92	0.67	0	0	0	2	0	1
2	23	0.11	2.2	0.35	4	8	0	5	2	4
3	10	0.34	0.47	0.9	0	1	0	9	0	0
4	487	0.039	0.62	0.9	9	4	4	30	439	1
5	1	1	0	1	0	0	0	0	1	0

CID Most Descriptive 5 Features

0	counti:1.4%	polic:1.3%	brief: 0.84%	compani: 0.76%	million: 0.72%
1	courag:14%	unless: 6%	lewi: 5.4%	resum: 5.3%	randi: 5.2%
2	mexico:19%	quake:12%	earthquak: 7.9%	beer: 4%	peru: 3.4%
3	lotteri:43%	lotto:11%	jackpot: 7.2%	bonu: 5.3%	woodard: 2.6%
4	game: 8%	scor:5.5%	team: 3.1%	plai: 2%	bowl: 2%
5	cushion:43%	murrai:12%	agoura: 6.1%	spoil: 4.3%	blank: 4.1%

Figure A.10. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for LA1 data set

Statistics of Complete-link Agglomerative Hierarchical Algorithm
(K=6) for LA1 Dataset

Number of: Documents = 2134 Topics = 6 Terms = 14363

Overall Similarity = 0.0216313 Overall F-measure = 0.283927

Overall Entropy = 2.36051 Overall Purity = 0.324742

Topics: 0:financial 1:foreign 2:national 3:metro 4:sports
5:entertainment

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	38	0.15	0.4	0.92	35	0	3	0	0	0
1	9	0.19	1.9	0.33	3	0	0	1	2	3
2	10	0.17	0.47	0.9	9	0	0	1	0	0
3	35	0.082	1.1	0.74	26	0	3	6	0	0
4	9	0.18	1.9	0.33	1	3	2	3	0	0
5	2033	0.016	2.4	0.3	296	223	174	617	490	233

CID Most Descriptive 5 Features

0	bank: 11%	dollar: 4.6%	market: 4.3%	trad: 3.9%	stock:3.5%
1	bill: 10%	oiler: 8.3%	discount: 8%	seinfeld: 6.9%	rate:4.2%
2	busi: 4%	list: 3.6%	investor: 3.3%	firm: 3.1%	compani:2.9%
3	comput: 15%	micro: 4.1%	ibm: 3.1%	digit: 3.1%	store: 3%
4	space: 10%	rockwell: 10%	shuttle: 9.4%	force: 6.2%	air:6.2%
5	game:1.4%	counti: 1.1%	san:0.93%	police: 0.9%	scor:0.81%

Figure A.11. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for LA1 data set

Statistics of Single-link Agglomerative Hierarchical Algorithm
(K=6) for LA1 Dataset

Number of: Documents = 2134 Topics = 6 Terms = 14363

Overall Similarity = 0.0188719 Overall F-measure = 0.285809

Overall Entropy = 2.43514 Overall Purity = 0.295689

Topics: 0:financial 1:foreign 2:national 3:metro 4:sports
5:entertainment

CID	Size	Sim	Entropy	Purity	0	1	2	3	4	5
0	2	0.67	0	1	0	0	0	0	0	2
1	1	1	0	1	0	0	0	1	0	0
2	2	0.94	0	1	0	0	0	2	0	0
3	2127	0.016	2.4	0.29	370	226	182	624	492	233
4	1	1	0	1	0	0	0	0	0	1
5	1	1	0	1	0	0	0	1	0	0

CID Most Descriptive 5 Features

0	seinfeld:32%	christon:7.3%	comedian: 6%	relev: 3.9%	jerri: 2.3%
1	rader:78%	murder:2.3%	harvei: 1.2%	prove: 1%	preliminari:0.84%
2	nash:35%	thorson: 28%	dile: 9.1%	holme: 5.9%	liberac: 3%
3	game:1.3%	counti:1.1%	san:0.87%	polic:0.85%	scor:0.74%
4	cotter:62%	anim:1.5%	stubborn: 1.5%	pet:0.99%	spe:0.92%
5	blaisdell:56%	postwar:6.1%	berkeleyi: 4.6%	truman: 2.3%	recoveri: 1.9%

Figure A.12. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for LA1 data set

Statistics of K-means Algorithm (K=10) for Reuters-21578 Dataset									
Number of: Documents = 9603 Topics = 90 Terms = 12772									
Overall Similarity = 0.103544 Overall F-measure = 0.159476									
Overall Entropy = 2.3644 Overall Purity = 0.514943									
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT			
0	506	0.057	3.7	0.65	crude	328			
1	1464	0.19	0.03	1	earn	1461			
2	753	0.38	4	0.2	acq	150			
3	797	0.084	0.4	0.019	earn	15			
4	642	0.13	0.017	1	earn	641			
5	552	0.1	2.3	0.43	interest	237			
6	1058	0.059	5	0.27	trade	287			
7	2210	0.02	1.6	0.62	acq	1365			
8	859	0.047	2.8	0.1	coffee	88			
9	762	0.055	6.6	0.49	grain	373			
CID	Most Descriptive 5 Features								
0	oil:34%	barrel:5.9%	crude:	4%	opec:	2.6%	price:	2.4%	
1	loss:23%	net:18%	shr:	16%	mln:	8%	rev:	7.9%	
2	blah:97%	fed:0.21%	billion:	0.15%	dlr:	0.1%	bank:0.076%		
3	bond:13%	issu:13%	debentur:	4.5%	manag:	4.3%	coupon:	3.1%	
4	div:17%	qtly:16%	record:	14%	dividend:	9.5%	prior:	9.1%	
5	stg:20%	bank:11%	rate:	6.9%	bill:	6.3%	monei:	5.6%	
6	billion:9.1%	trade:5.1%	januari:	5%	februari:	4.8%	dollar:	3.3%	
7	share:8.2%	compani:3.8%	dlr:	2.7%	offer:	2.6%	stock:	2.3%	
8	bank:12%	debt:11%	loan:	4.9%	brazil:	4.1%	coffe:	2.3%	
9	tonn:28%	wheat:7.5%	sugar:	4.5%	export:	3.6%	grain:	3.1%	

Figure A.13. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by *k*-means for Reuters-21578 data set

Statistics of Bisecting K-means Algorithm (K=10) for Reuters-21578 Dataset						
Number of: Documents = 9603 Topics = 90 Terms = 12772						
Overall Similarity = 0.103544 Overall F-measure = 0.159476						
Overall Entropy = 2.3644 Overall Purity = 0.514943						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	506	0.057	3.7	0.65	crude	328
1	1464	0.19	0.03	1	earn	1461
2	753	0.38	4	0.2	acq	150
3	797	0.084	0.4	0.019	earn	15
4	642	0.13	0.017	1	earn	641
5	552	0.1	2.3	0.43	interest	237
6	1058	0.059	5	0.27	trade	287
7	2210	0.02	1.6	0.62	acq	1365
8	859	0.047	2.8	0.1	coffee	88
9	762	0.055	6.6	0.49	grain	373
CID Most Descriptive 5 Features						
0	oil:34%	barrel:5.9%	crude:4%	opec:2.6%	price:2.4%	
1	loss:23%	net:18%	shr:16%	mln:8%	rev: 7.9%	
2	blah:97%	fed:0.21%	billion:0.15%	dlr:0.1%	bank:0.076%	
3	bond:13%	issu:13%	debentur:4.5%	manag:4.3%	coupon:3.1%	
4	div:17%	qtly:16%	record:14%	dividend:9.5%	prior:9.1%	
5	stg:20%	bank:11%	rate:6.9%	bill:6.3%	monei:5.6%	
6	billion:9.1%	trade:5.1%	januari:5%	februari:4.8%	dollar:3.3%	
7	share:8.2%	compani:3.8%	dlr:2.7%	offer:2.6%	stock:2.3%	
8	bank:12%	debt:11%	loan:4.9%	brazil:4.1%	coffe:2.3%	
9	tonn:28%	wheat:7.5%	sugar:4.5%	export:3.6%	grain:3.1%	

Figure A.14. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Reuters-21578 data set

Statistics of Divisive Hierarchical Algorithm (K=10) for
Reuters-21578 Dataset

Number of: Documents = 9603 Topics = 90 Terms = 12772

Overall Similarity = 0.0704095 Overall F-measure = 0.0987711

Overall Entropy = 3.13715 Overall Purity = 0.442466

CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	1457	0.026	6.1	0.16	grain	237
1	1366	0.14	2.8	0.55	earn	750
2	2620	0.034	5.2	0.16	money-fx	429
3	1555	0.033	1.2	0.4	acq	621
4	1667	0.16	0.18	0.97	earn	1621
5	59	0.058	2.4	0.58	ship	34
6	428	0.024	1.2	0.65	acq	280
7	309	0.033	0.91	0.69	acq	214
8	108	0.094	4	0.46	gold	50
9	34	0.12	3	0.38	ship	13

CID Most Descriptive 5 Features

0	tonn: 7.8%	oil: 3.4%	export: 2.7%	wheat: 2.2%	price: 1.7%
1	blah: 77%	qtly: 3.2%	div: 3.1%	record: 2.7%	dividend: 1.8%
2	bank: 8.7%	billion: 5.2%	rate: 2.8%	trade: 1.9%	stg: 1.8%
3	issu: 5.6%	offer: 5.6%	bond: 4.2%	share: 3.8%	debentur: 3.4%
4	loss: 23%	net: 18%	shr: 15%	mln: 8.5%	rev: 7.4%
5	ship: 17%	portland: 7.7%	port: 7.3%	load: 4.2%	coastal: 3.9%
6	sale: 4%	complet: 3.8%	unit: 3.2%	compani: 3%	acquisit: 2.9%
7	share: 11%	usair: 5.3%	merger: 4.8%	compani: 3.8%	earn: 3.1%
8	mine: 23%	gold: 19%	ounc: 7.2%	grade: 4.4%	ton: 3.9%
9	strike: 23%	seamen: 11%	union: 6.4%	port: 2.7%	marin: 2.6%

Figure A.15. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Reuters-21578 data set

Statistics of Average-link Algorithm (K=10) for Reuters-21578						
Number of: Documents = 9603 Topics = 90 Terms = 12772						
Overall Similarity = 0.0268052 Overall F-measure = 0.0567252						
Overall Entropy = 3.95038 Overall Purity = 0.321879						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	9085	0.024	4	0.31	earn	2839
1	222	0.064	3.5	0.36	gold	79
2	2	0.73	0	1	acq	2
3	4	0.59	0	1	acq	4
4	16	0.12	0.72	0.56	acq	9
5	216	0.047	2.7	0.63	ship	135
6	33	0.1	2.1	0.3	acq	10
7	20	0.12	1.5	0.45	acq	9
8	3	0.58	0.39	0.67	acq	2
9	2	0.64	0	1	acq	2
CID Most Descriptive 5 Features						
0	blah: 10%	loss: 5.7%	mln: 5.2%	net: 4.5%	shr: 3.7%	
1	gold: 17%	mine: 16%	ounc: 8.5%	ton: 7.7%	copper: 6.6%	
2	gerber: 65%	gst: 8%	cwt: 6.8%	buyout: 3%	fremont: 2.3%	
3	gate: 46%	learjet: 22%	interconnect:14%	norri: 5.6%	berri: 3.9%	
4	print: 12%	magazin: 9.4%	newspap: 8.3%	southam: 5.7%	printer: 5.3%	
5	ship: 12%	strike: 7.7%	gulf: 5.3%	port: 4.7%	seamen: 3.2%	
6	coastal:17%	court:8.3%	transamerican:6.6%	suffield:6.3%	bankruptci: 4.2%	
7	arco: 7.8%	ciba: 7.1%	cell: 5.8%	sandoz: 5.6%	quest: 4.4%	
8	imatron: 45%	mitsui: 29%	benigno: 3.8%	kidnap: 3.5%	businessman:2.2%	
9	saatchi: 55%	jwt: 25%	cleveland: 4.1%	advertis: 2.5%	ted: 2.2%	

Figure A.16. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Reuters-21578 data set

Statistics of Complete-link Algorithm (K=10) for Reuters-21578									
Number of: Documents = 9603 Topics = 90 Terms = 12772									
Overall Similarity = 0.052231 Overall F-measure = 0.0812158									
Overall Entropy = 3.80474 Overall Purity = 0.320837									
CID Size		Sim Entropy		Purity	MPT	# of Docs from MPT			
0	90	0.23	0.88	0.92	coffee	83			
1	37	0.16	3.8	0.97	grain	36			
2	66	0.24	2.9	1	grain	66			
3	21	0.11	3.9	0.43	crude	9			
4	6	0.53	1.3	0.67	acq	4			
5	10	0.3	0.88	0.7	acq	7			
6	345	0.23	0	1	earn	345			
7	28	0.093	1.2	0.5	acq	14			
8	478	0.4	0	1	earn	478			
9	8522	0.021	4.2	0.24	earn	2039			
CID	Most Descriptive 5 Features								
0	coffe:46%	quota: 8.1%	ico: 7.6%	bag: 4.2%	export: 3.9%				
1	grain:44%	load: 6.9%	portland: 6.1%	ship: 2.9%	gulf: 2.8%				
2	wheat:49%	tonn: 22%	export: 1.9%	soviet: 1.5%	depart: 1.1%				
3	usx:9.1%	tax: 7.1%	field: 4.5%	roderick: 4%	hog:3.8%				
4	champion:73%	claremont: 7.6%	echlin: 2.1%	blah: 1.1%	board: 1%				
5	comput:55%	microfilm: 8.9%	wavehil: 4.3%	comi: 2.8%	person:2.5%				
6	net 29%	shr: 24%	mln: 12%	rev: 10%	qtr: 5.3%				
7	seali:7.8%	ohio: 5.9%	triton: 5.2%	mln: 3.9%	neoax:3.9%				
8	loss:59%	profit: 14%	shr: 6.5%	net: 5.5%	rev: 4.3%				
9	blah:14%	mln: 3.2%	bank: 2.8%	dlr: 2.2%	billion: 2.2%				

Figure A.17. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Reuters-21578 data set

Statistics of Single-link Algorithm (K=10) for Reuters-21578						
Number of: Documents = 9603 Topics = 90 Terms = 12772						
Overall Similarity = 0.0236582 Overall F-measure = 0.0286495						
Overall Entropy = 4.13923 Overall Purity = 0.300427						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	1	1	0	1	acq	1
1	1	1	0	1	acq	1
2	1	1	0	1	acq	1
3	1	1	0	1	acq	1
4	1	1	0	1	acq	1
5	1	1	0	1	acq	1
6	1	1	0	1	pet-chem	1
7	1	1	0	1	acq	1
8	9594	0.023	4.1	0.3	earn	2876
9	1	1	0	1	earn	1
CID Most Descriptive 5 Features						
0	amsouth:65% tuskaloosa: 10% affili: 10% aso: 2.6% approv: 2.2%					
1	mcfarland:88% approv: 1.3% santa: 1.2% spring: 1.1% exchang:0.89%					
2	avia:66% reebok: 14% stockhold: 3.6% portland: 1.5% complaint:1.4%					
3	dumez:65% westburn: 23% unicorp: 1.7% share: 1% common:0.94%					
4	lazer:42% fidelcor: 38% ficr: 2.4% acquir: 2% bkne: 1.8%					
5	seton:82% sel: 2.3% newark: 1.4% member: 1.3% chairman: 1%					
6	protocol:15% aerosol: 9.1% layer: 7.7% scientist: 6.7% earth:6.4%					
7	frick:44% frigid: 25% coil: 8.9% refriger: 7.4% compressor:2.5%					
8	blah:10% loss: 5.5% mln: 5.1% net: 4.3% shr: 3.6%					
9	ncr:67% proced: 2.2% major: 1.6% strongest: 1.6% deliver:1.4%					

Figure A.18. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Reuters-21578 data set

Statistics of K-Means Algorithm (K=10) for Wap Dataset						
Number of: Documents = 1560 Topics = 20 Terms = 8061						
Overall Similarity = 0.0618784 Overall F-measure = 0.16397						
Overall Entropy = 1.8379 Overall Purity = 0.570513						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	191	0.068	2.9	0.27	business	52
1	97	0.075	0.56	0.92	sports	89
2	461	0.035	2.5	0.39	film	180
3	96	0.1	1.6	0.64	politics	61
4	193	0.049	2.5	0.43	music	83
5	71	0.12	0	1	health	71
6	140	0.053	2.6	0.49	people	69
7	24	0.18	0	1	health	24
8	245	0.054	0.068	0.99	health	243
9	42	0.17	1.8	0.43	culture	18
CID Most Descriptive 5 Features						
0	walter:8.1%	gail: 4.3%		militari: 3.2%	liquor: 2.8%	loan: 2.6%
1	diaz:11%	neurosurgeri: 2.1%		mira: 1.7%	agnieszka: 1.7%	vehicl:1.6%
2	elit:5.5%	craig: 1.8%		degen: 1.4%	uncompl: 1.3%	chiffon: 1.3%
3	thrive:11%	virsa: 7%		construct: 5.6%	trevor: 4.2%	duplic: 3.3%
4	academi:5.1%	dar: 3.4%		slovakia: 2.8%	cathol: 2.4%	arbitr: 1.9%
5	hornbi:26%	marshal: 4.7%		incen: 3.4%	bess: 3%	sheila: 2.3%
6	evidenc:12%	neonat: 3.4%		widescreen: 2.9%	statem: 2.5%	backward:2.1%
7	kendrick:10%	lesson: 8.6%		fetch: 5.8%	rees: 5.6%	suspicion: 4.4%
8	grammi:4.6%	nato: 4.1%		confidentia: 3.6%	jude: 3.4%	toll: 2.7%
9	plant: 11%	restag: 9.4%		sofa: 6.7%	edinburgh: 5.8%	lousi: 5.1%

Figure A.19. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by k -means for Wap data set

Statistics of Bisecting K-Means Algorithm (K=10) for Wap Dataset						
Number of: Documents = 1560 Topics = 20 Terms = 8061						
Overall Similarity = 0.0599456 Overall F-measure = 0.154444						
Overall Entropy = 1.95732 Overall Purity = 0.548077						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	101	0.087	2.2	0.59	politics	60
1	43	0.14	1.1	0.56	health	24
2	264	0.038	2.6	0.36	film	96
3	301	0.04	2.7	0.4	people	119
4	144	0.061	1.8	0.47	television	68
5	106	0.075	1.7	0.74	film	78
6	225	0.058	0.14	0.98	health	221
7	242	0.058	3.1	0.3	business	72
8	96	0.095	0	1	health	96
9	38	0.098	1.9	0.55	sports	21
CID Most Descriptive 5 Features						
0	thrive: 11%	virsa: 6.3%	construct: 6.1%	trevor: 4.5%	duplic:3.8%	
1	plant: 12%	restag: 12%	sofa: 7.7%	daunt: 6.4%	refrain: 5.7%	
2	elit: 4.2%	degen: 1.6%	berendt: 1.5%	veto: 1.5%	flaherti: 1.1%	
3	evidenc: 4.2%	academi: 2.3%	dar: 1.9%	explan: 1.5%	cathol: 1.3%	
4	ell: 6.8%	uncompl: 4.5%	conver: 3.6%	diaz: 3.1%	melod: 2.9%	
5	craig: 11%	elit: 9.6%	roddi: 5.9%	carolyn: 4.6%	donatella: 3.8%	
6	grammi: 5%	confidentia: 4.1%	nato: 3.6%	sheila: 3.2%	toll: 2.9%	
7	walter: 5.1%	fisher: 4.7%	gail: 3.9%	militari: 2.1%	item: 2%	
8	hornbi: 19%	marshal: 4.2%	spa: 3.4%	incen: 3.2%	lesson: 2.8%	
9	mira: 8.4%	letter: 5%	gardner: 4%	parad: 2.5%	bang: 2.4%	

Figure A.20. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by bisecting k -means for Wap data set

Statistics of Divisive Hierarchical Algorithm (K=10) for WAP						
Number of: Documents = 1560 Topics = 20 Terms = 8061						
Overall Similarity = 0.0396669 Overall F-measure = 0.209029						
Overall Entropy = 2.80873 Overall Purity = 0.35641						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	9	0.24	2.1	0.44	art	4
1	43	0.073	3	0.28	music	12
2	796	0.027	3.2	0.19	film	151
3	69	0.11	1.9	0.52	technology	36
4	618	0.031	2.4	0.55	health	338
5	10	0.18	1.3	0.6	culture	6
6	2	0.97	0	1	film	2
7	2	0.68	0	1	people	2
8	2	0.54	0	1	cable	2
9	9	0.26	2.2	0.33	film	3
CID Most Descriptive 5 Features						
0	recollect:11% ethan:6.1% georgetown: 4.2% chelsea:3.8% kafelnikov: 2.3%					
1	incid:5.5% libel: 5.4% slovakia: 4.9% bang: 4.6% rowan: 3%					
2	elit:3.1% degen: 1.1% evidenc: 1% restag: 0.92% ell:0.84%					
3	walter:14% militari: 5.9% loan: 4.9% shuttl: 3% gail:2.7%					
4	nato: 2% confidentia: 1.9% hornbi: 1.9% grammi: 1.7% marshal:1.7%					
5	portabl:15% omar: 8.2% subsid: 6.1% buena: 2.9% budget: 2.6%					
6	psychiatr:46% jesper: 16% marijuana: 5% elit: 2% dad: 1.6%					
7	lousi:9.6% mckinlei: 5.9% nathan: 5.8% feloni: 5.5% nchant:4.9%					
8	decor:9.6% prostat: 7.2% distort: 5.8% moonlight: 3.8% meer:2.7%					
9	specter:34% exacerb: 9.6% paxson: 3.6% vaccin: 2% steam: 2%					

Figure A.21. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by divisive hierarchical clustering for Wap data set

Statistics of Average-link Algorithm (K=10) for Wap Dataset

Number of: Documents = 1560 Topics = 20 Terms = 8061

Overall Similarity = 0.0507655 Overall F-measure = 0.122892

Overall Entropy = 2.3274 Overall Purity = 0.447436

CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	699	0.031	3.1	0.24	film	168
1	12	0.22	2	0.5	culture	6
2	8	0.25	1.6	0.38	film	3
3	36	0.19	2.3	0.5	culture	18
4	5	0.25	0.72	0.8	review	4
5	12	0.18	2	0.5	film	6
6	326	0.054	0.2	0.98	health	319
7	9	0.3	2.1	0.44	multimedia	4
8	349	0.043	3.5	0.21	business	73
9	104	0.076	0.43	0.93	sports	97

CID Most Descriptive 5 Features

0	elit: 4.1%	evidenc: 1.3%	craig: 1.1%	moonlight: 1.1%	degen: 1.1%
1	bang: 15%	game: 14%	libel: 10%	incid: 6.3%	wallac: 5.3%
2	mccourt:14%	masterpiec: 11%	tycoon: 6.6%	macdonald: 6.3%	transport: 6.1%
3	restag: 16%	plant: 13%	sofa: 8.3%	refrain: 5.9%	argum: 5.6%
4	headquart: 8.8%	featur: 7.1%	elisabeth: 5.8%	diva: 5.5%	braill: 4.2%
5	fbi: 8.4%	andrea: 8.2%	gather: 6.7%	danni: 6.6%	qualiti: 4.7%
6	hornbi: 3.9%	confidentia: 3.7%	nato: 3.6%	grammi: 3.4%	marshal: 3.4%
7	orphan: 35%	reform: 15%	preterm: 12%	coffin: 4.7%	trash: 4.6%
8	walter: 3.7%	fisher: 2.9%	gail: 2.3%	thrive: 2%	militari: 1.5%
9	diaz: 9.4%	edinburgh: 3.3%	snub: 2.2%	neurosurgeri: 1.9%	agnieszka: 1.6%

Figure A.22. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by average-link for Wap data set

Statistics of Complete-link Algorithm (K=10) for Wap Dataset						
Number of: Documents = 1560 Topics = 20 Terms = 8061						
Overall Similarity = 0.04484 Overall F-measure = 0.111988						
Overall Entropy = 2.97154 Overall Purity = 0.352564						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	24	0.099	2.2	0.38	music	9
1	15	0.18	2.3	0.4	culture	6
2	13	0.2	0.39	0.92	music	12
3	237	0.059	0.039	1	health	236
4	23	0.098	2.9	0.26	television	6
5	61	0.12	1.4	0.77	television	47
6	10	0.22	2	0.4	people	4
7	20	0.12	2.7	0.4	business	8
8	51	0.17	1.7	0.71	politics	36
9	1106	0.023	3.8	0.17	film	186
CID Most Descriptive 5 Features						
0	kurfuerstendam:6% vincent: 5.6% georgetown:3.1% dairi: 3.1% academi:2.1%					
1	game: 17% bang: 12% libel: 7.3% incid: 4.2% wallac: 4.2%					
2	viral: 13% academi: 9.8% simplifi: 6.6% screwbal: 6.3% robin: 5.8%					
3	hornbi: 4.9% nato: 4.2% grammi: 3.8% marshal: 3.6% confidentia: 3.5%					
4	amanda: 7.8% haven: 6.6% yugoslav: 5.2% below: 3.4% slovakia: 2.3%					
5	conver: 9.2% melod: 7.7% uncompl: 7.2% unopen: 6.3% pierr: 6%					
6	specter: 21% exacerb: 6.7% chemic: 5.3% modern: 4.1% breakthrough:3.5%					
7	sporad: 8.8% intern: 6.3% kick: 5.8% henman: 5.8% ordinari: 3.8%					
8	thrive: 15% construct: 9.4% virsa: 8.4% trevor: 7.8% duplic: 5.6%					
9	elit: 2.2% carolyn:0.85% walter:0.79% gail: 0.7% degen:0.62%					

Figure A.23. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by complete-link for Wap data set

Statistics of Single-link Algorithm (K=10) for Wap Dataset						
Number of: Documents = 1560 Topics = 20 Terms = 8061						
Overall Similarity = 0.0264066 Overall F-measure = 0.0324178						
Overall Entropy = 3.6985 Overall Purity = 0.224359						
CID	Size	Sim	Entropy	Purity	MPT	# of Docs from MPT
0	1551	0.021	3.7	0.22	health	341
1	1	1	0	1	review	1
2	1	1	0	1	review	1
3	1	1	0	1	review	1
4	1	1	0	1	review	1
5	1	1	0	1	people	1
6	1	1	0	1	review	1
7	1	1	0	1	review	1
8	1	1	0	1	review	1
9	1	1	0	1	review	1
CID Most Descriptive 5 Features						
0	elit: 1.5%	carolyn:0.65%	melod:0.63%	restag:0.58%	nato:0.53%	
1	frazier: 12%	agent: 8.3%	tandem: 5.2%	octob: 4.1%	costa: 2.9%	
2	highwai: 16%	brown: 9.6%	encompass: 4.8%	undermin: 2.8%	produc: 2.7%	
3	rental: 8.2%	stupid: 3.6%	dog: 3.6%	appreci: 3.6%	jovi: 3.6%	
4	barri: 8.8%	westport: 6.1%	sleep: 4.2%	assum: 3.2%	enqvist: 3.2%	
5	rule: 21%	held: 16%	elig: 14%	judith: 6%	exclu: 3.2%	
6	talli: 8.4%	reiner: 4.1%	physiolog: 3.5%	entitl: 3.5%	rusedski: 3.1%	
7	furiou: 14%	petrikin: 8.8%	morton: 7.2%	lar: 4.2%	rough: 2.6%	
8	shutout: 22%	wareh: 18%	buck: 5.5%	pen: 2.9%	stumbl: 2.1%	
9	basqu: 11%	blunt: 6.7%	light: 5.3%	illumin: 3.9%	garri: 3.7%	

Figure A.24. The performance, cluster-class distribution, and most descriptive 5 features of each cluster, of the clustering solution obtained by single-link for Wap data set

REFERENCES

1. Sebastiani, F., “Machine Learning in Automated Text Categorization”, *ACM Computing Surveys*, Vol. 34, No. 1, pp. 1–47, 2002.
2. Masand, B., G. Linoff, and D. Waltz, “Classifying News Stories Using Memory Based Reasoning”, *15th Ann Int ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 59–64, 1992.
3. Yang, Y., “An Evaluation of Statistical Approaches to Text Categorization”, *Journal of Information Retrieval*, Vol. 1, No. 1–2, pp. 69–90, 1999.
4. Yang, Y. and J. P. Pedersen, “A Comparative Study on Feature Selection in Text Categorization”, *The Fourteenth International Conference on Machine Learning*, pp. 412–420, 1997.
5. Joachims, T., “Text Categorization with Support Vector Machines: Learning with Many Relevant Features”, *European Conference on Machine Learning (ECML)*, 1998.
6. Ng, H. T., W. B. Goh, and K. L. Low, “Feature Selection, Perceptron Learning, and a Usability Case Study for Text Categorization”, *20th Ann Int ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 67–73, 1997.
7. Wiener, E., J. O. Pedersen, and A. S. Weigend, “A Neural Network Approach to Topic Spotting”, *Proceedings of the Fourth Annual Symposium on Document Analysis and Information Retrieval*, 1995.
8. Yang, Y. and X. Liu, “A Re-examination of Text Categorization Methods”, *Proceedings of SIGIR-99, 22nd ACM International Conference on Research and Development in Information Retrieval*, Berkeley, US, 1999.

9. McCallum, A. and K. Nigam, “A Comparison of Event Models for Naive Bayes Text Classification”, *AAAI-98 Workshop on Learning for Text Categorization*, 1998.
10. Koller, D. and M. Sahami, “Hierarchically Classifying Documents Using Very Few Words”, *Proceedings of 14th International Conference on Machine Learning*, pp. 170–178, Nashville, US, 1997.
11. Larkey, L. S., “A Patent Search and Classification System”, *Proceedings of 4th ACM Conference on Digital Libraries*, pp. 179–187, Berkely, US, 1999.
12. Ozgur, L., *Adaptive Anti-Spam Filtering Based on Turkish Morphological Analysis, Artificial Neural Networks and Bayes Filtering*, Master’s Thesis, Bogazici University, Turkey, 2003.
13. Sahami, M., S. Dumais, D. Heckerman, and E. Horvitz, “A Bayesian Approach to Filtering Junk E-Mail”, Sahami, M., editor, *Proceedings of AAAI-98 Workshop on Learning for Text Categorization*, 1998.
14. Dumais, S. T. and H. Chen, “Hierarchical Classification of Web Content”, *Proceedings of SIGIR-00, 23rd ACM International Conference on Research and Development in Information Retrieval*, pp. 256–263, Athens, GR, 2000.
15. Zamir, O., O. Etzioni, O. Madani, and R. M. Karp, “Fast and Intuitive Clustering of Web Documents”, *KDD’97*, pp. 287–290, 1997.
16. Jain, A. K., M. N. Murty, and P. J. Flynn, “Data Clustering: A Review”, *ACM Computing Surveys*, Vol. 31, No. 3, pp. 264–323, September 1999.
17. Zhao, Y. and G. Karypis, “Evaluation of Hierarchical Clustering Algorithms for Document Datasets”, *Proceedings of CIKM*, 2002.
18. Steinbach, M., G. Karypis, and V. Kumar, “A Comparison of Document Clustering Techniques”, *KDD Workshop on Text Mining*, 1999.

19. Salton, G., C. Yang, and A. Wong, “A Vector-Space Model for Automatic Indexing”, *Communications of the ACM*, Vol. 18, No. 11, pp. 613–620, 1975.
20. Cohen, W. W. and Y. Singer, “Context-Sensitive Learning Methods for Text Categorization”, *Proceedings of the 19th Annual ACM SIGIR Conference*, 1996.
21. Fuernkranz, J., T. Mitchell, and E. Riloff, “A Case Study in Using Linguistic Phrases for Text Categorization on the WWW”, Sahami, M., editor, *In Learning for Text Categorization: Papers from the 1998 AAAI Workshop (Technical Report WS-98-05)*, 1998.
22. Dumais, S., T. Platt, D. Heckermann, and M. Sahami, “Inductive Learning Algorithms and Representations for Text Categorization”, *Proceedings of the Seventh International Conference on Information and Knowledge Management*, 1998.
23. Sahami, M., *Using Machine Learning To Improve Information Access*, Ph.D. Thesis, Stanford University, 1998.
24. Porter, M. F., “An Algorithm for Suffix Stripping”, *Program*, Vol. 14, pp. 130–137, 1980.
25. <ftp://ftp.cs.cornell.edu/pub/smart/>, 2004.
26. <http://www.tartarus.org/~martin/PorterStemmer/>, 2004.
27. Aas, L. and L. Eikvil, “Text Categorisation: A Survey”, 941, Norwegian Computing Center, June 1999.
28. Salton, G. and C. Buckley, “Term Weighting Approaches in Automatic Text Retrieval”, *Information Processing and Management*, Vol. 24, No. 5, pp. 513–523, 1988.
29. Manning, C. D. and H. Schutze, *Foundations of Statistical Natural Language Processing*, The MIT Press, Cambridge, Massachusetts, USA, 1999.

30. Cohen, *Empirical Methods in AI*, MIT Press, 1999.
31. Spitters, M., “Comparing Feature Sets for Learning Text Categorization”, *Proceedings of RIAO 2000*, April 2000.
32. Hartigan, J., *Clustering Algorithms*, John Wiley & Sons, New York, NY, 1975.
33. Berkhin, P., “Survey of Clustering Data Mining Techniques”, Research paper, Accrue Software, <http://www.accrue.com/products/researchpapers.html>, 2002.
34. Forgy, E., “Cluster Analysis of Multivariate Data: Efficiency Versus Interpretability of Classification”, *Biometrics*, Vol. 21, pp. 768–780, 1965.
35. Kaufman, L. and P. Rousseeuw, *Finding Groups in Data*, Wiley, New York, NY, 1990.
36. Mitchell, T. M., *Machine Learning*, McGraw Hill, 1997.
37. Joachims, T., “A Probabilistic Analysis of the Rocchio Algorithm with TFIDF for Text Categorization”, *ICML-97*, 1997.
38. Burges, C. J. C., “A Tutorial on Support Vector Machines for Pattern Recognition”, *Data Mining and Knowledge Discovery*, Vol. 2, No. 2, pp. 121–167, 1998.
39. Joachims, T., “Making Large-Scale SVM Learning Practical”, Report LS-8, Universität Dortmund, 1998.
40. Joachims, T., *Advances in Kernel Methods-Support Vector Learning*, chapter Making Large-Scale SVM Learning Practical, MIT-Press, 1999.
41. “TREC. Text Retrieval Conference”, <http://trec.nist.gov>, 1999.
42. Lewis, D. D., “Reuters-21578 Document Corpus V1.0”, <http://kdd.ics.uci.edu/databases/reuters21578/reuters21578.html>, 2004.

43. Han, E.-H. S., D. Boley, M. Gini, R. Gross, K. Hastings, G. Karypis, V. Kumar, B. Mobasher, and J. Moore, “Webace: A Web Agent for Document Categorization and Exploration”, *Proceedings of the 2nd International Conference on Autonomous Agents*, 1998.
44. Karypis, G., “Cluto 2.0 Clustering Toolkit”, <http://www.users.cs.umn.edu/~karypis/cluto>, 2004.
45. van Rijsbergen, C. J., *Information Retrieval*, Butterworths, London, 1979.