

On the Relative Value of Cross-Company and Within-Company Data for Defect Prediction

Burak Turhan · Tim Menzies · Ayse Bener · Justin Distefano

Received: date / Accepted: date

Abstract We propose a practical defect prediction approach for companies that do not track defect related data. Specifically, we investigate the applicability of cross-company (CC) data for building localized defect predictors using static code features.

Firstly, we analyze the conditions, where CC data can be used as is. These conditions turn out to be quite few. Then we apply principles of analogy-based learning (i.e. nearest neighbor (NN) filtering) to CC data, in order to fine tune these models for localization. We compare the performance of these models with that of defect predictors learned from within-company (WC) data. As expected, we observe that defect predictors learned from WC data outperform the ones learned from CC data. However, our analyses also yield defect predictors learned from NN-filtered CC data, with performance close to, but still not better than, WC data. Therefore, we perform a final analysis for determining the minimum number of local defect reports in order to learn WC defect predictors. We demonstrate in this paper that the minimum number of data samples required to build effective defect predictors can be quite small and can be collected quickly within a few months.

Hence, for companies with no local defect data, we recommend a two-phase approach that allows them to employ the defect prediction process instantaneously. In phase one, companies should use NN-filtered CC data to initiate the defect prediction process and simultaneously start collecting WC (local) data. Once enough WC data is collected (i.e. after a few months), organizations should switch to phase two and use predictors learned from WC data.

Burak Turhan
Department of Computer Engineering, Bogazici University, Istanbul, Turkey
E-mail: turhanb@boun.edu.tr

Tim Menzies
Lane Department of Computer Science and Electrical Engineering, West Virginia
E-mail: tim@menzies.us

Ayşe Bener
Department of Computer Engineering, Bogazici University, Istanbul, Turkey
E-mail: bener@boun.edu.tr

Justin Distefano
Lane Department of Computer Science and Electrical Engineering, West Virginia
E-mail: jdistefano@ismwv.com

KEYWORDS: defect prediction; learning; metrics (product metrics); cross-company; within-company; nearest-neighbor filtering

1 Introduction

Defect prediction studies usually focus on building models with available local data (i.e. within company predictors). To employ these models, a company should have a data repository, where project metrics and defect information from past projects are stored. However, few companies apply this practice. We suspect that a common reason for not using defect predictors in practice, is the lack of local data repositories. Constructing such a repository requires keeping track of project metrics together with related defect information. When automated tools are not used, manual effort is inevitable to maintain these repositories.

On the other hand, there are public data repositories including projects from companies such as NASA [1]. In this context, we investigate whether these public project data can be helpful for other companies for building localized defect predictors¹, especially for those with limited or no defect data repository.

Defect prediction literature contains many examples where predictors are learned from within company data [2, 3, 4, 5, 6, 7, 8, 9]. However, there exists no case that attempts to harness cross company data. In this paper, we focus on *binary defect prediction* and perform analyses to check if we can reach a conclusion in favor of either CC or WC data. Specifically, this paper assesses the relative merits of cross-company (CC) vs within-company (WC) data for defect prediction.

We perform a series of analyses on 10 publicly available project data from different companies. Our analyses use the Naive Bayes data miner and are based on static code features².

Our analyses aim at answering the following research questions:

- *Are CC data ever useful for organizations?:*

Our goal is to identify the conditions under which cross-company data should be preferred to within-company for the purposes of structuring software inspections using outcomes of defect predictors. In the first analysis, we compare the performance of defect predictors learned from WC data to those learned from CC data.

- *How can companies filter CC data for local tuning?:*

CC data includes information from many diverse projects and are heterogeneous compared to WC data. The goal of the second analysis is to select a subset of the available CC data that is similar to WC data and to investigate the effect of data homogeneity on the defect prediction performance. We apply a simple nearest neighbor (NN) filtering to CC data for automatically constructing a locally tuned repository. We use the Euclidean distance between static code features of WC and CC data for measuring similarity and automated selection of neighbors.³

¹ Throughout the paper, the following notation is used: a defect predictor (or predictor) means a binary classification method that categorizes software modules as either defective or defect-free; data refers to $M \times N$ matrices of raw measurements of N metrics from M software modules; these N metrics are referred to as features.

² Therefore, throughout the paper, the term "data" refers to static code features.

³ We should carefully note that we do not make use of any conceptual similarities, since our analysis is based on static code features. As to the issue of conceptual connections within the code, we refer the reader to the concept location and cohesion work of Marcus et al. [10].

- *What is the smallest amount of local data needed for constructing a model?:*

In the third analysis, we focus on WC models and determine the smallest amount of local data needed for constructing a model. We employ an incremental learning approach to WC data in order to determine the number of samples in local repositories for building defect prediction models.

- *Can our results be generalized?:*

We initially use only NASA projects in our three analyses to answer the above questions. In order to check the external validity of our results, our last analysis replicates all three analyses on projects from a company that has no ties with NASA: specifically, a Turkish company writing software controllers for home appliances.

The contributions of this paper are relevant for both academia and practice. Prior to this paper, no study was performed to investigate the relative merits of using cross company or within company data for constructing defect predictors. In our analysis, in favor of the common belief, we empirically show that within company data is better for defect prediction problem. Further, on the contrary to the common belief, we empirically show that, the required local data for constructing defect predictors can be easily and quickly collected. For our analysis, we have collected data from industry and made it publicly available for the use of other researchers and practitioners. As a practical contribution, our analysis of cross company data and the proposed methodology allow the construction of defect predictors even for companies with no local defect data. In summary, we observe that the defect detectors learned from site A are found to be useful at site B, at least for the datasets we analyzed. Furthermore, it is a result of major significance that other project's data can be used locally (after applying a very simple filter). Nevertheless, there may be some reluctance on the part of programmers and managers to use data collected at other sites. Such caution is to be expected amongst professionals striving to deliver quality products to their clients. In this paper, we suggest that it can be quite beneficial to use data from other sites, particularly when no local data is available. Professional developers should apply our advice cautiously—perhaps via an initial case study on a small part of their system.

This paper is organized as follows: The next section motivates our work with examples of industrial groups that use these detectors. Then, in Section 3, we explain our methodology by describing the data sources, data mining algorithm and the performance evaluation criteria. In Sections 4, 5 and 6, we describe our analyses, results and discussions for the first three research questions stated above. In Section 7, we replicate all three analyses in order to check the external validity of our results. Sections 8, 9 and 10, discuss the related work, practical implications and threats to validity respectively. We conclude our work in Section 11.

2 Motivation

2.1 Why Study WC vs CC?

The genesis of this work was several frustrating years working at NASA trying to collect data on new projects. This proved difficult to do, for a variety of reasons (including feuding NASA centers and contractors trying to maintain some control over the information flow to the client). Several times in that process, we asked “Is it necessary to collect new data? Can't we just apply old data?”.

In the middle of 2007, we had an opportunity to evaluate CC data usage on a Turkish software development company. That company worked in a highly competitive market and

a decrease in their bug rate of even 1% had an economic benefit to them. However, that company had no historical data with which to train its defect detectors. We therefore trained a defect detector on NASA data and applied it to the Turkish software. Much to our surprise, that detector could “adequately” predict errors within Turkish software (“adequate” in the sense that the Turkish client looked at the code that triggered our detectors, found bugs, and was subsequently interested enough to request more studies of this nature).

Based on that small informal experience, we saw:

- The business benefits in using CC data: no need to wait for local data collection;
- The adequacy of the learned detectors;

A single case study is not enough to justify a general principle. Accordingly, we conducted the analyses described in this paper.

2.2 Why Study Defect Predictors?

Defect predictors are subtly different to other quality assurance practices. This section highlights those differences and explains why we elect to study them.

It is insightful to contrast defect prediction from Musa-style software reliability engineering (SRE) [11]. In standard SRE, there is a track record of post-release failures and the goal is to predict the probability of failure on demand in the next call on the software. One of the major results of the SRE literature is that as the software is developed and bugs are removed, this probability forms a reliability growth curve with predictable properties.

The standard defect prediction task is different to standard SRE. Our task is to predict defective modules during development and before delivery. Such early lifecycle detection is important since, as Boehm reports, repair cost increases exponentially, the longer a bug remains in the system [12]. Hence, we learn defect predictors from logs of pre-release defects found by a verification and validation (V&V) team. Such defects may be seen in code from systems that is not yet complete. Hence, many of the features used in this kind of analysis are static code features that can be extracted from as-yet-un-executing code.

The patterns found in pre-release defects or post-release failures can be used in different ways. Reliability growth curves can be used to predict how long before the running software will reach some required reliability of, i.e. one failure in 10,000 calls. Defect predictors, on the other hand, can be used to bias the ordering of modules to be inspected by V&V teams:

- In the case where there is insufficient resources to inspect all code (which is a very common situation in industrial developments), defect predictors can be used to increase the odds that the inspected code will have more defects.
- In the case where all the code is to be inspected, but that inspection process will take weeks to months to complete, defect predictors can be used to increase the chances that defective modules will be inspected earlier. This is useful since it gives the development team earlier notification of what modules require rework, hence giving them more time to complete that rework prior to delivery.

The above two cases are quite common. The authors of this paper are from the United States and Turkey and only began exploring defect detectors when they realized that their user community was making extensive use of these detectors.

Furthermore, these detectors are widely used despite the fact that they are quite “coarse-grained”. In the following study, our only knowledge of modules will be that they are defective or defect-free. This work makes no use of other knowledge about a module such as its criticality or the dates on which it failed.

- *Regarding failure occurrence logs:* When working with large teams of contractors, it is often difficult to access precise logs on when failures occurred since such information is often corporate critical. The NASA coarse-grained information used in this study, for example, took years of careful negotiation to acquire. All the data passed to us was highly sanitized; e.g. this research team could not even access project or module names.
- *Regarding criticality knowledge:* Our goal is to predict defective modules during development and before delivery. Goseva and Hamill [13] report that pre-delivery perceptions of module criticality can be widely misleading. In their work, they trace errors back to the modules that caused runtime faults. They find that a remarkably small set of modules are implicated in all faults and that set can change, radically, if the input space to the program alters.

For further notes on related work on defect predictors from static code attributes, see Section 8.3.

3 Methodology

3.1 Data

The analyses of this paper use the static code features of 10 projects tabulated in Table 1 and Table 2, which are downloaded from the PROMISE repository⁴. Available static code features are shown in Table 3. An advantage of static code features is that they can be quickly and automatically collected from the source code, even if no other information is available.

Table 1 Descriptions of ten software projects used in the paper. The rows labeled “NASA” come from NASA aerospace projects while the rows labeled “SOFTLAB” come from a Turkish software company writing applications for domestic appliances.

source	project	language	description
NASA	pc1	C++	Flight software for earth orbiting satellite
NASA	kc1	C++	Storage management for ground data
NASA	kc2	C++	Storage management for ground data
NASA	cm1	C++	Spacecraft instrument
NASA	kc3	JAVA	Storage management for ground data
NASA	mw1	C++	A zero gravity experiment related to combustion
SOFTLAB	ar4	C	Embedded controller for white-goods
SOFTLAB	ar3	C	Embedded controller for white-goods
NASA	mc2	C++	Video guidance system
SOFTLAB	ar5	C	Embedded controller for white-goods

Our project data are taken from software developed in different geographical locations across North America (NASA) and Turkey (SOFTLAB). Therefore, the static code features that are available for each project vary. While Table 3 shows the features available for each project, Table 4 shows the common features for both sources. While NASA and SOFTLAB are one single *source* of data, there are several projects within each source. For example, NASA is really an umbrella organization used to co-ordinate and fund a large and diverse set of projects:

⁴ <http://promisedata.org/repository>

Table 2 Summary of data from ten software projects of Table 1, sorted in order of number of functional units.

source	(# modules)		features	loc	%defective
	project	examples			
NASA	pc1	1,109	21	25,924	6.94
NASA	kc1	845	21	42,965	15.45
NASA	kc2	522	21	19,259	20.49
NASA	cm1	498	21	14,763	9.83
NASA	kc3	458	39	7749	9.38
NASA	mw1	403	37	8341	7.69
SOFTLAB	ar4	107	29	9196	18.69
SOFTLAB	ar3	63	29	5624	12.70
NASA	mc2	61	39	6134	32.29
SOFTLAB	ar5	36	29	2732	22.22
				4,102	

- The NASA data was collected from across the United States over a period of five years from numerous NASA contractors working at different geographical centers.
- These projects represent a wide array of projects, including satellite instrumentation, ground control systems and partial flight control modules (i.e. Altitude Control).
- The data sets also represent a wide range of code reuse: some of the projects are 100% new, and some are modifications to previously deployed code.

That is, even if we explore just the NASA data sets, we can still examine issues of cross- vs within- company data use. Nevertheless, using our connections with the Turkish software industry, we collected new data sets in the format of Table 3 from a Turkish white-goods manufacturer. The SOFTLAB datasets ($\{ar3, ar4, ar5\}$) in Table 1, are the controller software for:

- A washing machine;
- A dishwasher;
- And a refrigerator.

In summary, 7 datasets are from NASA projects developed at different sites by different teams, hence we treat each of them as if they were from 7 different companies. Remaining 3 datasets are from a Turkish company collected from software for domestic home appliances. Therefore we use 10 projects from 8 different companies.

In all analyses, we used all available features to learn defect predictors. One question is whether it might be better to focus on just some special subset of those features. In previous work [14], we have found that there is no best single set of features. In fact, this *best* set is highly unstable and changes dramatically from data set to data set. That study is summarized, below. The main message of that study is that it is “best” to never assume a *best* set of features. Rather, give them all to a learner and let it work out which ones are best to use or ignore for a particular application.

Feature subset selection is a technique that explores subsets of the available features. The simplest and fastest subsetting method is to rank attributes from the most informative to least informative. After discretizing numeric data⁵ then if F is a set of features, the number of bits required to encode a class after observing a feature is:

$$H(C|F) = -\sum_{f \in F} p(f) \sum_{c \in C} p(c|f) \log_2(p(c|f)) \quad (1)$$

⁵ E.g. given an attribute’s minimum and maximum values, replace a particular value n with $(n - \min)/((\max - \min)/10)$. For more on discretization, see [15].

Table 3 Static code features available in Table 2 projects.

#	Feature	pc1	kc1	kc2	cm1	kc3	mw1	ar4	ar3	mc2	ar5
1	<i>branchcount</i>	X	X	X	X	X	X	X	X	X	X
2	<i>codeandcommentloc</i>	X	X	X	X	X	X	X	X	X	X
3	<i>commentloc</i>	X	X	X	X	X	X	X	X	X	X
4	<i>cyclomaticcomplexity</i>	X	X	X	X	X	X	X	X	X	X
5	<i>designcomplexity</i>	X	X	X	X	X	X	X	X	X	X
6	<i>halsteaddifficulty</i>	X	X	X	X	X	X	X	X	X	X
7	<i>halsteadeffort</i>	X	X	X	X	X	X	X	X	X	X
8	<i>halsteaderror</i>	X	X	X	X	X	X	X	X	X	X
9	<i>halsteadlength</i>	X	X	X	X	X	X	X	X	X	X
10	<i>halsteadtime</i>	X	X	X	X	X	X	X	X	X	X
11	<i>halsteadvolume</i>	X	X	X	X	X	X	X	X	X	X
12	<i>totaloperands</i>	X	X	X	X	X	X	X	X	X	X
13	<i>totaloperators</i>	X	X	X	X	X	X	X	X	X	X
14	<i>uniqueoperands</i>	X	X	X	X	X	X	X	X	X	X
15	<i>uniqueoperators</i>	X	X	X	X	X	X	X	X	X	X
16	<i>executableloc</i>	X	X	X	X	X	X	X	X	X	X
17	<i>totalloc</i>	X	X	X	X	X	X	X	X	X	X
18	<i>halsteadcontent</i>	X	X	X	X	X				X	
19	<i>essentialcomplexity</i>	X	X	X	X	X	X			X	
20	<i>halsteadvocabulary</i>	X	X	X	X			X	X		X
21	<i>blankloc</i>	X	X	X	X		X	X	X	X	X
22	<i>callpairs</i>					X	X	X	X	X	X
23	<i>conditioncount</i>					X	X	X	X	X	X
24	<i>cyclomaticdensity</i>					X	X	X	X	X	X
25	<i>decisioncount</i>					X	X	X	X	X	X
26	<i>decisiondensity</i>					X	X	X	X	X	X
27	<i>halsteadlevel</i>					X	X	X	X	X	X
28	<i>multipleconditioncount</i>					X	X	X	X	X	X
29	<i>designdensity</i>					X	X	X	X	X	X
30	<i>normcyclomaticcomp.</i>					X	X	X	X	X	X
31	<i>formalparameters</i>					X	X	X	X	X	X
32	<i>modifiedconditioncount</i>					X	X			X	
33	<i>maintenanceseverity</i>					X	X			X	
34	<i>edgecount</i>					X	X			X	
35	<i>nodecount</i>					X	X			X	
36	<i>essentialdensity</i>					X	X			X	
37	<i>globaldatacomplexity</i>					X				X	
38	<i>globaldatadensity</i>					X				X	
39	<i>percentcomment</i>					X	X			X	
40	<i>numberoflines</i>					X	X			X	
	Total	21	21	21	21	39	37	29	29	39	29

The highest ranked feature F_i is the one with the largest *information gain*; i.e the one that most reduces the encoding required for the data *after* using that feature; i.e.

$$InfoGain(F_i) = H(C) - H(C|F_i) \quad (2)$$

where $H(C)$ comes from Equation 1. In *iterative InfoGain subsetting*, predictors are learned using the $i = 1, 2, \dots, N$ -th top-ranked attributes. Subsetting terminates when $i + 1$ attributes perform no better than i . In *exhaustive InfoGain subsetting*, the attributes are first ranked using iterative subsetting. Next, predictors are built using all subsets of the top j ranked attributes. For both iterative and exhaustive subsetting, the process is repeated 10 times using 90% of the data (randomly selected). Iterative subsetting takes time linear on the number of attributes N while exhaustive subsetting takes time 2^j (so it is only practical for small $j \leq N$).

Figure 1 shows the feature subset selection results reported previously [14]. Note that this study was run on many of the same NASA data sets used in this study. The right-hand-

data	N	%		selected attributes (see Figure 2)	selection method
		pd	pf		
pc1	100	48	17	3, 35, 37	exhaustive subsetting
mw1	100	52	15	23, 31, 35	iterative subsetting
kc3	100	69	28	16, 24, 26	iterative subsetting
cm1	100	71	27	5, 35, 36	iterative subsetting
pc2	100	72	14	5, 39	iterative subsetting
kc4	100	79	32	3, 13, 31	iterative subsetting
pc3	100	80	35	1, 20, 37	iterative subsetting
pc4	100	98	29	1, 4, 39	iterative subsetting
all	800	71	25		

Fig. 1 Best defect predictors learned in [14]. Mean results from Naive Bayes after a 10 repeats of (i) randomize the order of the data; (ii) divide that data into ten 90%:10% splits for training:test. Prior to learning, all numerics were replaced with logarithms. InfoGain was then used to select the best two or three attributes shown in the right-hand column (and if “three” performed as well as “two”, then this table shows the results using “two”).

ID	frequency in Figure 1	what	type
1	2	loc_blanks	locs
3	2	call_pairs	misc
4	1	loc_code_and_command	locs
5	2	loc_comments	locs
13	1	edge_count	misc
16	1	loc_executable	locs
20	1	I	H (derived Halstead)
23	1	B	H (derived Halstead)
24	1	L	H (derived Halstead)
26	1	T	H (derived Halstead)
31	2	node_count	misc
35	3	μ_2	h (raw Halstead)
36	1	μ_1	h (raw Halstead)
37	2	number_of_lines	locs
39	2	percent_comments	misc

Fig. 2 Attributes used in Figure 1.




Fig. 3 InfoGain for KC3 attributes. Calculated from Equation 2. Lines show means and t-bars show standard deviations after 10 trials on 90% of the training data (randomly selected).

Table 4 Static code features shared by NASA and SOFTLAB projects.

#	Feature	NASA shared	All Shared
1	<i>branchcount</i>	X	X
2	<i>codeandcommentloc</i>	X	X
3	<i>commentloc</i>	X	X
4	<i>cyclomaticcomplexity</i>	X	X
5	<i>designcomplexity</i>	X	X
6	<i>halsteaddifficulty</i>	X	X
7	<i>halsteadeffort</i>	X	X
8	<i>halsteaderror</i>	X	X
9	<i>halsteadlength</i>	X	X
10	<i>halsteadtime</i>	X	X
11	<i>halsteadvolume</i>	X	X
12	<i>totaloperands</i>	X	X
13	<i>totaloperators</i>	X	X
14	<i>uniqueoperands</i>	X	X
15	<i>uniqueoperators</i>	X	X
16	<i>executableloc</i>	X	X
17	<i>totalloc</i>	X	X
18	<i>halsteadcontent</i>	X	
19	<i>essentialcomplexity</i>	X	
	Total	19	17

side column of that figure reports which attributes used in that study were selected by feature subset selection. In a result consistent with many of the static code measures being highly correlated, note that most of the features were removed and using 3 features was often as good as using many more.

In terms of “which features are best to use?”, the key feature of Figure 1 is that different data sets selected very different “best” features. Figure 2 tallies the number of times each feature was selected: no feature was selected all the time and many of the features were never selected at all. This aspect can be explained by Figure 3 which shows the InfoGain of all the features in one of our data sets (KC3). Note how the highest ranked attributes (those on the left-hand-side) offer very similar information. That is, there are no clear winners so minor changes in the training sample (the 90% sub-sampling used in subsetting or a cross-validation study) can result in the selection of very different “best” attributes.

The pattern of InfoGain values of Figure 3 (where there are many alternative “best” features) repeats in all our data sets. This pattern explains a prior observation of Shepperd & Ince who found 18 publications where an equal number of studies reporting that the McCabe cyclomatic complexity is the same; is better; or is worse than lines of code in predicting defects [16]. Figure 3 motivates the following principles:

- Do not seek “best” subsets of static code attributes.
- Rather, seek instead for learning methods that can combine multiple partial defect indicators like the statistical methods of Naive Bayes.

Therefore, we do not perform feature subset selection, rather we use all available features in our analyses.

3.2 Naive Bayes Classifier

Each project in Table 2 contains information from many modules; the smallest unit of functionality⁶. To learn defect predictors, the project data are augmented with one dependent variable holding boolean values for “defects detected”. The data mining task is to find combinations of static code features that predict the boolean value of the dependent variable.

In prior work we have explored a range of data mining methods for defect prediction and found that classifiers based on Bayes theorem yields better performance than rule based methods (i.e. decision trees, oneR), for the Table 2 NASA data⁷ [14]. An extensive study by Lessmann et.al also shows that Naive Bayes performs equivalently well with 15 other methods [17]. Therefore, we find Naive Bayes as a viable choice as a classifier to use in our analysis.

A Naive Bayes classifier [18] is based on Bayes’ Theorem. Informally, the theorem says $next = old * new$ i.e. what we’ll believe $next$ comes from how new evidence effects old beliefs. More formally:

$$P(H|E) = \frac{P(H)}{P(E)} \prod_i P(E_i|H) \quad (3)$$

i.e. given fragments of evidence E_i and a prior probability for a class $P(H)$, the theorem lets us calculate a posterior probability $P(H|E)$. For example, in our data sets, there are two hypotheses: modules are either defective or not: $H \in \{defective, nonDefective\}$. Also, if a particular module has $numberOfSymbols = 27$ and $LOC = 40$ and was previously classified as “defective” then

$$\begin{aligned} E_1 &: numberOfSymbols = 27 \\ E_2 &: LOC = 40 \\ H &: defective \end{aligned}$$

When building defect predictors, the prior probability of each class (“defective” or “defect-free”) is calculated, given the defect content of each module. So, if a data set has 100 modules and 25 of them are faulty, then:

$$P(defective) = 0.25$$

When testing new data, a module is assigned to the class with the higher posterior probability, calculated from Equation 3.

Naive Bayes classifiers are called “naive” since they assume independence of each feature. While this assumption simplifies the implementation (only the frequency counts are required for each feature), it is possible that correlated events are missed by this “naive” approach. Domingos and Pazzani show theoretically that the independence assumption is a problem in a vanishingly small percent of cases [19]. This explains the repeated empirical result that, on average, Naive Bayes classifiers perform as well as other seemingly more sophisticated schemes⁸.

Equation 3 offers a simple method for handling missing values. Generating a posterior probability means of tuning a prior probability to new evidence. If that evidence is missing, then no tuning is needed. In this case Equation 3 sets $P(E_i|H) = 1$ which, in effect, makes no change to $P(H)$.

⁶ In other languages, modules may be called “function” or “method”.

⁷ SOFTLAB data were not available at that time.

⁸ Details of this issue are out of the scope of this paper. For more, please see Table 1 in [19]

When estimating the prior probability of hypothesis H , it is common practice [20] to use an M -estimate as follows. Given that the total number of hypothesis is C , the total number of training instances is I , and $N(H)$ is the frequency the hypothesis H within I , then

$$P(H) = \frac{N(H) + m}{I + m \cdot C} \quad (4)$$

Here m is a small non-zero constant (often, $m = 1$). Three special cases of Equation 4 are:

- For high frequency hypothesis in large training sets, $N(H)$ and I are much larger than m and $m \cdot C$, so Equation 4 simplifies to $P(H) = \frac{N(H)}{I}$, as one might expect.
- For low frequency classes in large training sets, $N(H)$ is small, I is large, and the prior probability for a rare class is never less than $\frac{1}{I}$; i.e. the inverse of the number of instances. If this were not true, rare classes would never appear in predictions.
- For very small data sets, I is small and $N(H)$ is even smaller. In this case, Equation 4 approaches the inverse of the number of classes; i.e. $\frac{1}{C}$. This is a useful approximation when learning from very small data sets when all the data relating to a certain class has not yet been seen.

The prior probability calculated in Equation 4 is a useful lower bound for $P(E_i|H)$. If some value v is seen $N(f = v|H)$ times in feature f 's observations for hypothesis H , then

$$P(E_i|H) = \frac{N(f = v|H) + l \cdot P(H)}{N(H) + l} \quad (5)$$

Here, l is the L -estimate and is set to a small constant (Yang & Webb [20] recommend $l = 2$). Two special cases of are:

- A common situation is when there are many examples of an hypothesis and numerous observations have been made for a particular value. In that situation, $N(H)$ and $N(f = v|H)$ are large and Equation 5 approaches $\frac{N(f=v|H)}{N(H)}$, as one might expect.
- In the case of very little evidence for a rare hypothesis, $N(f = v|H)$ and $N(H)$ are small and Equation 5 approaches $\frac{l \cdot P(H)}{l}$; i.e. the default frequency of an observation in a hypothesis is a fraction of the probability of that hypothesis. This is a useful approximation when very little data is available.

For numeric features it is common practice for Naive Bayes classifiers to use the Gaussian probability density function [21]:

$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (6)$$

where $\{\mu, \sigma\}$ are the feature's {mean, standard deviation}, respectively. To be precise, the probability of a continuous feature having exactly the value x is zero, but the probability that it lies within a small region, say $x \pm \epsilon/2$, is $\epsilon \times g(x)$. Since ϵ is a constant that weighs across all possibilities, it cancels out and needs not be computed.

All the static code features of Table 3 are numeric and are highly skewed. Therefore, we replace all numeric values with a "log-filter", i.e. N with $\ln(N)$. This spreads out skewed curves more evenly across the space from the minimum to maximum values (to avoid numerical errors with $\ln(0)$, all numbers under 0.000001 are replaced with $\ln(0.000001)$). This "spreading" can significantly improve the effectiveness of data mining, since the distribution of log-filtered feature values fits better to the normal distribution assumption [14].

3.3 Performance Evaluation

Data mining effectiveness was measured using *probability of detection* (pd), *probability of false alarm* (pf) and *balance* (bal) [5, 14]. If $\{A, B, C, D\}$ are the true negatives, false negatives, false positives, and true positives (respectively) found by a defect predictor, then:

$$pd = recall = D/(B + D) \quad (7)$$

$$pf = C/(A + C) \quad (8)$$

$$bal = balance = 1 - \frac{\sqrt{(0 - pf)^2 + (1 - pd)^2}}{\sqrt{2}} \quad (9)$$

All these values range zero to one. Better and *larger* balances fall *closer* to the desired zone of no false alarms ($pf = 0$) and 100% detection ($pd = 1$).

Other measures such as *accuracy* and *precision* were not used since, as shown in Table 2, the percent of defective examples in our projects was usually very small (median value around 13%). Accuracy and precision are poor indicators of performance for data where the target class (i.e. defective) is so rare (for more on this issue, see [5, 14]).

The results were visualized using *quartile charts*. To generate these charts, the performance measures for an analysis are sorted to isolate the median and the lower and upper quartile of numbers. For example:

$$\overbrace{\{4, 7, 15, 20, 31\}}^{q1}, \quad \overbrace{40}^{median}, \quad \overbrace{52, 64, 70, 81, 90}^{q4}$$

In our quartile charts, the upper and lower quartiles are marked with black lines; the median is marked with a black dot; and vertical bars are added to mark the 50% percentile value. The above numbers would therefore be drawn as follows:

$$0\% \mid \text{---} \quad \bullet \quad \text{---} \mid 100\%$$

The Mann-Whitney U test [22] was used to test for statistical difference between results. This non-parametric test replaces performance measure values (i.e. pd , pf , bal) with their rank inside the population of all sorted values. Such non-parametric tests are recommended in data mining since many of the performance distributions are non-Gaussian [23].

4 Analysis #1: Are CC data ever useful for organizations?

4.1 Design

Our goal is to determine whether using cross company data is beneficial for constructing defect predictors and to identify the conditions under which cross-company data should be preferred to within-company data. Our first WC-vs-CC analysis follows the pseudo code given in Table 5, for all 7 NASA projects of Table 1. For each project, test sets were built from 10% of the data, selected at random. Defect predictors were then learned from:

- CC data: all data from the other 6 projects.
- WC data: remaining 90% data of that project;

Most of the Table 1 data come from systems written in “C/C++” but at least one of the systems was written in JAVA. For cross-company data, an industrial practitioner may not have access to detailed meta-knowledge (e.g. whether it was developed in “C” or JAVA). They may only be aware that data, from an unknown source, are available for download from a certain url. To replicate that scenario, we will make no use of our meta-knowledge about Table 1.

In order to control for *order effects* (where the learned theory is unduly affected by the order of the examples) our procedure was repeated 20 times, randomizing the order of data in each project each time. In all, we ran 280 tests to compare WC-vs-CC:

$$(2 \text{ data sources}) * (20 \text{ randomized orderings}) * (7 \text{ projects})$$

Table 5 Pseudocode for Analysis 1

```

DATA = [PC1, KC1, KC2, CM1, KC3, MW1, MC2]           // all available data
LEARNER = [Naive Bayes]                               // defect predictor

C_FEATURES <- Find common features IN DATA
FOR EACH data IN DATA
    data = Select C_FEATURES in data                  // use common features
END
REPEAT 20 TIMES
    FOR EACH data in DATA
        CC_TRAIN = DATA - data                       // cross company training data
        WC_TRAIN = random 90% of data                 // within company training data
        TEST = data - WC_TRAIN                       // shared test data

        //construct predictor from CC data
        CC_PREDICTOR = Train LEARNER with CC_TRAIN
        // construct predictor from WC data
        WC_PREDICTOR = Train LEARNER with WC_TRAIN
        //Evaluate both predictors on the same test data
        [cc_pd, cc_pf, cc_bal] = CC_PREDICTOR on TEST
        [wc_pd, wc_pf, wc_bal] = WC_PREDICTOR on TEST
    END
END

```

The projects’ data come from different sources and, hence, have different features. For this analysis, only the features that are common in all NASA projects are used, a total of 19 features. These features are marked in “NASA Shared” column of Table 4.

4.2 Results from Analysis #1

Figure 4 shows the $\{pd, pf\}$ quartile charts for CC vs WC data averaged over seven NASA projects. The pattern is very clear: CC data dramatically increases *both* the probability of detection and the probability of false alarms. The *pd* results are particularly striking.

For cross-company data:

- 50% of the *pd* values are at or above 97%
- 75% of the *pd* values are at or above 83%;
- And all the *pd* values are at or over 50%.




Fig. 4 Analysis #1 results averaged over seven NASA tables. Numeric results on left; quartile charts on right. “Q1” and “Q3” denote the 25% and 75% percentile points (respectively). The upper quartile of the first row is not visible since it runs from 100% to 100%; i.e. it has zero length.

To the best of our knowledge, Figure 4 are the largest pd values ever reported from these data. However, these very high pd values come at some considerable cost. We observe in Figure 4 that the median false alarm rate has changed from 29% (with WC) to 64% (with CC) and the maximum pf rate now reaches 100%. Note that a 100% pf rate means that all defect-free modules are classified as defective, which yields inspection of all these modules unnecessarily and contradicts with the purpose of defect prediction. However, it is not right to assess the general behavior of the CC defect predictors with such an extreme case. We mention this issue in order to clarify that high false alarm rates may be prohibitive in the practical application of defect predictors.

We explain these increases in pd, pf with the *extraneous factors* in CC data. More specifically, using a large training set (e.g. seven projects in Table 2) informs of not only the sources of errors, but also numerous irrelevancies. For example, the defect characteristics of software modules with different complexity levels or sizes may differ [24]. In this context, a complicated search algorithm’s metrics are irrelevant to the defective behavior of a simple *sum* function. Since there are few modules with extreme characteristics (i.e. complexity, size) in a single project, their effect on the overall model are limited. However, when data from multiple companies are combined, the number of these extreme cases, and hence their cumulative effect on the overall model increase significantly. Therefore, factors such as programming language constructs (i.e. object oriented vs. procedural) and project specific requirements (i.e. availability, speed) that have impacts on the module characteristics can be considered as extraneous factors.

Hence, large training sets increase the error detection rates (i.e. pd) since there are more known sources of errors. However, they also increase the probability of false alarms (pf) since there are more *extraneous factors* introduced to the analysis. We will test the validity of this claim in the next analysis.

4.3 Checking the Analysis #1 Results

Once a *general result* is defined (e.g. defect predictors learned from CC projects dramatically increase both pf and pd), it is good practice to check for specific exceptions to that pattern. Table 6 shows a summary of results when U tests with $\alpha = 0.05$ were applied to test results from each of the 7 projects, *in isolation* and Figure 5 shows the $\{pd, pf\}$ quartile charts for Analysis #1 for each NASA project:

- For six projects, the general result holds (i.e. both (pd, pf) increases if defect predictors learned from CC projects are used rather than defect predictors learned from WC projects. See group *a* in Table 6).

Table 6 Summary of U-test results (95% confidence): moving from WC to CC. For all projects' results, see Figure 5.

group	pd $WC \rightarrow CC$	pf $WC \rightarrow CC$	tables	$ tables $
a	increased	increased	CM1 KC1 KC2 MC2 MW1 PC1	6
b	same	same	KC3	1

- For one project, there is no difference in the performances of the defect predictors learned from CC and WC projects (see group *b* in Table 6).

4.4 Discussion of Analysis #1

When practitioners use defect predictors with high false alarm rates (e.g. the 64% reported above), they must allocate a large portion of their debugging budget to the unfruitful exploration of erroneous alarms.

In our industrial work, we have sometimes seen several situations where detectors with high false alarms are useful:

- *When the cost of missing the target is prohibitively expensive.* In mission critical or security applications, the goal of 100% detection may be demanded in all situations, regardless of the cost of chasing false alarms.
- *When only a small fraction of the data is returned.* Hayes, Dekhtyar, & Sundaram call this fraction *selectivity* and offer an extensive discussion of the merits of this measure [25].
- *When there is little or no cost in checking false alarms.*

Having said that, we have shown these results to five of our user groups in the United States and Turkey and none of them could accept false alarm rates as high as 64%. Therefore, the conditions under which the benefits of CC data (high probabilities of detection) outweigh their high costs (high false alarm rates) are not common.

In summary, for most software applications, very high pf rates like the CC results of Figure 4 make the predictors impractical to use.

5 Analysis #2: How can companies filter CC data for local tuning?

The results of the first analysis restricts the use of CC data to a limited domain (i.e. mission critical) and may look discouraging at first glance. In the first analysis we explained our observations with the *extraneous factors*, that the results are affected by the irrelevant factors in CC data. In this section we hypothesize this claim and test for its validity.

5.1 Design

In this analysis, we try to construct more homogeneous *defect* datasets from CC data in an automated manner. For this purpose we use a simple filtering method (i.e. NN) which is to be described next. The analysis design is given in Table 7. For this analysis, we use all




Fig. 5 Project-wise Analysis #1 results for NASA projects.

common features available in NASA projects which is a total of 19 features. These features are marked in NASA Shared column of Table 4.

Table 7 Pseudocode for Analysis 2

```

DATA = [PC1, KC1, KC2, CM1, KC3, MW1, MC2]
LEARNER = [Naive Bayes]

C_FEATURES <- Find common features IN DATA
FOR EACH data IN DATA
    data = Select C_FEATURES in data
END
REPEAT 100 TIMES
    FOR EACH data in DATA
        WC_TRAIN = random 90% of data
        TEST = data - WC_TRAIN
        CC_TRAIN = DATA - data

        //NN-Filtering: Select k=10 neighbors for each validation data

        FOR EACH test IN TEST
            dist = L2_DISTANCE{test, CC_TRAIN}
            NNCC_TRAIN <- 10 Samples in CC_TRAIN with min{dist}
        END

        //Use only unique samples from the union of all selected neighbors

        NNCC_TRAIN = UNIQUE(NNCC_TRAIN)

        NNCC_PREDICTOR = Train LEARNER with NNCC_TRAIN
        CC_PREDICTOR = Train LEARNER with CC_TRAIN
        WC_PREDICTOR = Train LEARNER with WC_TRAIN
        [nncc_pd, nncc_pf, nncc_bal] = NNCC_PREDICTOR on TEST
        [cc_pd, cc_pf, cc_bal] = CC_PREDICTOR on TEST
        [wc_pd, wc_pf, wc_bal] = WC_PREDICTOR on TEST
    END
END

```

5.1.1 Nearest Neighbor (NN) Filtering

Our idea behind filtering is to collect similar instances together in order to construct a learning set that is homogeneous with the validation set. More formally, we try to introduce a bias in the model by using training data that are similar to the validation data characteristics. While a bias is not desired in general, it is what we seek on purpose, since in our case we can control the bias against removing the noise in CC data (i.e. extraneous factors). We simply use the k-Nearest Neighbor (k-NN) method to measure the similarity between the validation set and the training *candidates*. The similarity measure is the Euclidean distance between the static code features of validation and training candidate sets. As mentioned before, we do not perform any feature selection and use all available common features for filtering. The expected outcome of the filtering part is to obtain a subset of available CC data that shows similar characteristics to the local code culture.

The details of NN filter are as follows: We calculate the pairwise Euclidean distances between the validation set and the *candidate* training set (i.e. all CC data). Let N be the number of validation set size. For each validation data, we pick its $k = 10$ nearest neighbors from candidate training set⁹. Then we come up with a total of $10 \times N$ similar data points (i.e. module feature vectors). These $10 \times N$ samples may not be unique (i.e. a single data sample can be a nearest neighbor of many data samples in the validation set). Using only unique ones, we form the training set and use a random 90% of it for training a predictor.

⁹ Caveat: We did not optimize the value of k for each project. We simply used a constant $k = 10$. We consider it as a future work to dynamically set the value of k for a given project [26].




Fig. 6 Analysis #2 PD results where $NN_{pd} \geq WC_{pd}$. Rankings computed via Mann-Whitney (95% confidence) comparing each row to all other rows.

Please note that we do not use the class information (i.e. a module is defective or defect-free) while measuring similarity and selecting neighbors. This corresponds to a real life case, where the development of some modules are completed and they are ready for testing: there is no defect related information available, however static code features are collected with automated tools.

5.2 Results

If the results of Analysis #1 are due to the extraneous factors in CC data (as we suspect), then we would expect lower pf 's in NN results than CC results. Therefore, for pd , we define the null hypothesis as:

$$H_0 : NN_{pf} \geq CC_{pf}$$

H_0 is rejected by the U test with $\alpha = 0.05$. Therefore, using NN filtered CC data significantly decreases the false alarms compared to CC data. Yet, we observe that pd 's have also decreased. However, false alarm rates are more dramatically decreased than detection rates as seen in Figure 6 to Figure 9. For example, in *CM1* project, the median false alarm rate decreases nearly one third, from 91% to 33%, whereas the median detection rate slightly decreases from 94% to 82%. In all cases, predictors learned from NN data dramatically reduce the high false alarm rates associated with the use of cross-company data. Often that reduction halves the false alarm rate. For example, in *MW1*, the median false alarm rate drops from 68% (CC) to 33% (NN).

Showing that defect predictors learned from NN-filtered CC data are significantly better than the ones learned from CC data, the ideal result would be that models based on NN-filtered can be used as a surrogate for WC data based models. This, in turn, would mean that developers could avoid the tedious and expensive work of local data collection. We now investigate the relation between the NN-filtered CC data and the WC data based defect predictors.

If defect predictors learned from NN-filtered CC data out-performed the ones learned from WC data then we would expect to notice two observations:

- *Observation1*: NN would have pd values above or equal to WC's pd . The examples displaying *Observation1* are shown in Figure 6.
- *Observation2*: NN would have pf values below or equal to WC's pf . The examples displaying *Observation2* are shown in Figure 8.




Fig. 7 Analysis #2 PD results where $NN_{pd} < WC_{pd}$.




Fig. 8 Analysis s #2 PF results where $NN_{pf} \leq WC_{pf}$.

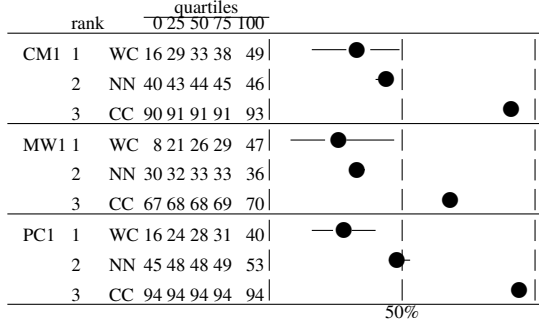


Fig. 9 Analysis #2 PF results where $NN_{pf} > WC_{pf}$.

Please note that the conjunction of *Observation1* and *Observation2* is uncommon. In fact, our results suggest that *Observation1* and *Observation2* are somewhat mutually exclusive:

- As shown in Figures 6 and 9, the examples where *NN* increases the probability of detection are also those where it increases the probability of false alarms.

Hence, we cannot recommend *NN* as a replacement for *WC*. Nevertheless, if local *WC* data are not available, then we would recommend processing foreign *CC* data with *NN*.

5.3 Discussions

In Analysis #1 we have used random samples of *CC* data and observed that the false alarm rates substantially increased compared to the *WC* models. Our new analysis shows that *NN* filtering *CC* data removes the increased false alarm rates. Now we argue that, using *NN* filtering instead of using all available *CC* data, helps choosing training examples that are similar to problem at hand. Thus, the irrelevant information in non-similar examples is avoided. However, this also removes the rich sample base and yields a slight decrease in detection rates. Mann-Whitney tests reveal that *NN* filtering is

- far better than random sampling cross company data,
- and still worse than using within company data.

The performances of defect predictors based on the *NN*-filtered *CC* vs. *WC* data do not give necessary empirical evidence to make a strong conclusion. Sometimes *NN* data based models may perform better than *WC* data based models. A possible reason may be hidden in the processes that are implemented those projects. May be, a group of new developers were working together for the first time and corresponding *WC* data included more heterogeneity, which is reduced by *NN*. May be the development methodology changed during the project, producing a different code culture. However, we do not have access to the internals of these projects that allows a discussion of these observations.

6 Analysis #3: What is the smallest amount of local data needed for constructing a model?

Our results of Analyses #1 and #2 reveal that *WC* data models are better if data are available. In this section, we will show that defect predictors can be learned from very small samples of *WC* data.

6.1 Design

An important aspect of the Analyses #1 and #2 results is that defect predictors were learned using only a handful of defective modules. For example, consider a 90%/10% train/test split on *pc1* with 1,109 modules, only 6.94% of which are defective. On average, the training set will only contain $1109 * 0.9 * 6.94/100 = 69$ defective modules. Despite this, *pc1* yields an adequate median $\{pd, pf\}$ results of $\{63, 27\}\%$.

Analysis #3 was therefore designed to check the smallest amount of data needed to learn defect predictors. The design is given in Table 8. Analysis #3 is essentially the same as the first analysis, but without the cross-company study. Instead, analysis #3 took the 7 NASA projects of Table 1 and learned predictors using:

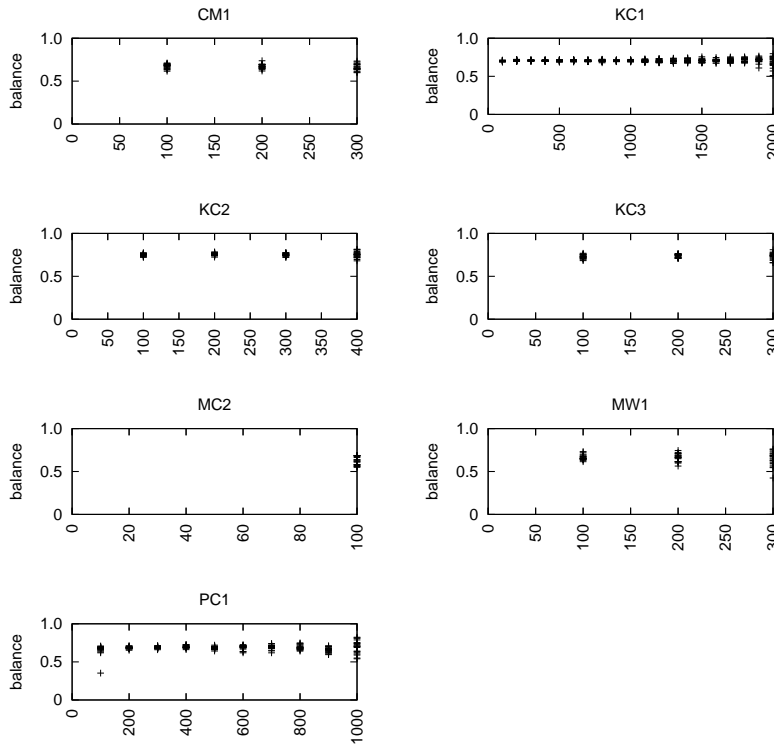


Fig. 10 Results from analysis #3. Training set size grows in units of 100 examples, moving left to right over the x-axis. The MC2 results only appear at the maximum x-value since MC2 has less than 200 examples.

Table 8 Pseudocode for Analysis 3

```

DATA = [PC1, KC1, KC2, CM1, KC3, MW1, MC2]
LEARNER = [Naive Bayes]

REPEAT 100 TIMES
  FOR EACH data IN DATA
    WC_TRAIN = random 90% of data
    TEST = data - WC_TRAIN
    FOR i IN {100, 200, 300, ...}
      WC_TRAIN_INCREMENTAL <- Random i Examples from WC_TRAIN
      WC_INC_PREDICTOR = Train LEARNER with WC_TRAIN_INCREMENTAL
      [iwc_pd, iwc_pf, iwc_bal] = WC_INC_PREDICTOR on TEST
    END
  END
END

```

- reduced WC data: a randomly selected subset of up to 90% of each project data.

After randomizing the order of the data, training sets were built using just the first 100, 200, 300, ... data samples in the project. After training the defect predictor, its performance is tested on the remaining data samples not used in training.




Fig. 11 Y-axis shows plateau point after learning from data sets that have up to X examples (from [28]). The left plot shows results from using Naive Bayes (nbk) or a decision tree learner (j48) [29] to predict for discrete classes. Right plot shows results from using linear regression (lsr) or model trees (m5) [30] to learn predictors for continuous classes. In this study, data sets were drawn from the UC Irvine data repository [31].

Analysis #1 only used the features found in all NASA projects. For this analysis, we imposed no such restrictions and used whatever features were available in each data set.

6.2 Results from Analysis #3

Equation 9 defined “balance” to be a combination of $\{pd, pf\}$ that decreases if pd decreases or pf increases. As shown in Figure 10, there was very little change in balanced performance after learning from 100,200,300,... examples. Indeed, there is some evidence that learning from larger training sets had detrimental effects: the more training data, the larger the variance in the performance of the learned predictor. Observe how, in $kc1$ and $pc1$, as the training set size increases (moving right along the x-axis) the dots showing the balance performance start spreading out. We conjecture that this effect is due to the presence of occasional larger outliers in the training data (the probability of discovering these increases as the training set grows inside and, if the learner trains on them, then the resulting theory has more variance.)

The Mann-Whitney U test was applied to check the visual trends seen in Figure 10. For each project, all results from training sets of size 100,200,300... were compared to all other results from the same project. The issue was “how much data are enough?” i.e. what is the *minimum* training set size that never lost to other training set of a larger size. Usually, that *min* value was quite small:

- In five tables $\{cm1, kc2, kc3, mc2, mw1\}$, $min = 100$;
- In $\{kc1, pc1\}$, $min = \{200, 300\}$ instances, respectively.

We explain the analysis #3 results as follows. This analysis used simplistic static code features such as lines of code, number of unique symbols in the module, etc. Such simplistic static code features are hardly a complete characterization of the internals of a function. We would characterize such static code features as having *limited information content* [27]. Limited content is soon exhausted by repeated sampling. Hence, such simple features reveal all they can reveal after a small sample.

Table 9 An estimate of the effort required to build and test 100 modules.

100 modules may take as little as two to four person months to construct. This estimate was generated as follows:

- In the *cm1* data base, the median module size is 17 lines. 100 randomly selected modules would therefore use 1700 LOC.
- To generate an effort estimate for these modules, we used the on-line COCOMO [33] effort estimator (http://sunset.usc.edu/research/COCOMOII/expert_cocomo/expert_cocomo2000.html). Estimates were generated assuming 1700 LOC and the required reliability varying from very low to very high.
- The resulting estimates ranged from between 2.4 and 3.7 person months to build and test those modules.

6.3 Checking the Analysis #3 Results

There is also some evidence that the results of Analysis #3 (that performance improvements stop after a few hundred examples) have been seen previously in the data mining literature. To the best of our knowledge, this is the first report of this effect in the defect prediction literature.

- In their discussion on how to best handle numeric features, Langley and John offer plots of the accuracy of Naive Bayes classifiers after learning on 10,20,40,...200 examples. In those plots, there is little change in performance after 100 instances [32].
- Orrego [28] applied four data miners (including Naive Bayes) to 20 data sets to find the *plateau point*: i.e. the point after which there was little net change in the performance of the data miner. To find the plateau point, Orrego used t-tests to compare the results of learning from Y or $Y + \Delta$ examples. If, in a 10-way cross-validation, there was no statistical difference between Y and $Y + \Delta$, the plateau point was set to Y . As shown in Figure 11, many of those plateaus were found at $Y \leq 100$ and most were found at $Y \leq 200$. Please note that these plateau sizes are consistent with the results of Analysis #3.

6.4 Discussion of Analysis #3

In the majority of cases, predictors learned from as little as one hundred examples perform as well as predictors learned from many more examples. This suggests that the effort associated with learning defect predictors from within-company data may not be overly large. For example, Table 9 estimates that the effort required to build and test 100 modules may be as little as 2.4 to 3.7 person months. However, practitioners should use this approach cautiously. The populations of one hundred examples in our experiments are randomly selected from completed projects *with* stratification. Therefore, in practice, any one hundred sample may not necessarily reflect the company characteristics and constructing this initial set may take longer than expected.

7 Replication: Can our results be generalized?

Analyses # 1 to # 3 were based on NASA projects. For the external validity of the conclusions of those Analyses, we replicate the same analyses on SOFTLAB projects of Table 1.

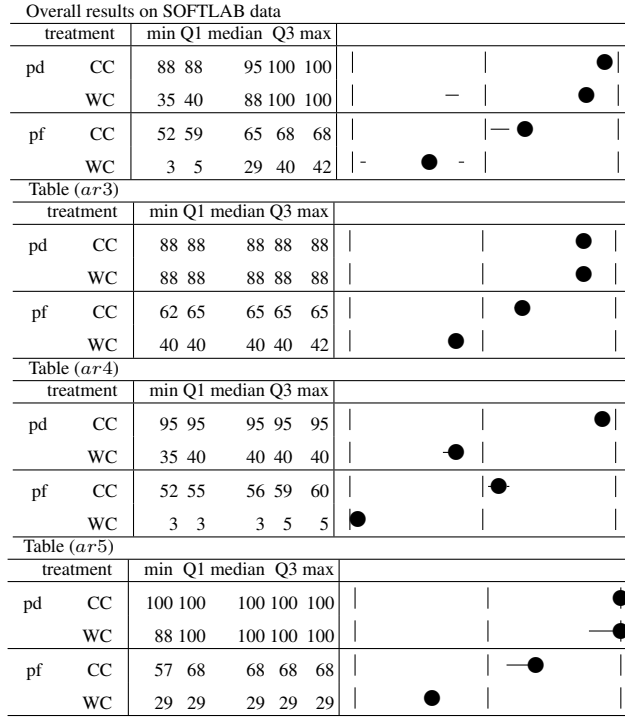


Fig. 12 Analysis #1 results for the SOFTLAB projects. Overall and individual results are shown respectively.

For each SOFTLAB project, we follow the same procedure as in Analyses # 1 and #2; i.e. 10% of the rows of each data set are selected at random for constructing test sets. Then three different types of defect predictors are constructed. First type are defect predictors trained with cross company data (i.e. all 7 NASA projects). Second type of defect predictors are trained with within company data (i.e. random 90% rows of remaining SOFTLAB projects¹⁰). Finally, the third type are defect predictors trained with nearest neighbor filtered cross company data (i.e. *similar* rows from 7 NASA tables).

The SOFTLAB projects include 29 static code features, 17 of which are common with the NASA projects. In order to simplify the comparison between these new projects and Analyses #1 and #2, we used only these shared attributes in our CC analyses. On the other hand we use all available features in WC analyses for SOFTLAB projects. In the following external validity analysis, we treated each NASA project as cross- company data for SOFTLAB projects.

Figure 12 shows the results:

- The *pd* values for CC data increase compared to WC data with the cost of increased *pf*.
- CC data shifts {Q1, median} of *pf* from {5, 29} to {59, 65}.

¹⁰ In order to reflect the use in practice, we do not use the remaining 90% of the same project for training, we rather use a random 90% of data from other projects. Please note that all WC analysis in this paper reflects within-company, not within project simulations. Since SOFTLAB data are collected from a single company, learning a predictor on some projects and to test it on a different one does not violate within company simulation.




Fig. 13 Analysis #2 PD results for the SOFTLAB projects where $NN_{pd} \geq WC_{pd}$. Rankings computed via Mann-Whitney (95% confidence) comparing each row to all other rows.




Fig. 14 Analysis #2 PF results for the SOFTLAB projects where $NN_{pf} \leq WC_{pf}$.




Fig. 15 Analysis #2 PF results for the SOFTLAB projects where $NN_{pf} > WC_{pf}$.

- For CC data:
 - 25% of the pd values are at 100%.
 - 50% of the pd values are above 95%
 - And all the pd values are at or over 88%.

These results also provide evidence for the validity of our conclusions for Analysis # 3. In Analysis # 3, we conclude that the minimum number of instances for training a defect predictor is around 100 – 200 data samples. Please note that three SOFTLAB projects *ar3*, *ar4* and *ar5* have $\{63, 107, 36\}$ modules respectively. Thus, the minimum number of training samples occurs when a predictor is trained from (*ar3* + *ar5*) projects. In this case, only $(63 + 36) * 0.90 = 90$ training samples are used to construct a defect predictor for the *ar4* project. Similarly, the maximum number of training samples occurs when a predictor is trained from (*ar3* + *ar4*) projects. Then $(63 + 107) * 0.90 = 153$ training samples are used to construct a defect predictor for the *ar5* project. Therefore, the WC results in Figure 12 are achieved using a minimum of 90 and a maximum of 153 training samples.

Figure 13 to Figure 15 shows *Observation1* (i.e. $NN_{pd} \geq WC_{pd}$) and *Observation2* (i.e. $NN_{pf} \leq WC_{pf}$) for SOFTLAB projects. Please recall that these observations were mutually exclusive for NASA data. The pattern is similar in SOFTLAB projects:

- for *ar4* mutual exclusiveness hold: $NN_{pd} \geq WC_{pd}$ and $NN_{pf} > WC_{pf}$
- for *ar3* and *ar5*: $NN_{pf} \leq WC_{pf}$. If the observations were mutually exclusive, we would expect $NN_{pd} < WC_{pd}$. While this is not the case for *pd*, we observe that $NN_{pd} = WC_{pd}$ and $NN_{pd} \not\geq WC_{pd}$ (see Figure 13).

In summary, the WC, CC and NN patterns found in American NASA rocket software are also observed in software controllers of Turkish domestic appliances. While this is not the definitive proof of the external validity of our results, we find it a very compelling result that is reproducible in different companies.

8 Related Work

8.1 Neighborhood Reasoning

Our use of nearest neighbor is somewhat similar to *k-NN classifier* or *reasoning by analogy*. A *k-NN classifier* returns the majority class in the “k” nearest neighbors to the test instance. Lessman et.al. benchmark k-NN against 19 other learners and found that it was one of the four worst methods for defect prediction [17].

In *reasoning by analogy*, inference about some new case is made by partial match to prior cases, then applying some repair or generalization procedure to the matched cases. Like k-NN, reasoning by analogy makes use of some distance metric that finds old cases similar (i.e. nearer) to the new case. While reasoning by analogy is a common effort estimation method (e.g. [34]), it is far less common to see it applied to defect prediction. One reason for this is the typical size of the training set: effort estimation typically deals with dozens of prior examples while defect training sets can have thousands of examples. We conjecture that as the data set grows in size, simple induction over many instances may be more powerful than complex analogical reasoning over a handful of examples.

Our work is most similar to that of Khoshgoftaar and Seliya [35] who augmented k-NN with a counting procedure that was a slightly more sophisticated decision procedure than “return the majority class”. Our approach differs to theirs in that we bundle together the k-nearest neighbors, then pass that bundle to a full-fledged classifier (Naive Bayes).

8.2 On Defect Prediction Using Static Code Features

We study defect predictors learned from static code attributes since they are *useful*, *easy to collect*, and *widely-used*.

Useful: Defect predictors are considered useful, if they provide a prediction performance that is comparable to or better than manual reviews [36]. This paper finds defect predictors with a probability of detection of 80%, or higher. This is higher than currently-used industrial methods such as manual code reviews:

- A *IEEE Metrics 2002* [37] panel concluded that manual software reviews can find $\approx 60\%$ of defects¹¹

¹¹ That panel supported neither Fagan’s claim [38] that inspections can find 95% of defects before testing or Shull’s claim that specialized directed inspection methods can catch 35% more defects than other methods [39].

Table 10 Some representative pds and pfs for prediction problems from the UC Irvine machine learning database [31]. These values were generated using the standard settings of a state-of-art decision tree learner (J48). For each data set, ten experiments were conducted, where a decision tree was learned on 90% of the data, then tests are done of the remaining 10%. The numbers shown here are the average results across ten such experiments.

data	probability of	
	detection	false alarm
pima diabetes	60	19
sonar	71	29
horse-colic	71	7
heart-statlog	73	21
rangeseg	76	30
credit rating	88	16
sick	88	1
hepatitis	94	56
vote	95	3
ionosphere	96	18
mean	81	20

- In 2004, Raffo (personnel communication) reports that the defect detection capability of industrial review methods can vary from probability of detection: $pd = TR(35, 50, 65)\%$ ¹². for full Fagan inspections [40] to $pd = TR(13, 21, 30)\%$ for less-structured inspections.

Further, defect predictors based on static code features are considered as static analysis, since they do not require the execution of code. Zheng et.al. compares automated static analysis (ASA) with manual inspections on Nortel software and conclude that *...our results indicate that ASA is an economical complement to other verification and validation techniques*

Easy to collect: Employing defect predictors in practice should not take too much time [36] for data collection and constructing the models themselves. An advantage of static code features is that they can be quickly and automatically collected from the source code, even if no other information is available. Static code attributes like lines of code and the McCabe/Halstead attributes can be automatically and cheaply collected, even for very large systems [41]. By contrast, other methods such as manual code reviews are labor-intensive; e.g. Each member of a review team can inspect 8 to 20 LOC/minute [42]. Furthermore, other features (e.g. number of developers, the software development practices used to develop the code) may be unavailable or hard to characterize.

Widely used: Many researchers use static features to guide software quality predictions (see [41, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54]). Verification and validation (V&V) textbooks advise using static code features to select modules worthy of manual inspections [55].

Nevertheless, there are many reasons to doubt the value of static code attributes for defect prediction. Static code attributes are hardly a complete characterization of the internals of a function. Fenton offers an insightful example where *the same* functionality is achieved using *different* programming language constructs resulting in *different* static measurements for that module [56]. Fenton uses this example to argue the uselessness of static code attributes. Further, Fenton & Pfleeger note that the main McCabe’s attribute (cyclomatic complexity, or $v(g)$) is highly correlated with lines of code [56]. Shepperd & Ince repeated that result, commenting that “for a large class of software it (cyclomatic complexity) is no more than a proxy for, and in many cases outperformed by, lines of code” [16].

¹² $TR(a, b, c)$ is a triangular distribution with min/mode/max of a, b, c .

However, if the above criticisms are correct then we would predict that, in general, the performance of a defect predictor learned by a data miner should be very poor. More specifically, the supposedly better static code attributes such as Halstead and McCabe should perform no better than just simple thresholds on lines of code.

Neither of these predictions are true, at least for the data sets used in this study. The defect predictors learned from static code attributes perform surprisingly well. Formally, learning a defect predictor is a *binary prediction problem* where each module in a project has been labeled as "defect-free" or "defective". The learning problem is to build some predictor which guesses the labels for as-yet-unseen modules. In this paper we find predictors (i.e. for SOFTLAB projects) with a probability of detection (pd) and probability of false alarm (pf) of

$$median(pd) = 88; median(pf) = 29$$

Table 10 lets us compare our new results against standard binary prediction results from the UC Irvine machine learning repository of standard test sets for data miners [31]. Our $median(pd, pf)$ are very close to the standard results of $mean(pd, pf) = (81\%, 20\%)$ which is noteworthy in two ways:

1. It is unexpected. If static code attributes capture so little about source code (as argued by Shepherd, Ince, Fenton and Pfleeger), then we would expect lower probabilities of detection and much higher false alarm rates.
2. These (pd, pf) results are better than currently used industrial methods such as the $pd \approx 60\%$ reported at the 2002 IEEE Metrics panel or the $median(pd) = [21..50]$ reported by Raffo¹³.

9 Practical Implications

While Table 10 shows that our defect detectors work nearly as well as standard data mining methods, it does not necessarily demonstrate that false alarm rates of around 29% are useful in an industrial context. For example, Arisholm and Briand have certain concerns on the practical usage of defect predictors [57]. They argue that if $X\%$ of the modules are predicted to be faulty and if those modules contain less than $X\%$ of the defects, then the costs of generating the defect predictor is not worth the effort.

Let us analyze the testing efforts on MW1 project from Arisholm and Briand's point of view. For MW1, there are a total of 403 modules with 31 defective and 372 defect-free ones. CC model yields 90% pd and 68% pf , and one should examine 280 modules, which is around a 31% reduction in inspection efforts compared to examining all modules. Yet, we argue that 68% pf rate is quite high and using NN we are able to reduce it to 33% along with 68% pd . This corresponds to examining 144 modules, a reduction of 47% compared to exhaustive testing (and we assume an exhaustive test should examine 274 modules for detecting 68% defects, as Arisholm and Briand suggests).

¹³ Please note that we can only compare the defect *detection* properties of automatic vs manual methods. Unlike automatic defect prediction via data mining, the above manual inspection methods don't just report "true,false" on a module. Rather, the manual methods also provide specific debugging information. Hence, a complete comparison of automatic vs manual defect prediction would have to include both an analysis of the time to *detect* potential defects *and* the time required to *fix* them. Manual methods might score higher to automatic methods since they can offer more clues back to the developer about what is wrong with the method. However, such an analysis is beyond the scope of this paper. Here, we focus only on the relative merits of different methods for *predicting* error prone modules.

This analysis can be extended for all projects used in this paper. For instance, the company from which SOFTLAB data in Table 1 are collected is keen to use our detectors, arguing that they operate in a highly competitive market segment where profit margins are very tight. Therefore reducing the cost of the product even by 1% can make a major difference both in market share and profits. Their applications are embedded systems where, over the last decade, the software components have taken precedence over the hardware. Hence their problem is a software engineering problem. According to Brooks [58], half the cost of software development is in unit and systems testing. The company also believes that their main challenge is the testing phase and they seek predictors that indicate where the defects might exist *before* they start testing. Their expectation from the predictor is not to detect all defects, but to guide them to the problematic modules so that they can detect more defects in shorter times. Hence, any reduction in their testing efforts allows them to efficiently use their scarce resources.

Considering the results of our analyses, if a company lacks local data, we would suggest a two-phase approach. In phase one, that organization uses imported CC data, filtered via NN. Also, during phase one, the organization should start a data collection program to collect static code attributes. Phase two commences when there is *enough* local WC data to learn defect predictors. During phase two, the organization would switch to new defect predictors learned from the WC data.

An important issue worth more mentioning is the concern about the time required for setting up a metric program (i.e. in order to collect data for building actual defect predictors). Our incremental WC results suggest that, in the case of defect prediction, this concern may be less than previously believed. Kitchenham *et al.* [59] argue that organizations use cross-company data since within-company data can be so hard to collect:

- The time required to collect enough data on past projects from within a company may be prohibitive.
- Collecting within-company data may take so long that technologies change and older projects do not represent current practice.

In our analysis we observe that as few as 100 modules are enough to learn adequate defect predictors. When so few examples are enough, it is possible that projects can learn local defect predictors that are relevant to their current technology in just a few months.

Further, our experiences with our industry partners show that data collection is not necessarily a major concern. Static code attributes can be automatically and quickly collected with relatively little effort. We have found that when there is high level management commitment, it becomes a relatively simple process. For the three projects of SOFTLAB data, neither the static code attributes, nor the mapping of defects to software modules were available when the authors attempted to collect these data. Since these were smaller scale projects, it was sufficient to spend some time with the developers and going through defect reports. Although not all projects have 100 modules individually, the company has a growing repository from several projects and enough data to perform defect prediction.

We also have experience with a large scale telecommunication company, where a long-term metric program for monitoring complex projects (around 750.000 lines of code) requires introducing automated processes. Again with high level management support, it was possible to employ appropriate tool support and these new processes were introduced easily and invisibly to the staff. For that project, we have now a growing repository of defects mapped with source code (around 25 defects per month). Software in that project are being developed for more than 10 years and have very low defect rates. We have obtained the first results in the 8th month of the project and it is planned to be completed in 12 months. In




Fig. 16 x-axis shows the rank of the data miners given in the y-axis. All the methods whose top ranks are 4 to 12 are statistically insignificantly different. From [17].

summary setting up a metric program for defect prediction can be done more quickly than it is perceived.

10 Threats to validity

In prior work we have explored a range of data mining methods for defect prediction and found that classifiers based on Bayes theorem work best for the Table 1 data [14]. Since that study, we have tried to find better data mining algorithms for defect prediction. To date, we have failed. Other researchers have also failed to improve our results. For example, Lessmann et al. investigate the statistical difference of the results between 19 learners, including naive Bayes, on the NASA projects [17]. Figure 16 shows that the simple Bayesian method discussed above ties in first place along with 15 other methods.

Our analyses do not involve a value-based approach. We do not consider that fixing defects in a module may be more critical than fixing defects in another one. A reason is that our data sources do not consistently include necessary data (i.e. the severity level of defects). If and when data are available, same kind of analysis should be repeated considering this issue.

Another point is that descriptions of software modules *only* in terms of static code attributes can overlook some important aspects of software including: the type of application domain; the skill level of the individual programmers involved in system development; contractor development practices; the variation in measurement practices; and the validation of the measurements and instruments used to collect the data. For this reason some researchers augment, or replace static code measures with other information such as the history of past faults or changes to code or number of developers who have worked on the code [60]. Yet again, we have used all available data in our data sources, which are the static code features of the projects.

The external validity of generalizing from NASA examples has been discussed elsewhere [14]. In summary, NASA uses contractors who are contractually obliged to demonstrate their understanding and usage of current industrial best practices. These contractors service many other industries; for example, Rockwell-Collins builds systems for many gov-

ernment and commercial organizations. For these reasons, other noted researchers such as Basili, Zelkowitz, et al. [61] have argued that conclusions from NASA data are relevant to the general software engineering industry.

Nevertheless, it is always wise to test claims of external validity. Hence:

- The SOFTLAB data were, initially, in reserve. Our first three analyses are based solely on the aerospace applications found in the NASA data.
- Our last analysis checked if the SOFTLAB data exhibit the same pattern as the NASA data.

SOFTLAB data sets were deliberately chosen to be as far removed as possible from American aerospace software applications (i.e. Turkish home appliances controller software). Please note that this software was developed via methods that are both culturally and organizationally different to NASA aerospace software. Turkish domestic appliances company software are developed by a small team of 2-3 people. The development is carried out in an ad-hoc, informal way rather than formal, process oriented approach in NASA. Furthermore, the company is a profit and revenue driven commercial organization, whereas NASA is a cost driven government entity. This implies that our approach is widely applicable among different development practices. More precisely, our approach is independent of the processes that yield the final product, at least for the wide range of projects that we have analyzed.

11 Conclusion

In this study, we have analyzed defect predictors of static code features, constructed by cross company vs. within company data and found clear, unambiguous conclusions:

- CC-data dramatically increase the probability of detecting defective modules (i.e. from median value 75 to 92);
- But CC-data also dramatically increase the false alarm rate (i.e. from median value 29 to 64).
- NN-filtering CC data avoids the high false alarm rates by removing irrelevancies in CC data (i.e. from median value 64 to 32). This removal takes place by automatically selecting similar project data in terms of available static code features and discarding non-similar ones.
- Yet WC-data models are still the best and they can be constructed with small amounts of data (i.e. 100 examples).

Interpreting these in terms of the posed questions in the introduction, we conclude that:

- CC data are useful in extreme cases such as mission critical projects, where the cost of false alarms can be afforded. Therefore, CC data should be used only when WC data are not available.
- Pruning CC data with NN-filter allows the use of CC data for constructing practical defect predictors in other domains. NN-filtered CC data yields much better results than raw CC data, yet closer but worse results than WC data.
- The best option of using WC data requires the collection of a mere hundred examples from within a company and can be done in a short time (i.e. a few months).
- We observe the same patterns not only in aerospace software from NASA, but also in software from a completely different company located in another country. While this is a strong evidence of generality, we take great care *not to* interpret it as a proof of external validity.

We conclude our findings by proposing a two-phase approach for initiating defect prediction process:

- *Phase1*: Use NN filtered CC data to make local predictions and start to collect WC data.
- *Phase2*: After a few hundred examples are available in the local repository (usually a few months), discard the predictor learned on CC data and switch to those learned from WC data.

Acknowledgements This research was supported in part by Turkish Scientific Research Council (TUBITAK) under grant #EEEAG 108E014 and was partially funded by a National Science Foundation grant #CCF-0810879.

References

1. G. Boetticher, T. Menzies, and T. Ostrand, "The PROMISE Repository of Empirical Software Engineering Data," 2007, <http://promisedata.org/repository>.
2. E. Arisholm and L. Briand, "Predicting fault-prone components in a java legacy system," *ISESE '06: Proceedings of the 2006 ACM/IEEE international symposium on empirical software engineering*, Sep 2006. [Online]. Available: <http://portal.acm.org/citation.cfm?id=1159733>.
3. R. Bell, T. Ostrand, and E. Weyuker, "Looking for bugs in all the right places," *ISSTA '06: Proceedings of the 2006 international symposium on Software testing and analysis*, Jul 2006. [Online]. Available: <http://portal.acm.org/citation.cfm?id=1146238>.
4. T. Ostrand, E. Weyuker, and R. Bell, "Automating algorithms for the identification of fault-prone files," *ISSTA '07: Proceedings of the 2007 international symposium on Software testing and analysis*, Jul 2007. [Online]. Available: <http://portal.acm.org/citation.cfm?id=1273463>.
5. T. Menzies, A. Dekhtyar, J. Distefano, and J. Greenwald, "Problems with precision," *IEEE Transactions on Software Engineering*, September 2007, <http://menzies.us/pdf/07precision.pdf>.
6. Z. Chen, T. Menzies, and D. Port, "Feature subset selection can improve software cost estimation," in *PROMISE'05*, 2005, available from <http://menzies.us/pdf/05fsscocomo.pdf>.
7. A. Dekhtyar, J. H. Hayes, and T. Menzies, "Text is software too," in *International Workshop on Mining Software Repositories (submitted)*, 2004, available from <http://menzies.us/pdf/04msrtext.pdf>.
8. Y. Jiang, B. Cukic, and T. Menzies, "Fault prediction using early lifecycle data," in *ISSRE'07*, 2007, available from <http://menzies.us/pdf/07issre.pdf>.
9. N. Nagappan and T. Ball, "Static analysis tools as early indicators of pre-release defect density," in *ICSE 2005, St. Louis*, 2005.
10. A. Marcus, D. Poshyvanyk, and R. Ferenc, "Using the conceptual cohesion of classes for fault prediction in object-oriented systems," *Software Engineering, IEEE Transactions on*, vol. 34, no. 2, pp. 287–300, March–April 2008.
11. J. Musa, A. Iannino, and K. Okumoto, *Software Reliability: Measurement, Prediction, Application*. McGraw Hill, 1987.
12. B. Boehm and P. Papaccio, "Understanding and controlling software costs," *IEEE Trans. on Software Engineering*, vol. 14, no. 10, pp. 1462–1477, October 1988.
13. "Architecture-based software reliability: Why only a few parameters matter?"
14. T. Menzies, J. Greenwald, and A. Frank, "Data mining static code attributes to learn defect predictors," *IEEE Transactions on Software Engineering*, January 2007, available from <http://menzies.us/pdf/06learnPredict.pdf>.
15. J. Dougherty, R. Kohavi, and M. Sahami, "Supervised and unsupervised discretization of continuous features," in *International Conference on Machine Learning*, 1995, pp. 194–202, available from <http://www.cs.pdx.edu/~timm/dm/dougherty95supervised.pdf>.
16. M. Shepperd and D. Ince, "A critique of three metrics," *The Journal of Systems and Software*, vol. 26, no. 3, pp. 197–210, September 1994.
17. S. Lessmann, B. Baesens, C. Mues, and S. Pietsch, "Benchmarking classification models for software defect prediction: A proposed framework and novel findings," *accepted for publication IEEE Transactions on Software Engineering*, 2009.
18. R. Duda, P. Hart, and N. Nilsson, "Subjective bayesian methods for rule-based inference systems," in *Technical Report 124, Artificial Intelligence Center, SRI International*, 1976.

19. P. Domingos and M. J. Pazzani, "On the optimality of the simple bayesian classifier under zero-one loss," *Machine Learning*, vol. 29, no. 2-3, pp. 103–130, 1997. [Online]. Available: citeseer.ist.psu.edu/domingos97optimality.html
20. Y. Yang and G. Webb, "Weighted proportional k-interval discretization for naive-bayes classifiers," in *Proceedings of the 7th Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD 2003)*, 2003, available from <http://www.csse.monash.edu/~webb/Files/YangWebb03.pdf>.
21. I. H. Witten and E. Frank, *Data mining. 2nd edition*. Los Altos, US: Morgan Kaufmann, 2005.
22. H. B. Mann and D. R. Whitney, "On a test of whether one of two random variables is stochastically larger than the other," *Ann. Math. Statist.*, vol. 18, no. 1, pp. 50–60, 1947, available online at <http://projecteuclid.org/DPubS?service=UI&version=1.0&verb=Display&hand%le=euclid.aoms/1177730491>.
23. J. Demsar, "Statistical comparisons of classifiers over multiple data sets," *Journal of Machine Learning Research*, vol. 7, pp. 1–30, 2006, available from <http://jmlr.csail.mit.edu/papers/v7/demsar06a.html>.
24. A. G. Koru, K. E. Emam, D. Zhang, H. Liu, and D. Mathew, "Theory of relative defect proneness," *Empirical Softw. Engg.*, vol. 13, no. 5, pp. 473–498, 2008.
25. J. H. Hayes, A. Dekhtyar, and S. K. Sundaram, "Advancing candidate link generation for requirements tracing: The study of methods," *IEEE Trans. Software Eng.*, vol. 32, no. 1, pp. 4–19, 2006. [Online]. Available: <http://doi.ieeecomputersociety.org/10.1109/TSE.2006.3>
26. D. Baker, "A hybrid approach to expert and model-based effort estimation," Master's thesis, Lane Department of Computer Science and Electrical Engineering, West Virginia University, 2007, available from <https://eidr.wvu.edu/etd/documentdata.eTD?documentid=5443>.
27. T. Menzies, B. Turhan, A. Bener, G. Gay, B. Cukic, and Y. Jiang, "Implications of ceiling effects in defect predictors," in *Proceedings of PROMISE 2008 Workshop (ICSE)*, 2008, available from <http://menzies.us/pdf/08ceiling.pdf>.
28. A. Orrego, "Sawtooth: Learning from huge amounts of data," Master's thesis, Computer Science, West Virginia University, 2004.
29. R. Quinlan, *C4.5: Programs for Machine Learning*. Morgan Kaufman, 1992, ISBN: 1558602380.
30. J. R. Quinlan, "Learning with Continuous Classes," in *5th Australian Joint Conference on Artificial Intelligence*, 1992, pp. 343–348, available from <http://citeseer.nj.nec.com/quinlan92learning.html>.
31. C. Blake and C. Merz, "UCI repository of machine learning databases," 1998, URL: <http://www.ics.uci.edu/~mlearn/MLRepository.html>.
32. G. John and P. Langley, "Estimating continuous distributions in bayesian classifiers," in *Proceedings of the Eleventh Conference on Uncertainty in Artificial Intelligence Montreal, Quebec: Morgan Kaufmann*, 1995, pp. 338–345, available from <http://citeseer.ist.psu.edu/john95estimating.html>.
33. B. Boehm, "Safe and simple software cost analysis," *IEEE Software*, pp. 14–17, September/October 2000, available from <http://www.computer.org/certification/beta/Boehm.Safe.pdf>.
34. M. Shepperd and C. Schofield, "Estimating software project effort using analogies," *IEEE Transactions on Software Engineering*, vol. 23, no. 12, November 1997, available from <http://www.utdallas.edu/~rbanker/SE.XII.pdf>.
35. T. M. Khoshgoftaar and N. Seliya, "Analogy-based practical classification rules for software quality estimation," *Empirical Softw. Engg.*, vol. 8, no. 4, pp. 325–350, 2003.
36. R. Moser, W. Pedrycz, and G. Succi, "A comparative analysis of the efficiency of change metrics and static code attributes for defect prediction," *ICSE '08: Proceedings of the 30th international conference on Software engineering*, May 2008. [Online]. Available: <http://portal.acm.org/citation.cfm?id=1368088.1368114>
37. F. Shull, V. Basili, B. Boehm, A. Brown, P. Costa, M. Lindvall, D. Port, I. Rus, R. Tesoriero, and M. Zelkowitz, "What we have learned about fighting defects," in *Proceedings of 8th International Software Metrics Symposium, Ottawa, Canada*, 2002, pp. 249–258, available from http://fc-md.umd.edu/fcmd/Papers/shull_defects.ps.
38. M. Fagan, "Advances in software inspections," *IEEE Trans. on Software Engineering*, pp. 744–751, July 1986.
39. F. Shull, I. Rus, and V. Basili, "How perspective-based reading can improve requirements inspections," *IEEE Computer*, vol. 33, no. 7, pp. 73–79, 2000, available from <http://www.cs.umd.edu/projects/SoftEng/ESEG/papers/82.77.pdf>.
40. M. Fagan, "Design and code inspections to reduce errors in program development," *IBM Systems Journal*, vol. 15, no. 3, 1976.
41. N. Nagappan and T. Ball, "Static analysis tools as early indicators of pre-release defect density," in *ICSE*, 2005, pp. 580–586. [Online]. Available: <http://doi.acm.org/10.1145/1062558>
42. T. Menzies, D. Raffo, S. on Setamanit, Y. Hu, and S. Tootoonian, "Model-based tests of truisms," in *Proceedings of IEEE ASE 2002*, 2002, available from <http://menzies.us/pdf/02truisms.pdf>.
43. M. Halstead, *Elements of Software Science*. Elsevier, 1977.

44. T. McCabe, "A complexity measure," *IEEE Transactions on Software Engineering*, vol. 2, no. 4, pp. 308–320, Dec. 1976.
45. M. Chapman and D. Solomon, "The relationship of cyclomatic complexity, essential complexity and error rates," 2002, proceedings of the NASA Software Assurance Symposium, Coolfont Resort and Conference Center in Berkley Springs, West Virginia. Available from <http://www.ivv.nasa.gov/business/research/osmasas/conclusion2002/Mike.C%20hchapman.The.Relationship.of.Cyclomatic.Complexity.Essential.Complexity.and.Erro%r.Rates.ppt>.
46. T. Menzies, J. DiStefano, A. Orrego, and R. Chapman, "Assessing predictors of software defects," in *Proceedings, workshop on Predictive Software Models, Chicago*, 2004, available from <http://menzies.us/pdf/04psm.pdf>.
47. "Polyspace verifier[®]," 2005, available from http://www.di.ens.fr/~cousot/projects/DAEDALUS/synthetic_summary/POLYSP%ACE/polyspace-daedalus.htm.
48. G. Hall and J. Munson, "Software evolution: code delta and code churn," *Journal of Systems and Software*, pp. 111 – 118, 2000.
49. A. Nikora and J. Munson, "Developing fault predictors for evolving software systems," in *Ninth International Software Metrics Symposium (METRICS'03)*, 2003.
50. T. Khoshgoftaar and N. Seliya, "Comparative assessment of software quality classification techniques: An empirical case study," *Empirical Software Engineering*, vol. 9, no. 3, pp. 229–257, 2004.
51. W. Tang and T. M. Khoshgoftaar, "Noise identification with the k-means algorithm," in *ICTAI*, 2004, pp. 373–378. [Online]. Available: <http://doi.ieeecomputersociety.org/10.1109/ICTAI.2004.93>
52. A. Porter and R. Selby, "Empirically guided software development using metric-based classification trees," *IEEE Software*, pp. 46–54, March 1990.
53. J. Tian and M. Zelkowitz, "Complexity measure evaluation and selection," *IEEE Transaction on Software Engineering*, vol. 21, no. 8, pp. 641–649, Aug. 1995.
54. K. Srinivasan and D. Fisher, "Machine learning approaches to estimating software development effort," *IEEE Trans. Soft. Eng.*, pp. 126–137, February 1995.
55. S. Rakitin, *Software Verification and Validation for Practitioners and Managers, Second Edition*. Artech House, 2001.
56. N. E. Fenton and S. Pfleeger, *Software Metrics: A Rigorous & Practical Approach (second edition)*. International Thompson Press, 1995.
57. E. Arisholm and L. Briand, "Predicting fault-prone components in a java legacy system," in *5th ACM-IEEE International Symposium on Empirical Software Engineering (ISESE)*, Rio de Janeiro, Brazil, September 21-22, 2006, available from <http://simula.no/research/engineering/publications/Arisholm.2006.4>.
58. F. P. Brooks, *The Mythical Man-Month, Anniversary edition*. Addison-Wesley, 1995.
59. B. A. Kitchenham, E. Mendes, and G. H. Travassos, "Cross- vs. within-company cost estimation studies: A systematic review," *IEEE Transactions on Software Engineering*, pp. 316–329, May 2007.
60. T. L. Graves, A. F. Karr, J. S. Marron, and H. P. Siy, "Predicting fault incidence using software change history," *IEEE Trans. Software Eng.*, vol. 26, no. 7, pp. 653–661, 2000, available on-line at www.niss.org/technicalreports/tr80.pdf.
61. V. Basili, F. McGarry, R. Pajerski, and M. Zelkowitz, "Lessons learned from 25 years of process improvement: The rise and fall of the NASA software engineering laboratory," in *Proceedings of the 24th International Conference on Software Engineering (ICSE) 2002, Orlando, Florida*, 2002, available from <http://www.cs.umd.edu/projects/SoftEng/ESEG/papers/83.88.pdf>.